

```
In [1]:  
import os  
import pathlib  
import time  
import re  
import pickle  
import numpy as np  
import pandas as pd  
import matplotlib  
import matplotlib.pyplot as plt  
import seaborn as sns  
import tensorflow as tf  
from tensorflow import keras  
from tensorflow.keras import layers  
print(tf.__version__)  
  
# plotting style options  
font = {'size': 18}  
matplotlib.rc('font', **font)  
sns.set_theme()  
sns.set(font_scale = 1.2)  
%matplotlib inline
```

2.7.0

# Contents

The notebook is structured as follows:

1. Introduction
  - 1.1. Problem statement
  - 1.2. Proposed solution
1. Exploration of the dataset
  - 2.1. Correlations
  - 2.2. Time series plots
  - 2.3. Histograms
1. Pre-processing
  - 3.1. Addition of higher-order derivatives of input features (velocity, acceleration, jerk, snap, crackle, pop)
  - 3.2. Scaling of input features and outputs
  - 3.3. Preparing input features for RNN training
  - 3.4. Splitting of data into train/validation/test sets
1. Modeling
  - 4.1. Linear regression
  - 4.2. DNN regression

- 4.3. RNN regression
- 1. Evaluation
  - 5.1. Loss on test sets
  - 5.2. Prediction error
  - 5.3. Time series plots of prediction vs. ground truth
- 1. Summary and prospects

# 1. Introduction

## 1.1. Problem statement

- Two opposing robots, R1 and R2, apply pressure onto a metal sheet to deform it.
- Each robot can move in  $x, y, z$ , and (roll, pitch, yaw)  $a, b, c$ .
- R1 (the "forming robot") pushes into the sheet ( $z$ ) while moving along a path in  $x/y$ -plane, parallel to the sheet metal surface.
- R2 (the opposing "support robot") pushes into the sheet from the other side.
- Both robots experience forces at the tool-tip in  $x, y$ , and  $z$  directions.
- Positions, angles, and tool-tip forces are captured for each robot.

**Deliverables:**

- Develop a model to predict forces experienced by tool-tips for any path traveled by the robot arms
- Provide a Jupyter notebook that includes all the steps

## 1.2. Proposed solution

The problem is to find a function  $G$  to predict tool-tip forces from position coordinates:

$$\vec{f} = G(x, y, z, a, b, c)$$

Models like these could be solved using a parametrized dynamical model with Lagrange or Newton-Euler methods assuming rigid body motion. However, instead of doing that, I'm going to **approximate  $G$  with a neural network**.

Here are a few sources I looked at for inspiration:

- [http://www.scholarpedia.org/article/Robot\\_dynamics](http://www.scholarpedia.org/article/Robot_dynamics)
- <https://ethz.ch/content/dam/ethz/special-interest/mavt/robotics-n-intelligent-systems/rsl-dam/documents/RobotDynamics2016/6-dynamics.pdf>
- <https://towardsdatascience.com/approximating-dynamic-models-of-industrial-robots-with-neural-networks-2474d1a2eecd>
- <https://www.nature.com/articles/s41598-021-97003-1.pdf>

- <https://www.tensorflow.org/tutorials/keras/regression>
- <https://www.tensorflow.org/guide/keras/rnn>

## 2. Exploration of the dataset

Before training a model, it is always good to look at the data. Let's have a look at time series plots, histograms, and correlations.

### Import the datasets into pd.DataFrame

In [124...]

```
workdir = pathlib.Path("./")

# create a timestamp for tagging saved models
timestamp = time.strftime("%Y%m%d_%H%M")

# output directory for saved models, plots, etc.
output_dir = workdir / 'olsson_solution_{}'.format(timestamp)
output_dir.mkdir(parents=True, exist_ok=True)

dataset_filenames = ["Test1", "Test2", "Test4"]

# list of dataframes
datasets = list()
print("loading datasets:")
for filename in dataset_filenames:
    print("-", workdir / str(filename+'.csv'))
    datasets.append(pd.read_csv(workdir / str(filename+'.csv')))
```

loading datasets:

- Test1.csv
- Test2.csv
- Test4.csv

### Look at a summary of what's in the dataframe

In [3]:

```
print("First five entries: \n", datasets[0].head(5).T)
```

First five entries:

	0	1	2	3	4
t	1.636580e+09	1.636580e+09	1.636580e+09	1.636580e+09	1.636580e+09
a_enc_1	-4.951100e+00	-4.951100e+00	-4.951100e+00	-4.951100e+00	-4.951100e+00
b_enc_1	1.830000e-02	1.830000e-02	1.830000e-02	1.830000e-02	1.830000e-02
c_enc_1	-7.190000e-02	-7.190000e-02	-7.190000e-02	-7.190000e-02	-7.190000e-02
x_enc_1	2.136337e+02	2.136337e+02	2.136337e+02	2.136337e+02	2.136337e+02
y_enc_1	3.241015e+02	3.241015e+02	3.241015e+02	3.241015e+02	3.241015e+02
z_enc_1	8.953528e+02	8.953528e+02	8.953528e+02	8.953528e+02	8.953528e+02
a_enc_2	-1.549772e+02	-1.549772e+02	-1.549772e+02	-1.549772e+02	-1.549772e+02
b_enc_2	2.023000e-01	2.024000e-01	2.024000e-01	2.024000e-01	2.024000e-01
c_enc_2	-1.798798e+02	-1.798798e+02	-1.798798e+02	-1.798798e+02	-1.798798e+02
x_enc_2	2.232210e+01	2.232040e+01	2.232040e+01	2.232040e+01	2.232040e+01
y_enc_2	7.831761e+02	7.831754e+02	7.831754e+02	7.831754e+02	7.831754e+02
z_enc_2	-7.725771e+02	-7.725771e+02	-7.725771e+02	-7.725771e+02	-7.725771e+02
fx_1	-2.326357e+00	-2.192611e+00	-2.103594e+00	-1.869649e+00	-2.336206e+00
9.639795e+00	9.531656e+00	9.776526e+00	9.100982e+00	9.058406e+00	
fz_1	-3.264595e+01	-3.307391e+01	-3.143578e+01	-3.171914e+01	-3.232948e+01

```
fx_2      1.180561e+01  1.169716e+01  1.166217e+01  1.141468e+01  1.122329e+01
fy_2      1.865609e+01  1.846252e+01  1.860119e+01  1.848982e+01  1.795298e+01
fz_2     -1.283101e+01 -1.225022e+01 -1.145559e+01 -1.253816e+01 -1.042543e+01
```

In [4]:

```
print("\ndescribe():\n", datasets[0].describe().T[['count', 'mean', 'std', 'min', 'max'])

describe():
   count          mean           std          min          max
t      20091.0  1.636590e+09  5800.523780  1.636580e+09  1.636600e+09
a_enc_1  20091.0 -8.924334e+01   7.811876 -9.001034e+01 -4.951004e+00
b_enc_1  20091.0  8.760533e-04   0.002961 -1.604445e-02  2.338158e-02
c_enc_1  20091.0  1.884124e-03   0.007390 -7.199558e-02  2.159353e-02
x_enc_1  20091.0  4.578387e+02  197.669145  8.937531e+01  8.306894e+02
y_enc_1  20091.0  1.965825e+02  103.805836 -1.773652e+00  3.672685e+02
z_enc_1  20091.0 -6.500421e+01  100.980291 -1.767849e+02  8.953528e+02
a_enc_2  20091.0  8.804365e+01   22.349876 -1.783886e+02  1.789519e+02
b_enc_2  20091.0  1.449096e-03   0.029150 -2.199676e+00  1.164570e+00
c_enc_2  20091.0 -4.858466e+01  173.319815 -1.800000e+02  1.800000e+02
x_enc_2  20091.0  4.560382e+02  202.637762  2.231891e+01  8.318730e+02
y_enc_2  20091.0  1.995663e+02  120.025141 -3.800164e+00  7.831773e+02
z_enc_2  20091.0 -7.605120e+01   81.255998 -7.725771e+02 -6.817898e-01
fx_1    20091.0  3.084250e+01   681.262919 -1.919500e+03  1.876367e+03
fy_1    20091.0  9.071373e+01  1192.944410 -1.841587e+03  2.238735e+03
fz_1    20091.0  2.518588e+03   582.206105 -7.799607e+01  3.374799e+03
fx_2    20091.0 -2.917874e+01   441.132354 -1.488688e+03  1.233724e+03
fy_2    20091.0 -1.031784e+02   729.100519 -1.584797e+03  1.303751e+03
fz_2    20091.0 -7.406174e+02   302.882751 -2.315789e+03 -5.323892e+00
```

**Make separate lists of input features (position coordinates and euler angles) and outputs (forces)**

In [5]:

```
#features = [k for k in datasets[0].keys() if not re.search('^f', k) and 't' not in k]
#outputs = [k for k in datasets[0].keys() if re.search('^f', k)]

features_1 = [i+'_enc_1' for i in 'xyzabc']
features_2 = [i+'_enc_2' for i in 'xyzabc']
outputs_1 = ['fx_1', 'fy_1', 'fz_1']
outputs_2 = ['fx_2', 'fy_2', 'fz_2']
features = features_1 + features_2
outputs = outputs_1 + outputs_2

print('features:', features)
print('outputs:', outputs)
```

```
features: ['x_enc_1', 'y_enc_1', 'z_enc_1', 'a_enc_1', 'b_enc_1', 'c_enc_1', 'x_enc_2', 'y_enc_2', 'z_enc_2', 'a_enc_2', 'b_enc_2', 'c_enc_2']
outputs: ['fx_1', 'fy_1', 'fz_1', 'fx_2', 'fy_2', 'fz_2']
```

## 2.1. Correlations

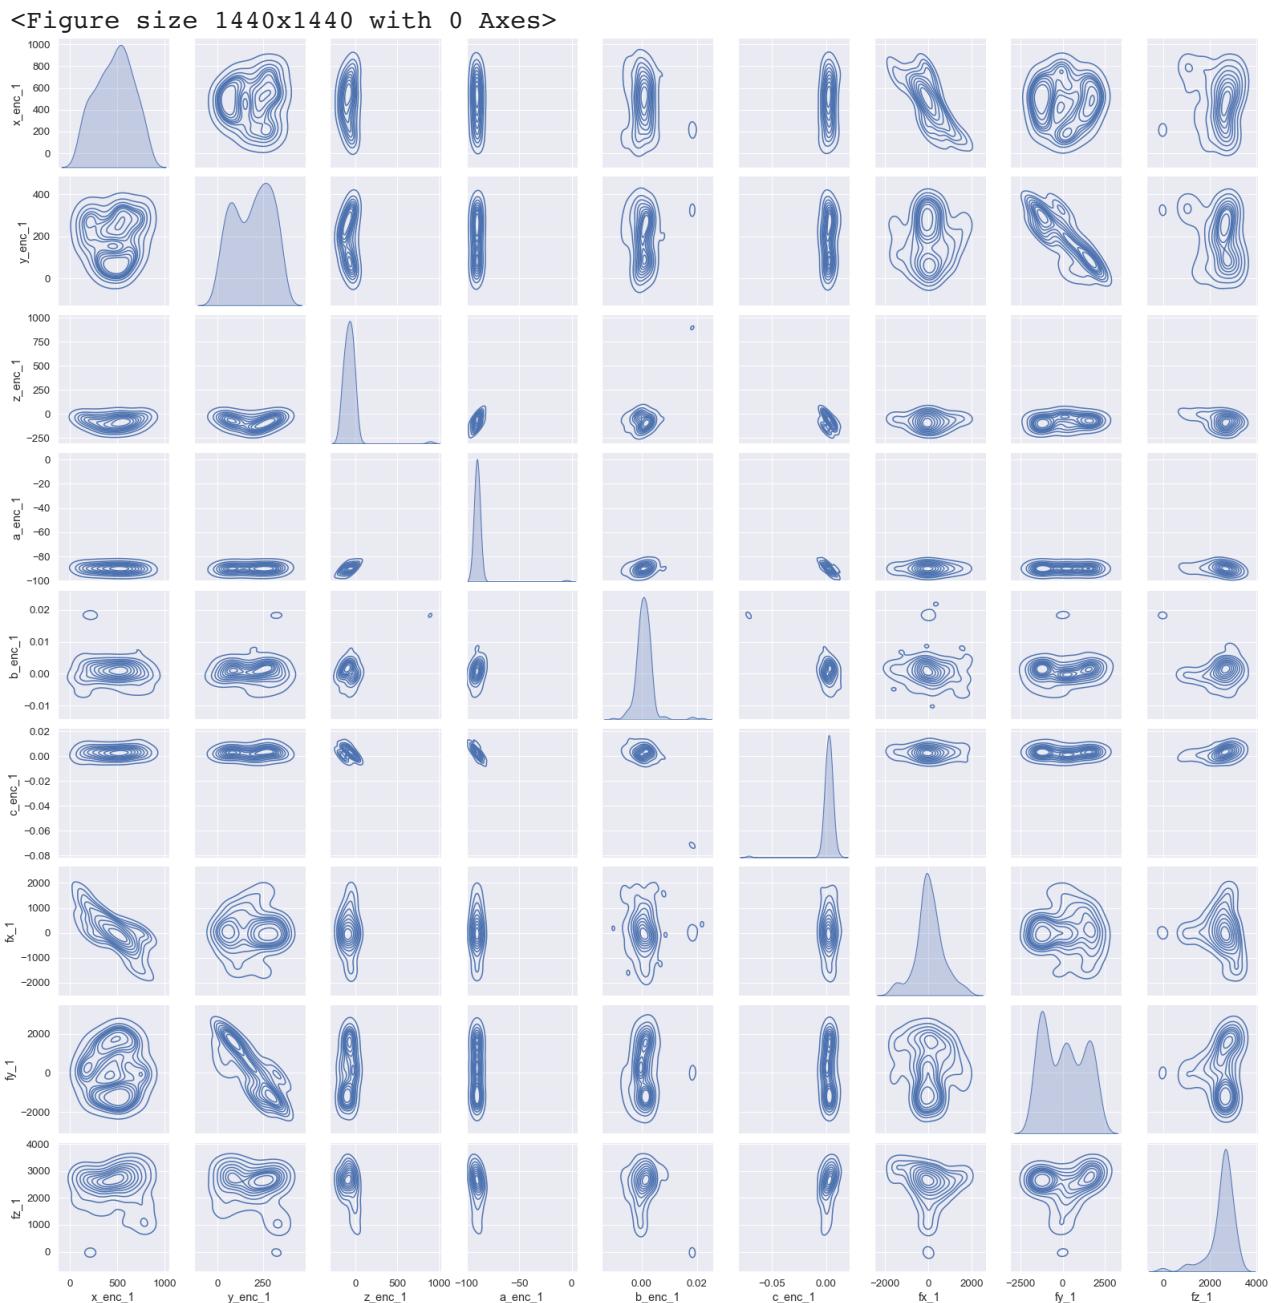
**First look at correlations for R1 and R2 separately, for simplicity**

seaborn.pairplot can be helpful to get an idea of correlations

In [6]:

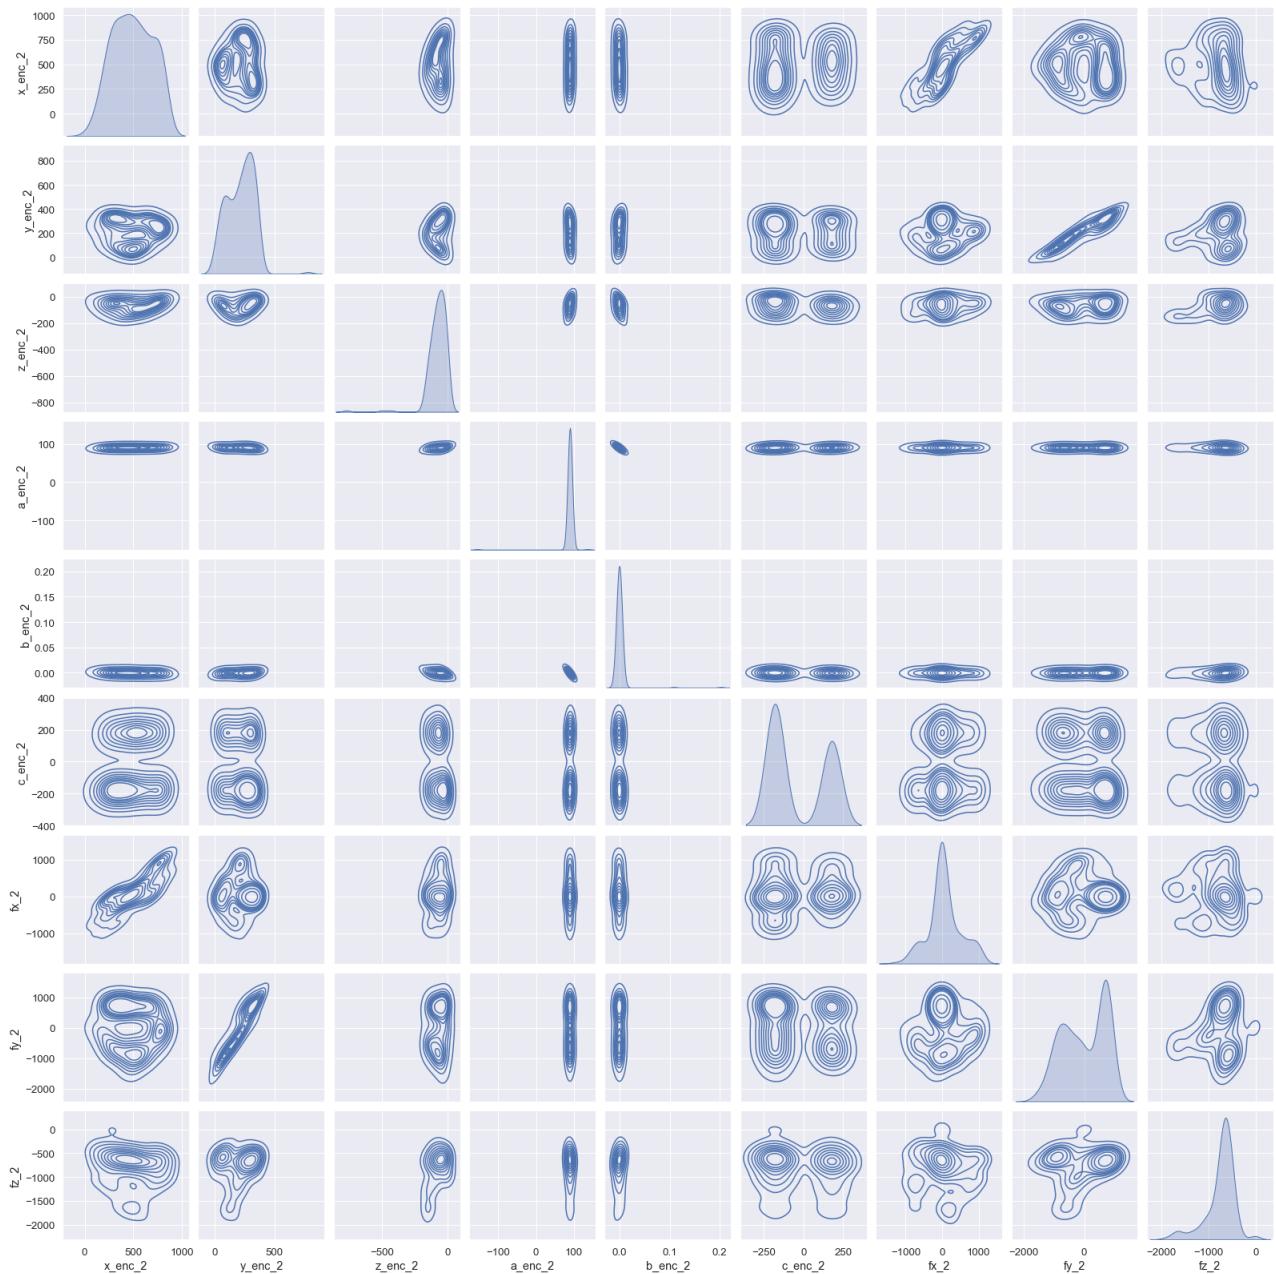
```
# pair plots (slow...)
labels = [features_1+outputs_1, features_2+outputs_2]
for robot_idx in range(2):
    fig = plt.figure(figsize=(20, 20))
```

```
sns.pairplot(datasets[0].sample(200)[labels[robot_idx]], kind='kde')
plt.savefig(output_dir/'pairplot_robot{}.pdf'.format(robot_idx+1))
```



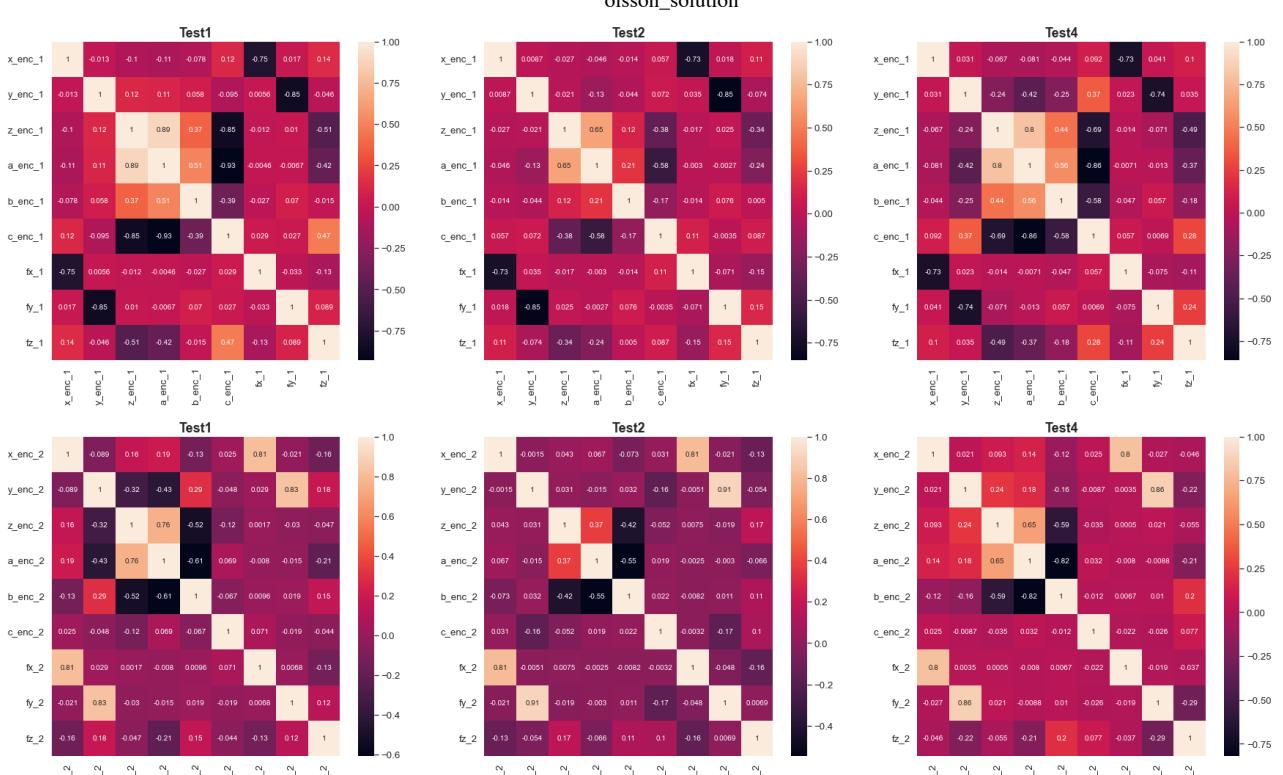
<Figure size 1440x1440 with 0 Axes>

## olsson\_solution



## Plot correlation matrices

```
In [7]:=
sns.set_style("whitegrid")
labels = [features_1+outputs_1, features_2+outputs_2]
for robot_idx in range(2):
    fig = plt.figure(figsize=(10*len(datasets), 8))
    for i,df in enumerate(datasets):
        corr = df[labels[robot_idx]].corr()
        ax = fig.add_subplot(1, len(datasets), i+1)
        ax.set_title(dataset_filenames[i], weight='bold').set_fontsize('18')
        sns.heatmap(corr, annot=True)
    plt.savefig(output_dir/'corr_robot{}.pdf'.format(robot_idx+1))
```



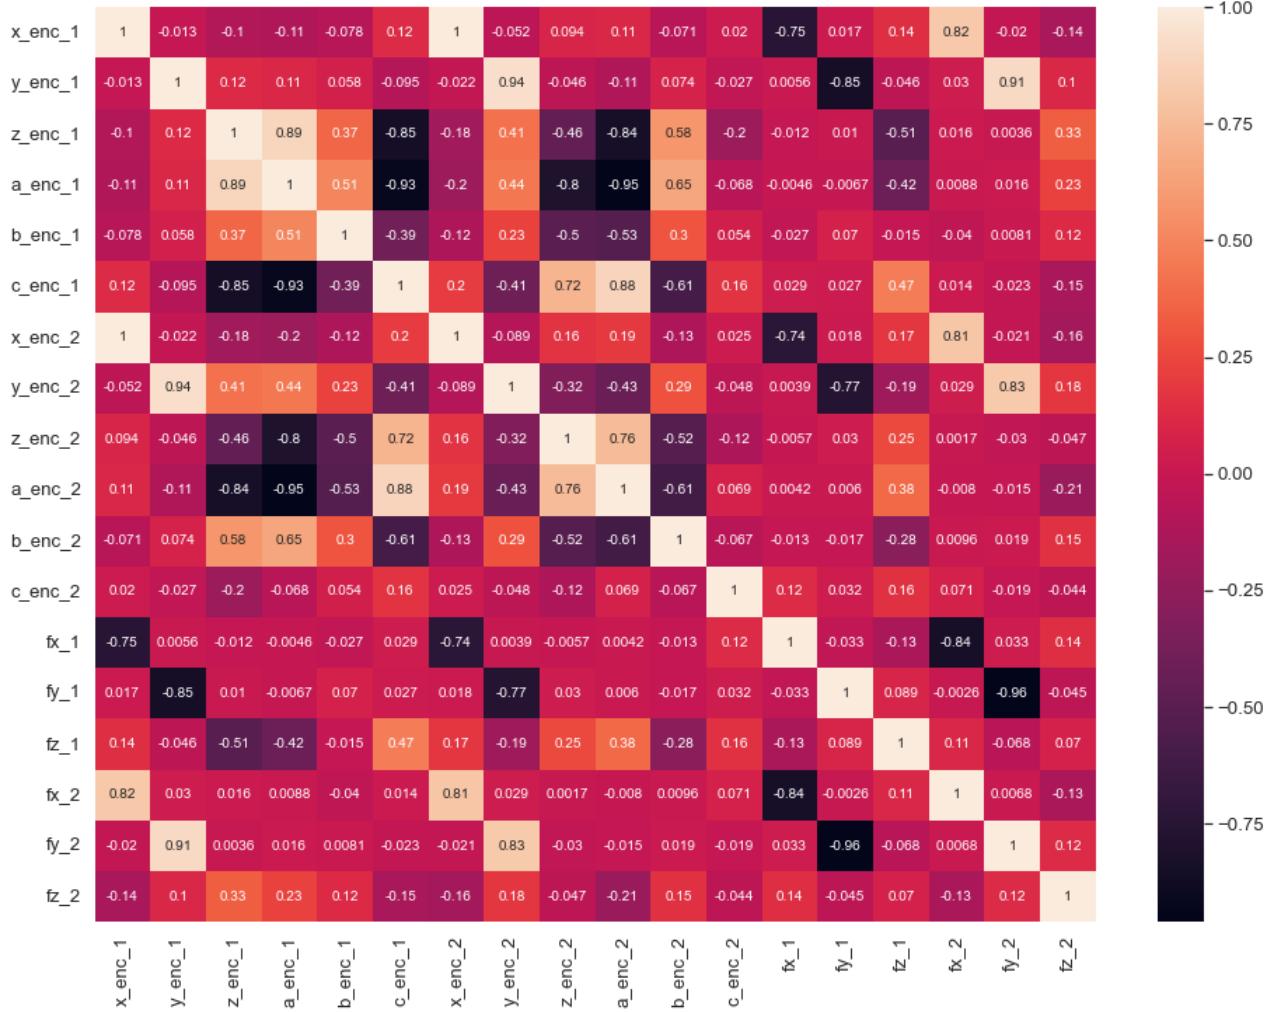
### Observations:

- Forces  $f_x$  and  $f_y$  are strongly correlated with positions  $x$  and  $y$  respectively and relatively weakly related to anything else.
- $f_z$  is mainly correlated to  $z$ , but also quite strongly correlated with  $a$  (roll) and  $c$  (yaw), especially for R1.

### Now let's look at correlations between R1 and R2

```
In [8]:=
dataset_idx = 0 # Test1
labels = features_1+features_2+outputs_1+outputs_2
fig = plt.figure(figsize=(16, 12))
corr = datasets[dataset_idx][labels].corr()
ax = fig.add_subplot(111)
ax.set_title(dataset_filenames[dataset_idx], weight='bold').set_fontsize('18')
sns.set_style("whitegrid")
sns.heatmap(corr, annot=True)
plt.savefig(output_dir/'corr_{0}_both_robots.pdf'.format(dataset_filenames[dataset_idx]))
```

Test1



### Observations:

- Strong correlations between  $x_1$ ,  $x_2$ ,  $f_{x1}$ , and  $f_{x2}$  and between  $y_1$ ,  $y_2$ ,  $f_{y1}$ , and  $f_{y2}$  which seems reasonable given the motion of the two robots.
- Relatively weak correlation between  $f_{z1}$  and  $f_{z2}$ . Forces in  $z$  have quite strong correlation to  $a$ ,  $b$ , and  $c$ .

## 2.2. Time series plots

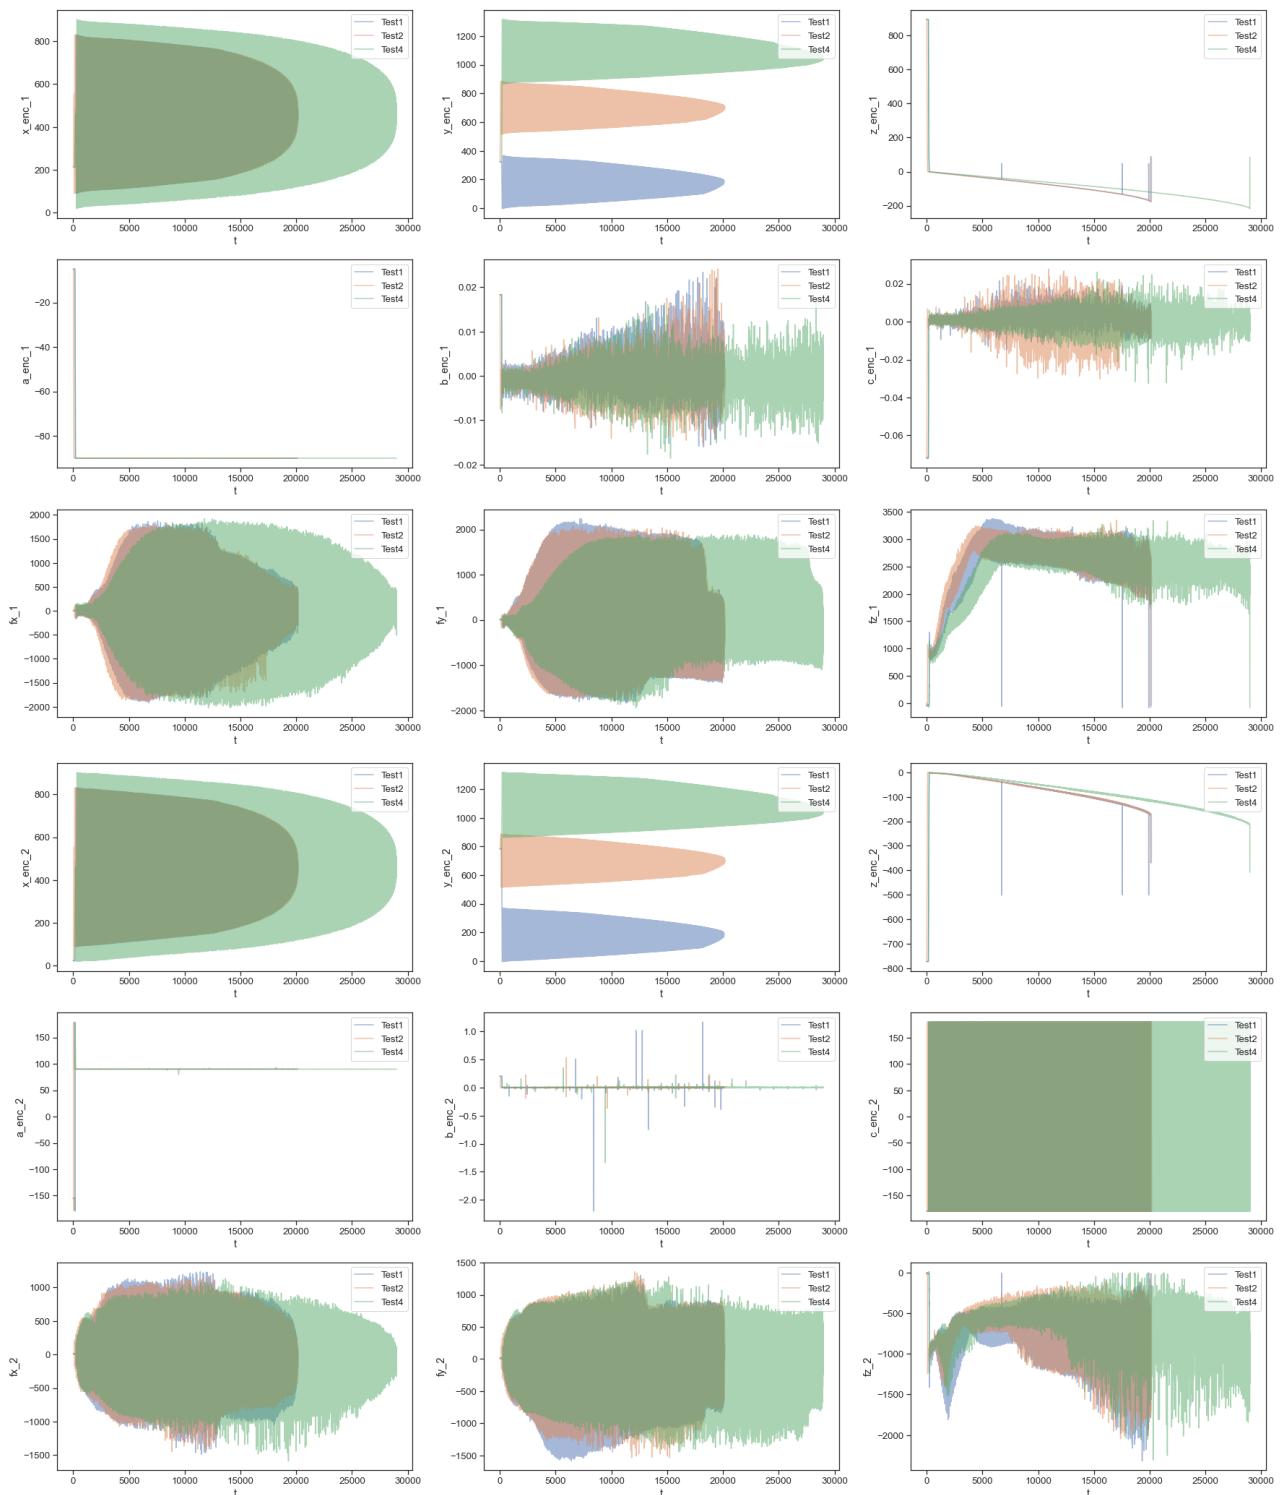
Plot variables in dataset to see what things look like

In [9]:

```

sns.set_style("ticks")
variables_to_plot = [i+'_enc' for i in 'xyzabc'] + ['fx', 'fy', 'fz']
for robot_idx in range(2):
    fig = plt.figure(figsize=(30,18))
    for i,k in enumerate(variables_to_plot):
        k += '_{}'.format(robot_idx+1)
        ax = fig.add_subplot(3,3, i+1)
        for j in range(len(datasets)):
            # shift 't' to start at 0 for each dataset
            ax.plot(datasets[j]['t']-min(datasets[j]['t']), datasets[j][k], label=k)
        ax.set_xlabel('t', y=0.5)
        ax.set_ylabel(k, y=0.5)
    
```

```
ax.legend(loc=1)
fig.savefig(output_dir/'1d_plot_variables_vs_time_robot{}.pdf'.format(robot_
```



## 2.3. Histograms

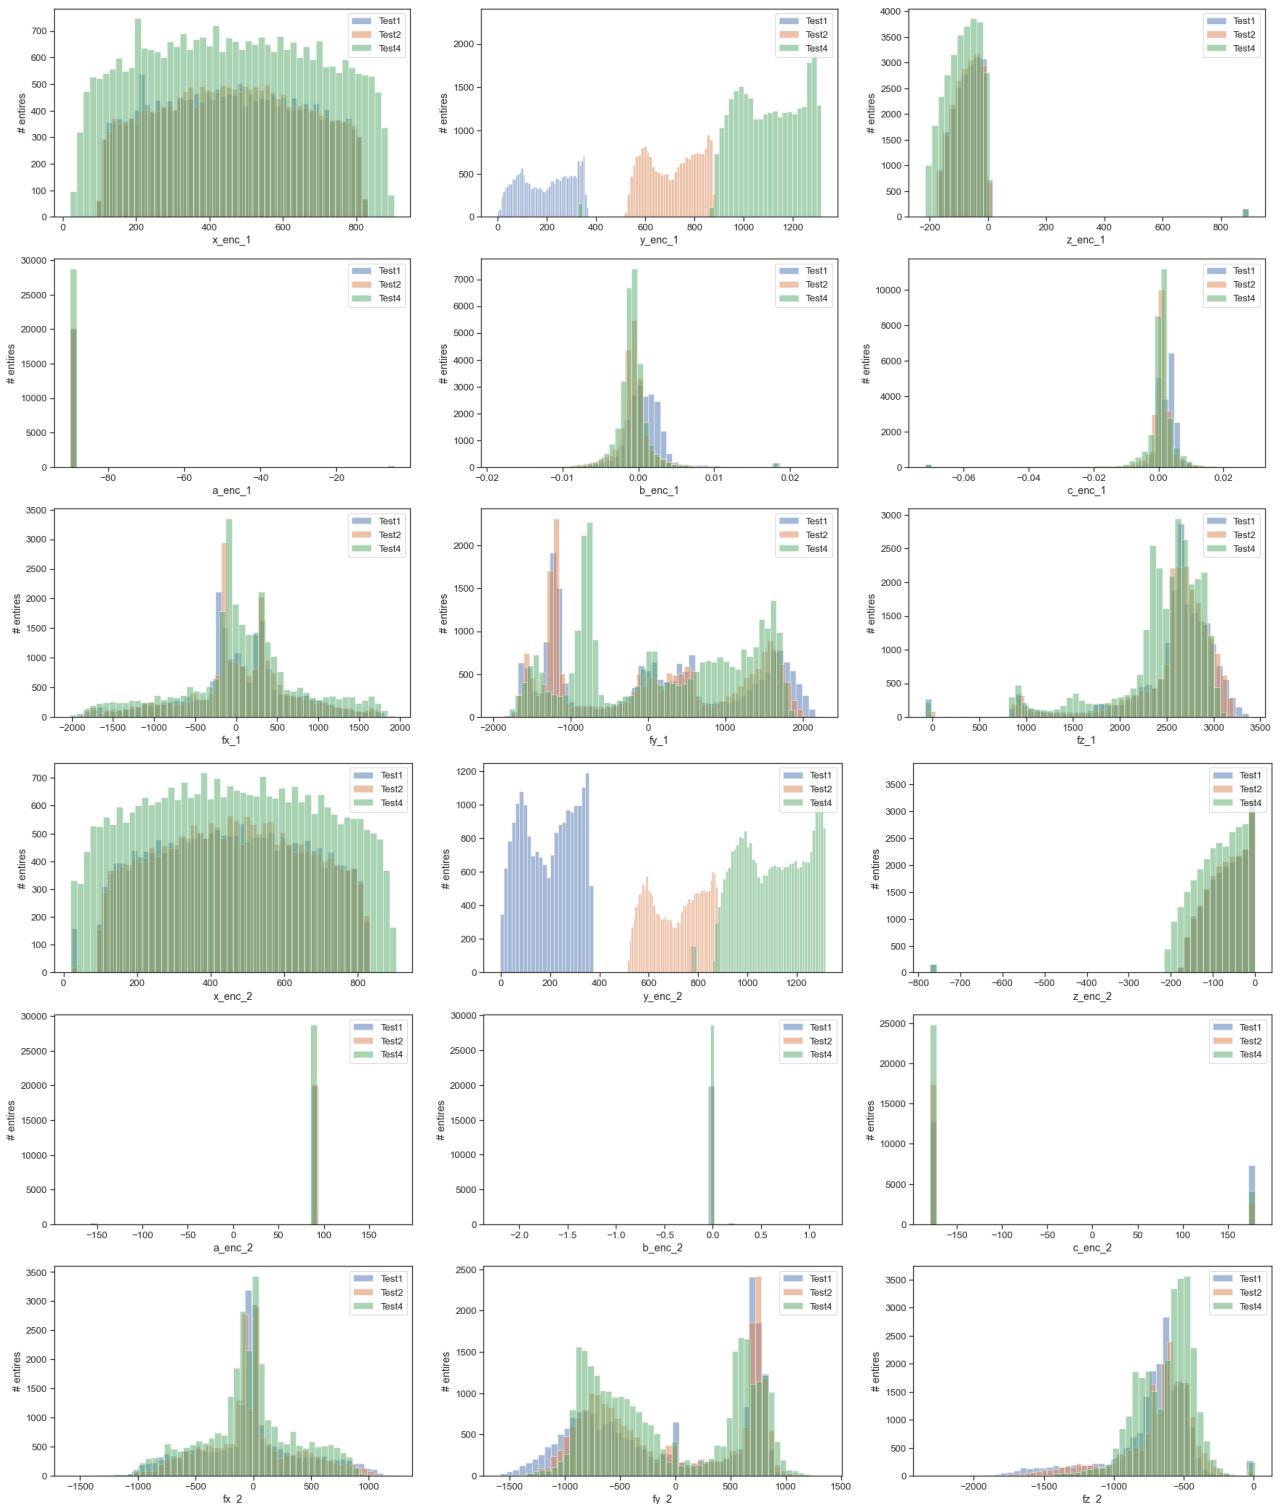
Make histograms of variables in dataset

```
In [10]:=
sns.set_style("ticks")
variables_to_plot = [i+'_enc' for i in 'xyzabc'] + ['fx', 'fy', 'fz']
for robot_idx in range(2):
    fig = plt.figure(figsize=(30,18))
```

```

for i,k in enumerate(variables_to_plot):
    k += '_{}'.format(robot_idx+1)
    ax = fig.add_subplot(3,3, i+1)
    for j in range(len(datasets)):
        ax.hist(datasets[j][k], label=dataset_filenames[j], bins=50, alpha=0
    ax.set_xlabel(k, y=0.5)
    ax.set_ylabel('# entires', y=0.5)
    ax.legend(loc=1)
fig.savefig(output_dir/'1d_hist_variables_vs_time_robot{}.pdf'.format(robot_

```



### 3. Pre-processing

### 3.1. Addition of higher-order derivatives of input features

I noticed that the models could more accurately predict forces from motion in datasets not seen during training when adding higher-order derivatives of the position coordinates as input features.

First (velocity) and second (acceleration) order derivatives had a notable effect on the performance (more about that in section 4). I also tried to include up to 6th order derivatives (3rd=jerk, 4th=snap, 5th=crackle, 6th=pop) [1]. These had a smaller impact but did help in some cases.

[1] [https://en.wikipedia.org/wiki/Fourth,\\_fifth,\\_and\\_sixth\\_derivatives\\_of\\_position](https://en.wikipedia.org/wiki/Fourth,_fifth,_and_sixth_derivatives_of_position)

```
In [11]: # differentiate variables in dataframe
def add_gradients(df, keys_to_diff, nth_order=1):
    for k in keys_to_diff:
        df['d'+str(nth_order)+'_'+re.sub('d\d_', '', k)] = np.gradient(df[k])

In [12]: # add derivatives of position and euler angles:
# 1=velocity, 2=acceleration, 3=jerk, 4=snap, 5=crackle, 6=pop
nth_order = 6
features_to_diff = features
for n in range(1, nth_order+1):
    for df in datasets:
        add_gradients(df, features_to_diff, n)
    features_to_diff = [k for k in datasets[0].keys() if re.search('^\d{:d}'.format(n), k)]
```

#### Create feature lists

```
In [13]: # x-axis only
features_x1 = ['x_enc_1'] + ["d{:d}_x_enc_1".format(i) for i in range(1,7)]
features_x2 = ['x_enc_2'] + ["d{:d}_x_enc_2".format(i) for i in range(1,7)]
outputs_x1 = ['fx_1']
outputs_x2 = ['fx_2']

In [14]: def generate_feature_list(arm_idx, nth_order=6):
    features = []
    for i in 'xyzabc':
        features.append('{}_enc_{}'.format(i, arm_idx))
    for i in 'xyzabc':
        for j in range(1, nth_order+1):
            features.append('d{}_{}_enc_{}'.format(j, i, arm_idx))
    return features
```

```
In [15]: # up to 6th order derivatives (velocity, acceleration, jerk, snap, crackle, pop)
features_1_nth = generate_feature_list(1, nth_order)
features_2_nth = generate_feature_list(2, nth_order)

# up to 2nd order derivatives (velocity, acceleration)
features_1_2nd = generate_feature_list(1, 2)
```

```
features_2_2nd = generate_feature_list(2, 2)

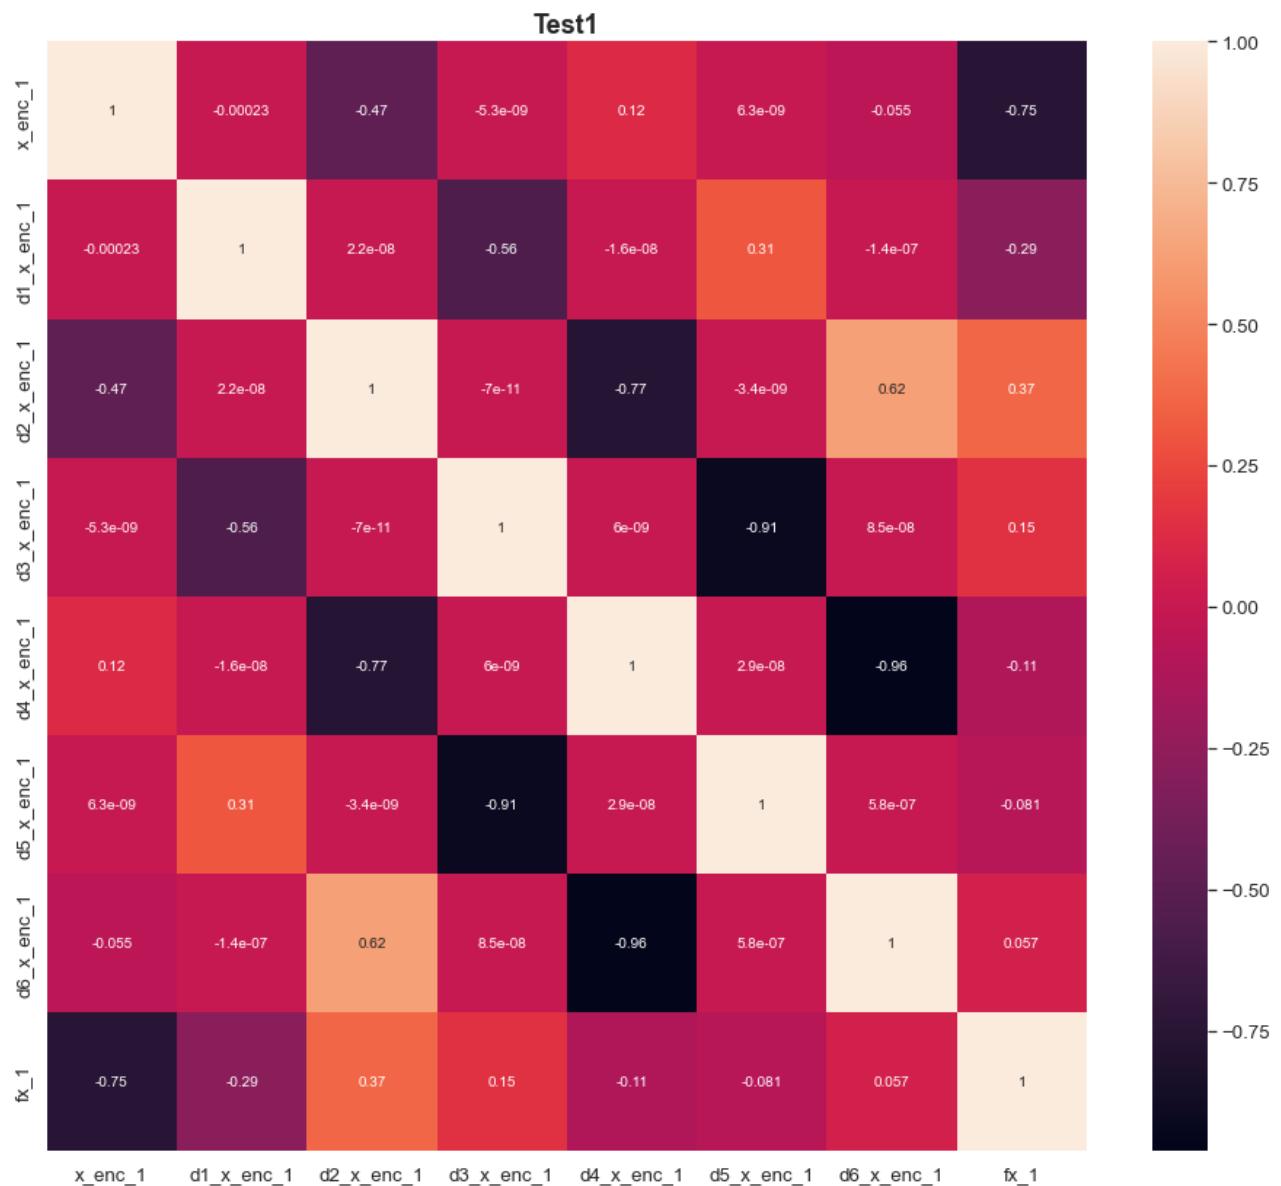
outputs_1 = ['fx_1', 'fy_1', 'fz_1']
outputs_2 = ['fx_2', 'fy_2', 'fz_2']
```

## Correlations to higher order derivatives

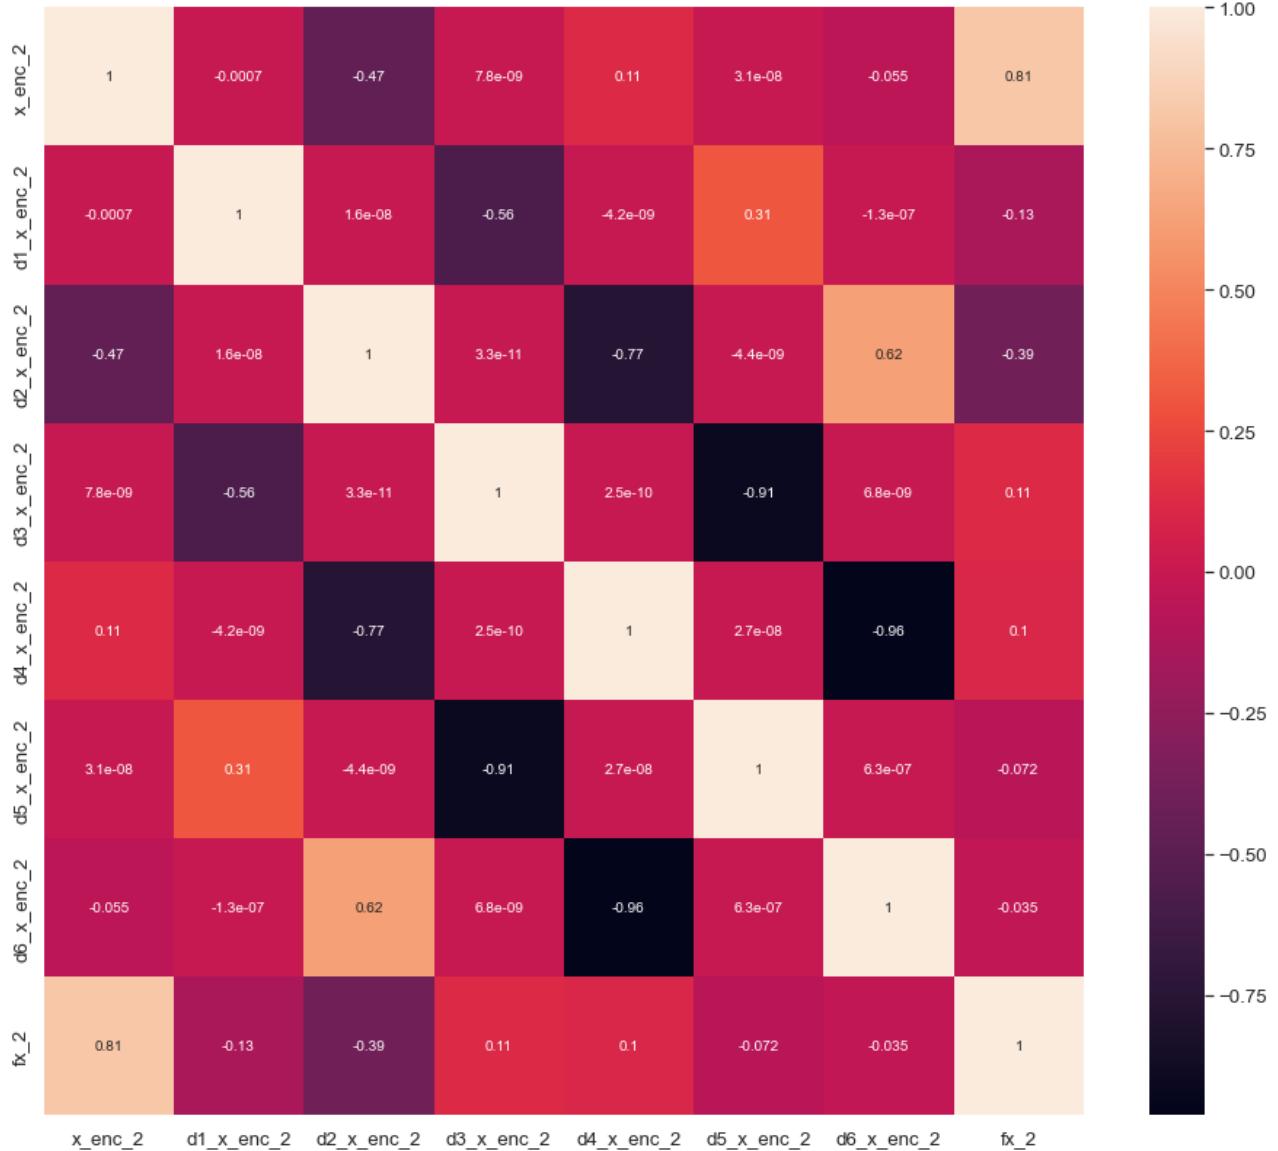
*x-axis only*

In [16]:

```
idx = 0 # Test1
labels = [features_x1+outputs_x1, features_x2+outputs_x2]
sns.set_style("whitegrid")
for robot_idx in range(2):
    fig = plt.figure(figsize=(16, 14))
    corr = datasets[idx][labels[robot_idx]].corr()
    ax = fig.add_subplot(111)
    ax.set_title(dataset_filenames[idx], weight='bold').set_fontsize('18')
    sns.heatmap(corr, annot=True)
    plt.savefig(output_dir/'corr_derivatives_x_robot{}.pdf'.format(robot_idx+1))
```



Test1



## All coordinates

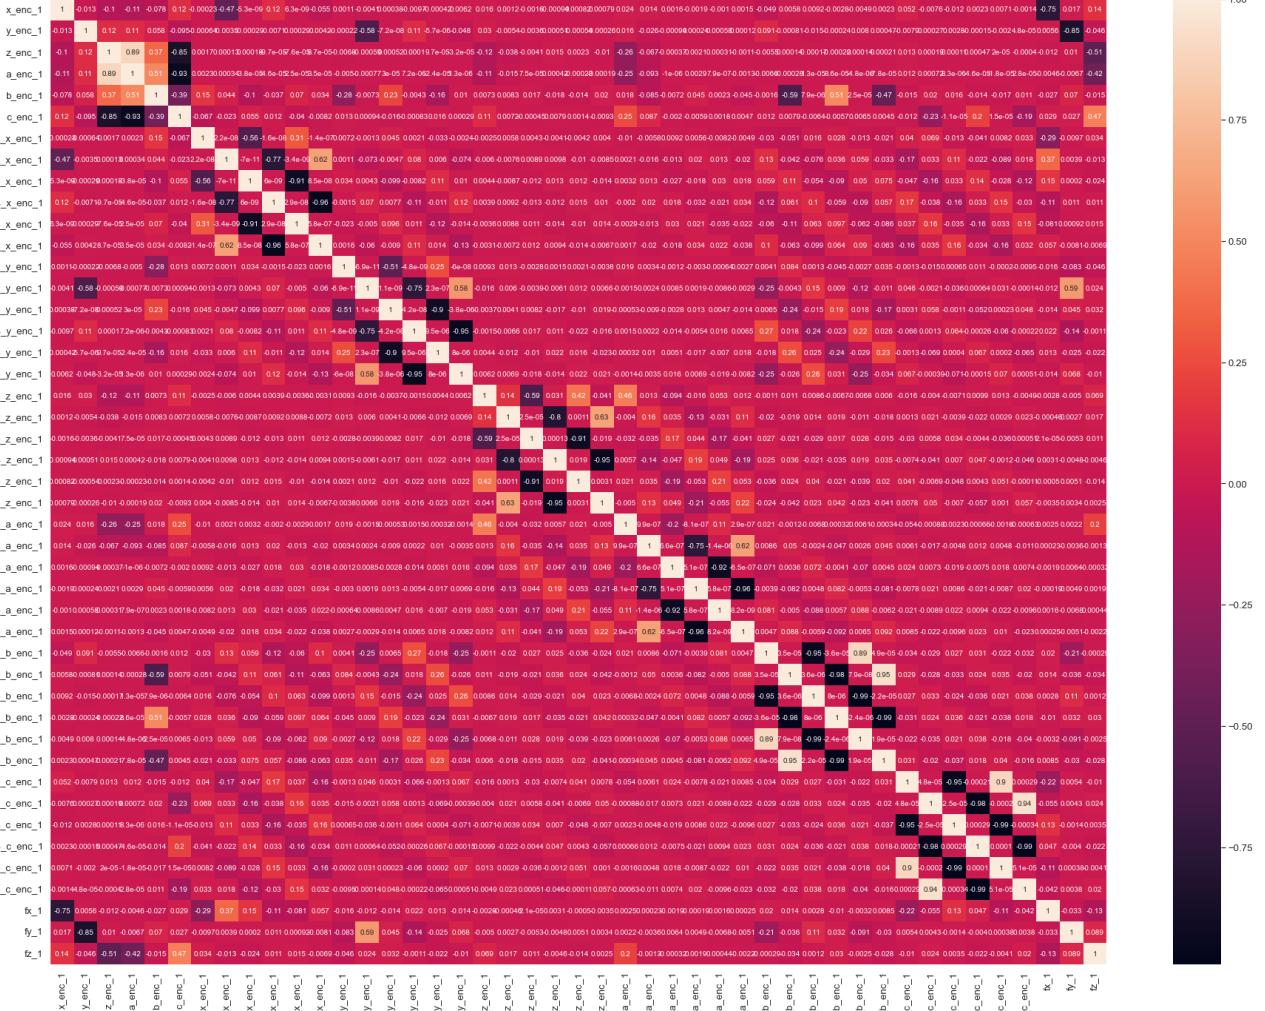
In [17]:

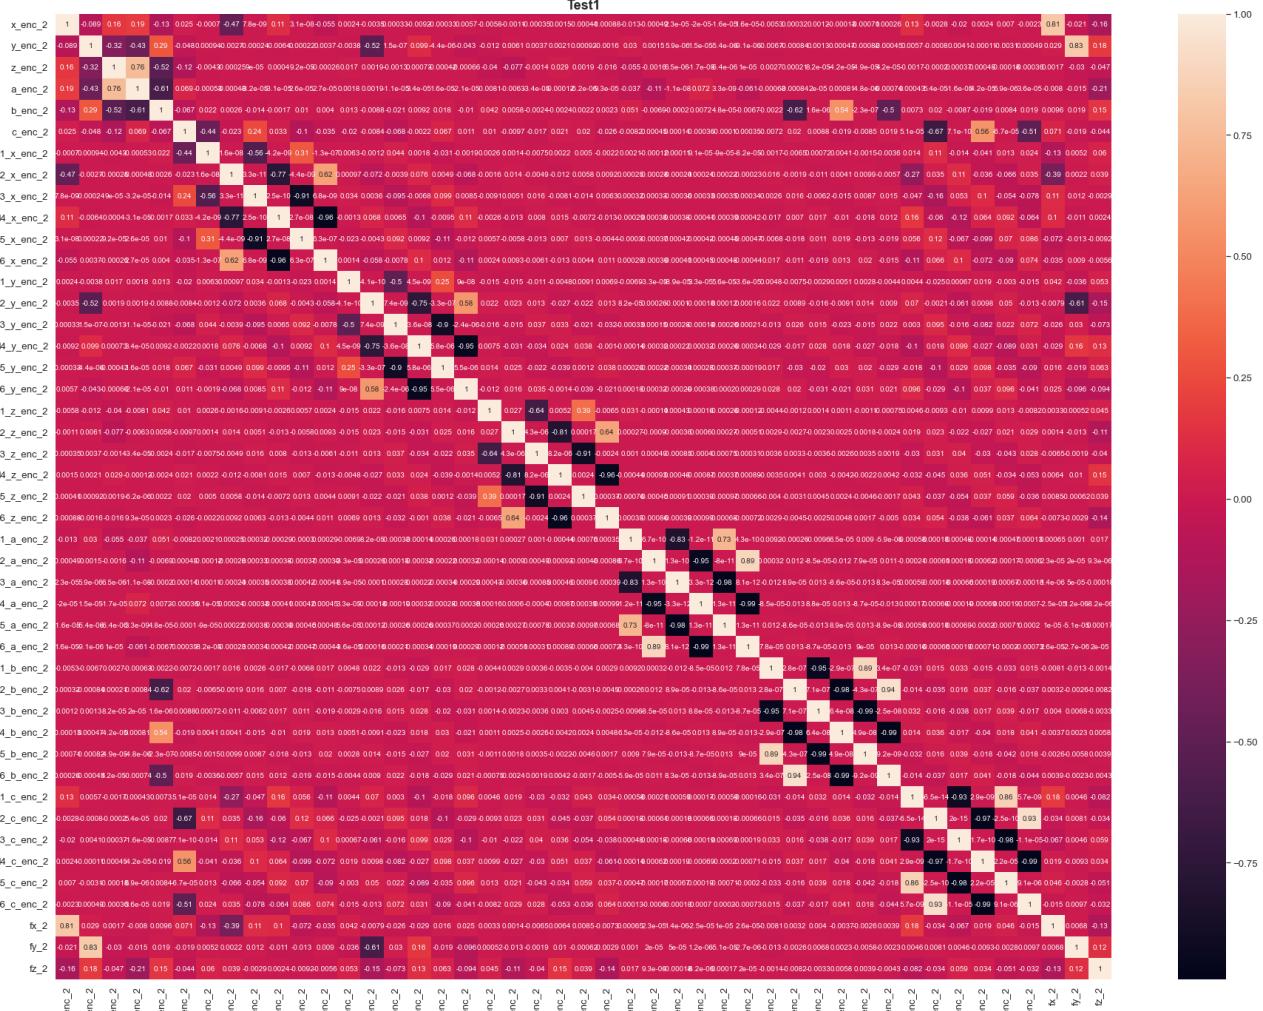
```

idx = 0 # Test1
labels = [features_1_nth+outputs_1, features_2_nth+outputs_2]
sns.set_style("whitegrid")
for robot_idx in range(2):
    fig = plt.figure(figsize=(32, 24))
    corr = datasets[idx][labels[robot_idx]].corr()
    ax = fig.add_subplot(111)
    ax.set_title(dataset_filenames[idx], weight='bold').set_fontsize('18')
    sns.heatmap(corr, annot=True)
    plt.savefig(output_dir/'corr_derivatives_all_robot{}.pdf'.format(robot_idx+1))

```

Test1





### 3.2. Scaling of input features and outputs

## Specify which datasets to use

Hold out 'Test2' for testing of model trained on data from 'Test1' and 'Test4'.

In [18]:

```
df1 = datasets[0] # Test1
df2 = datasets[1] # Test2
df4 = datasets[2] # Test4
df = df1.copy()
#df = df.append(datasets[1])
df = df.append(datasets[2])
```

Split dataframe into X (features) and Y (outputs)

In [19]:

```
features_nth = features_1_nth + features_2_nth
features_2nd = features_1_2nd + features_2_2nd
X = df[features_nth].to_numpy()
Y = df[outputs].to_numpy()
print(X.shape, Y.shape)
```

(49078, 84) (49078, 6)

## Feature scaling

Both MinMaxScaler and StandardScaler were tried, the former gave slightly more robust performance.

In [20]:

```
from sklearn.preprocessing import StandardScaler, MinMaxScaler
#scaler_x = StandardScaler()
#scaler_y = StandardScaler()
scaler_x = MinMaxScaler()
scaler_y = MinMaxScaler()
X_normed = scaler_x.fit_transform(X)
Y_normed = scaler_y.fit_transform(Y)
```

In [21]:

```
# sanity check
print(X[0:5,0])
print(X_normed[0:5,0])
print(scaler_x.inverse_transform(X_normed)[0:4,0])
```

```
[213.6337 213.6337 213.6337 213.6337 213.6337]
[0.21954808 0.21954808 0.21954808 0.21954808 0.21954808]
[213.6337 213.6337 213.6337 213.6337]
```

## 3.3. Preparing input features for RNN training

Including multiple time steps of the input features and training an RNN to predict forces improved the performance over using a DNN. After some experimentation, I concluded that about 20 timesteps worked pretty well.

In [22]:

```
def split_sequences(X, Y, n_steps):
    X_seq, Y_seq = list(), list()
    for i in range(len(X)):
        end_i = i + n_steps
        if end_i > len(X):
            break
        Xi, yi = X[i:end_i, :], Y[end_i-1, :]
        X_seq.append(Xi)
        Y_seq.append(yi)
    return (np.array(X_seq), np.array(Y_seq))
```

In [23]:

```
# prepare sequences for RNN with 20 time steps
n_steps = 20
X_seq, Y_seq = split_sequences(X_normed, Y_normed, n_steps)
print(X_seq.shape, Y_seq.shape)
```

```
(49059, 20, 84) (49059, 6)
```

## 3.4. Splitting of data into train/validation/test sets

In [24]:

```
from sklearn.model_selection import train_test_split
# if False: keep last events for testing
shuffle = True
# 70-10-20 train-validation-test split
```

```

train_frac = 0.7
val_frac = 0.1

# for dnn
X_train, X_val_test, Y_train, Y_val_test = train_test_split(X_normed, Y_normed,
X_val, X_test, Y_val, Y_test = train_test_split(X_val_test, Y_val_test, test_size=1)

print(X_train.shape, Y_train.shape)
print(X_val.shape, Y_val.shape)
print(X_test.shape, Y_test.shape)

# for rnn
X_seq_train, X_seq_val_test, Y_seq_train, Y_seq_val_test = train_test_split(X_seq,
X_seq_val, X_seq_test, Y_seq_val, Y_seq_test = train_test_split(X_seq_val_test,
Y_seq_val_test, test_size=1)

print(X_seq_train.shape, Y_seq_train.shape)
print(X_seq_val.shape, Y_seq_val.shape)
print(X_seq_test.shape, Y_seq_test.shape)

```

```

(34354, 84) (34354, 6)
(4907, 84) (4907, 6)
(9817, 84) (9817, 6)
(34341, 20, 84) (34341, 6)
(4905, 20, 84) (4905, 6)
(9813, 20, 84) (9813, 6)

```

Create train/val/test sets with 12 (positions and angles) and 36 input features (positions, angles, and up to 2nd order derivatives)

In [25]:

```

# sanity check
print(features)
print(features_2nd)
print(outputs)

```

```

['x_enc_1', 'y_enc_1', 'z_enc_1', 'a_enc_1', 'b_enc_1', 'c_enc_1', 'x_enc_2', 'y
_enc_2', 'z_enc_2', 'a_enc_2', 'b_enc_2', 'c_enc_2']
['x_enc_1', 'y_enc_1', 'z_enc_1', 'a_enc_1', 'b_enc_1', 'c_enc_1', 'd1_x_enc_1',
'd2_x_enc_1', 'd1_y_enc_1', 'd2_y_enc_1', 'd1_z_enc_1', 'd2_z_enc_1', 'd1_a_enc_
1', 'd2_a_enc_1', 'd1_b_enc_1', 'd2_b_enc_1', 'd1_c_enc_1', 'd2_c_enc_1', 'x_enc
_2', 'y_enc_2', 'z_enc_2', 'a_enc_2', 'b_enc_2', 'c_enc_2', 'd1_x_enc_2', 'd2_x_
enc_2', 'd1_y_enc_2', 'd2_y_enc_2', 'd1_z_enc_2', 'd2_z_enc_2', 'd1_a_enc_2', 'd
2_a_enc_2', 'd1_b_enc_2', 'd2_b_enc_2', 'd1_c_enc_2', 'd2_c_enc_2']
['fx_1', 'fy_1', 'fz_1', 'fx_2', 'fy_2', 'fz_2']

```

In [26]:

```

# indices of 12 position input features
feature_idx = [df[features_nth].columns.get_loc(col) for col in features]
# indices of position, velocity, and acceleration input features
feature_idx_2nd = [df[features_nth].columns.get_loc(col) for col in features_2nd]
# indices of all features
feature_idx_nth = [df[features_nth].columns.get_loc(col) for col in features_nth]

```

In [27]:

```

# sanity check
print(df[features][4000:4001].to_numpy())
print(scaler_x.inverse_transform(X_normed)[:, feature_idx][4000:4001, :])

```

```

[[ 7.01138030e+02   9.05163185e+01  -2.66579203e+01  -8.99984769e+01
 -4.16576996e-04   2.06331677e-03    7.04664055e+02   8.46991898e+01
 -2.02531592e+01   8.99942266e+01  -2.56791472e-04  -1.79998530e+02]]

```

```
[[ 7.01138030e+02  9.05163185e+01 -2.66579203e+01 -8.99984769e+01
-4.16576996e-04  2.06331677e-03  7.04664055e+02  8.46991898e+01
-2.02531593e+01  8.99942266e+01 -2.56791472e-04 -1.79998530e+02]]
```

In [28]:

```
# select 12 input features: x,y,z,a,b,c for robots 1 and 2
X_train_12 = X_train[:,feature_idx]
X_val_12 = X_val[:,feature_idx]
X_test_12 = X_test[:,feature_idx]

X_seq_train_12 = X_seq_train[:, :, feature_idx]
X_seq_val_12 = X_seq_val[:, :, feature_idx]
X_seq_test_12 = X_seq_test[:, :, feature_idx]

# select 36 input features: x,y,z,a,b,c + 1st and 2nd order derivatives for robots 1 and 2
X_train_36 = X_train[:, feature_idx_2nd]
X_val_36 = X_val[:, feature_idx_2nd]
X_test_36 = X_test[:, feature_idx_2nd]

X_seq_train_36 = X_seq_train[:, :, feature_idx_2nd]
X_seq_val_36 = X_seq_val[:, :, feature_idx_2nd]
X_seq_test_36 = X_seq_test[:, :, feature_idx_2nd]
```

In [29]:

```
# sanity check
print(X_train_12.shape)
print(X_seq_train_12.shape)

print(X_train_36.shape)
print(X_seq_train_36.shape)
```

```
(34354, 12)
(34341, 20, 12)
(34354, 36)
(34341, 20, 36)
```

## 4. Modeling

Finally, we're getting to the exciting part of training some neural nets.

In [165...]

```
# function to plot loss, to be used several times below
def plot_loss(history, name, title=''):
    fig = plt.figure(figsize=(24,10))
    fig.suptitle(title)

    # full range
    ax = fig.add_subplot(121)
    ax.plot(history.history['loss'], label='loss')
    ax.plot(history.history['val_loss'], label='val_loss')
    ax.set_xlabel('Epoch')
    ax.set_ylabel('Error')
    ax.legend()
    ax.grid(True)

    # last 30 percent of epochs
    zoom_frac = 0.7
    nepochs = len(history.history['loss'])
    ax = fig.add_subplot(122)
    ax.set_xlim(nepochs * (1 - zoom_frac), nepochs)
```

```

ax.plot(history.history[ 'loss' ], label='loss')
ax.plot(history.history[ 'val_loss' ], label='val_loss')
xmin = int(zoom_frac*nepochs)
xmax = nepochs
ax.set_xlim([xmin, xmax])
ymin = 0.99*np.min(history.history[ 'loss' ][int(zoom_frac*nepochs):]
                     + history.history[ 'val_loss' ][int(zoom_frac*nepochs):])
ymax = 1.01*np.max(history.history[ 'loss' ][int(zoom_frac*nepochs):]
                     + history.history[ 'val_loss' ][int(zoom_frac*nepochs):])
ax.set_ylim([ymin, ymax])
ax.set_xlabel('Epoch')
ax.set_ylabel('Error')
ax.legend()
ax.grid(True)

plt.savefig(output_dir/name)

```

In [142...]

```

# for comparing test results of different models
test_results = dict()

```

## 4.1. Linear Regression

Let's start with a simple linear regression.

Predict  $f_{x_1}$  from  $x_1$

In [32]:

```

x1_train = X_train[:,0].reshape(len(X_train),1)
fx1_train = Y_train[:,0].reshape(len(Y_train),1)
x1_val = X_val[:,0].reshape(len(X_val),1)
fx1_val = Y_val[:,0].reshape(len(Y_val),1)
x1_test = X_test[:,0].reshape(len(X_test),1)
fx1_test = Y_test[:,0].reshape(len(Y_test),1)

```

In [33]:

```

linear_model_x1 = keras.experimental.LinearModel()
linear_model_x1.compile(optimizer='adam', loss='mean_squared_error')
history_linear_x1 = linear_model_x1.fit(
    x1_train, fx1_train,
    validation_data=(x1_val, fx1_val),
    batch_size = 32,
    epochs=40)
with open(output_dir/'history_linear_x1.pickle', 'wb') as f:
    pickle.dump(history_linear_x1.history, f)

```

```

Epoch 1/40
1074/1074 [=====] - 2s 932us/step - loss: 0.0849 - val_
loss: 0.0388
Epoch 2/40
1074/1074 [=====] - 1s 908us/step - loss: 0.0279 - val_
loss: 0.0192
Epoch 3/40
1074/1074 [=====] - 1s 897us/step - loss: 0.0166 - val_
loss: 0.0152
Epoch 4/40
1074/1074 [=====] - 1s 922us/step - loss: 0.0150 - val_

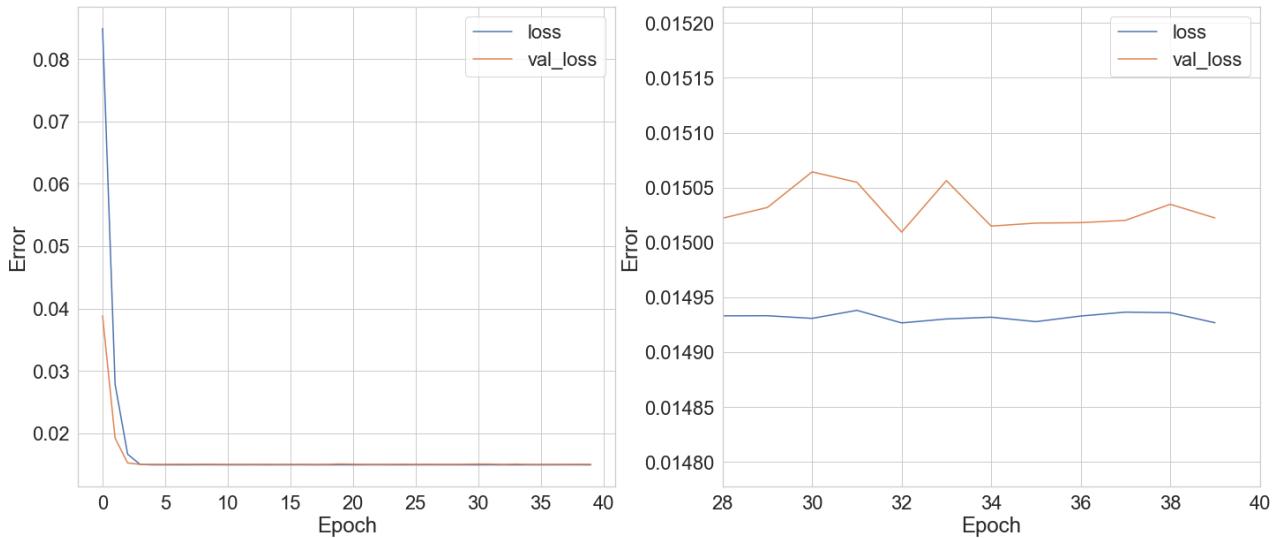
```

```
loss: 0.0150
Epoch 5/40
1074/1074 [=====] - 1s 844us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 6/40
1074/1074 [=====] - 1s 754us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 7/40
1074/1074 [=====] - 1s 771us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 8/40
1074/1074 [=====] - 1s 797us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 9/40
1074/1074 [=====] - 1s 860us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 10/40
1074/1074 [=====] - 1s 870us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 11/40
1074/1074 [=====] - 1s 826us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 12/40
1074/1074 [=====] - 1s 792us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 13/40
1074/1074 [=====] - 1s 945us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 14/40
1074/1074 [=====] - 1s 926us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 15/40
1074/1074 [=====] - 1s 900us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 16/40
1074/1074 [=====] - 1s 888us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 17/40
1074/1074 [=====] - 1s 818us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 18/40
1074/1074 [=====] - 1s 844us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 19/40
1074/1074 [=====] - 1s 839us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 20/40
1074/1074 [=====] - 1s 872us/step - loss: 0.0149 - val_
loss: 0.0151
Epoch 21/40
1074/1074 [=====] - 1s 890us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 22/40
1074/1074 [=====] - 1s 858us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 23/40
1074/1074 [=====] - 1s 853us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 24/40
1074/1074 [=====] - 1s 867us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 25/40
1074/1074 [=====] - 1s 849us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 26/40
```

```
1074/1074 [=====] - 1s 850us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 27/40
1074/1074 [=====] - 1s 867us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 28/40
1074/1074 [=====] - 1s 823us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 29/40
1074/1074 [=====] - 1s 861us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 30/40
1074/1074 [=====] - 1s 856us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 31/40
1074/1074 [=====] - 1s 867us/step - loss: 0.0149 - val_
loss: 0.0151
Epoch 32/40
1074/1074 [=====] - 1s 888us/step - loss: 0.0149 - val_
loss: 0.0151
Epoch 33/40
1074/1074 [=====] - 1s 865us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 34/40
1074/1074 [=====] - 1s 854us/step - loss: 0.0149 - val_
loss: 0.0151
Epoch 35/40
1074/1074 [=====] - 1s 858us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 36/40
1074/1074 [=====] - 1s 866us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 37/40
1074/1074 [=====] - 1s 850us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 38/40
1074/1074 [=====] - 1s 864us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 39/40
1074/1074 [=====] - 1s 831us/step - loss: 0.0149 - val_
loss: 0.0150
Epoch 40/40
1074/1074 [=====] - 1s 642us/step - loss: 0.0149 - val_
loss: 0.0150
```

## Plot loss vs. epoch

```
In [166]: plot_loss(history_linear_x1, 'loss_linear_x1.pdf', '1D linear model ($f_{x_1}$ v
```

olsson\_solution  
 1D linear model ( $f_{x_1}$  vs.  $x_1$ )


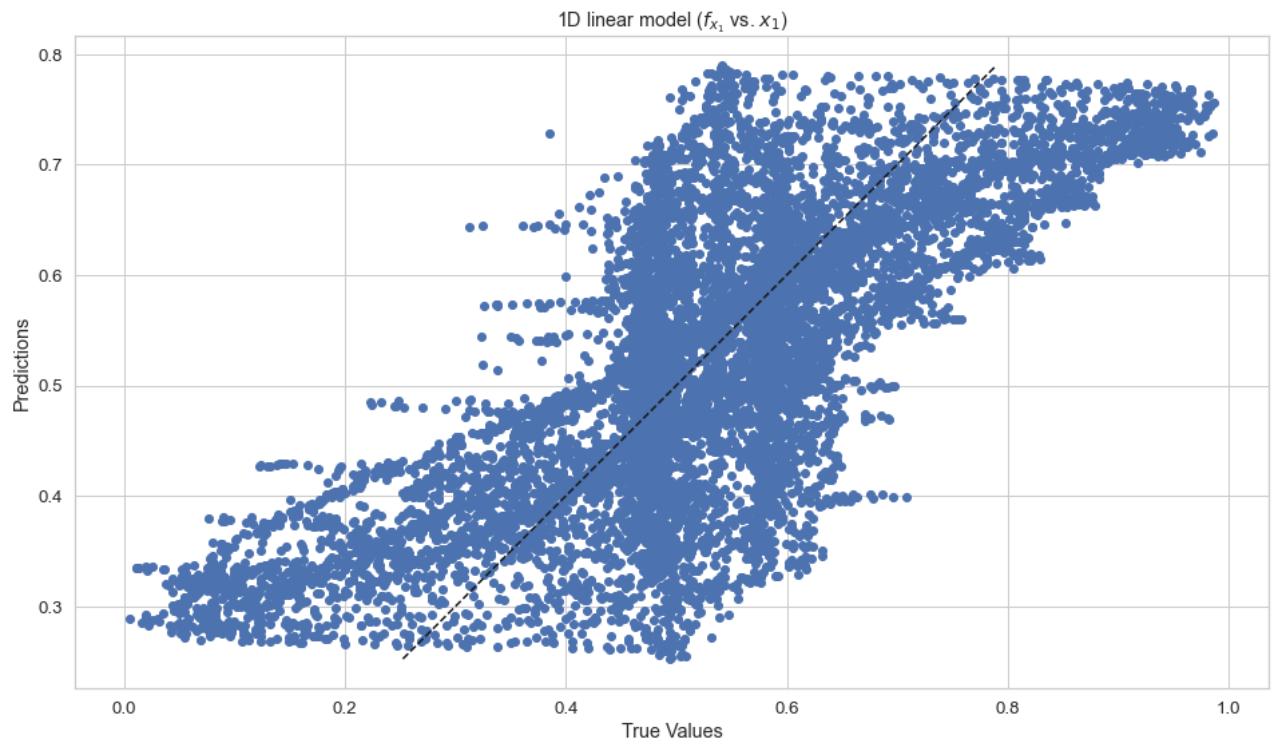
In [35]:

```
# save model loss on test set for evaluation section below
test_results['linear_x1'] = linear_model_x1.evaluate(x1_test, fx1_test, verbose=
```

## Plot predictions vs. true values

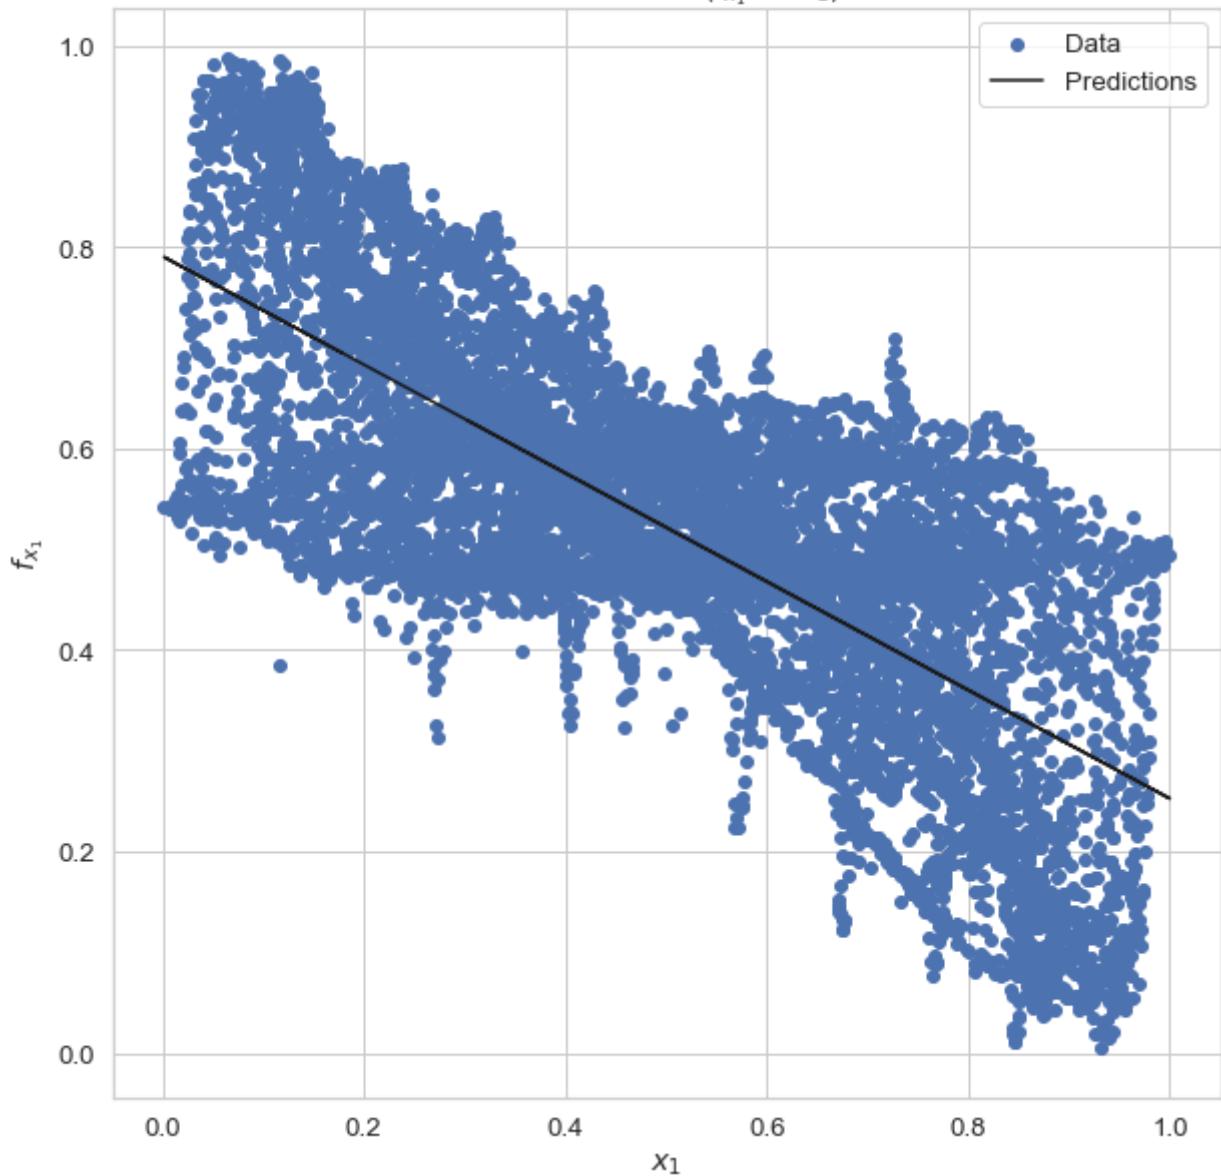
In [36]:

```
fx1_test_pred = linear_model_x1.predict(x1_test)
fig = plt.figure(figsize=(16,10))
ax = fig.add_subplot(111, aspect='equal')
ax.scatter(fx1_test, fx1_test_pred)
ax.plot([np.min(fx1_test_pred), np.max(fx1_test_pred)], [np.min(fx1_test_pred),
#ax.plot([0.2, 0.8], [0.2, 0.8], linestyle='dashed', label='Predictions', color=
ax.set_xlabel('True Values')
ax.set_ylabel('Predictions')
ax.set_title('1D linear model ($f_{x_1}$ vs. $x_1$)')
plt.savefig(output_dir/'pred_vs_true_linear_x1.pdf')
```

Plot  $f_{x_1}$  vs.  $x_1$  (scaled)

In [37]:

```
fig = plt.figure(figsize=(16,10))
ax = fig.add_subplot(111, aspect='equal')
ax.scatter(x1_test, fx1_test, label='Data')
ax.plot(x1_test, fx1_test_pred, label='Predictions', color='k')
ax.set_xlabel('$x_{\{1\}}$')
ax.set_ylabel('$f_{\{x_{\{1\}}\}}$')
ax.set_title('1D linear model ($f_{\{x_1\}}$ vs. $x_{\{1\}}$)')
ax.legend()
plt.savefig(output_dir/'linear_fx1_vs_x1.pdf')
```

1D linear model ( $f_{x_1}$  vs.  $x_1$ )

## Predict all 6 forces from 12 input features

```
In [38]: linear_model_12 = keras.experimental.LinearModel(units=6)
linear_model_12.compile(optimizer='adam', loss='mean_squared_error')
history_linear_12 = linear_model_12.fit(
    X_train_12, Y_train,
    validation_data=(X_val_12, Y_val),
    batch_size = 32,
    epochs=200)
with open(output_dir/'history_linear_12.pickle', 'wb') as f:
    pickle.dump(history_linear_12.history, f)
```

```
Epoch 1/200
1074/1074 [=====] - 1s 874us/step - loss: 0.0454 - val_loss: 0.0302
Epoch 2/200
1074/1074 [=====] - 1s 792us/step - loss: 0.0295 - val_loss: 0.0295
Epoch 3/200
1074/1074 [=====] - 1s 807us/step - loss: 0.0292 - val_
```

```
loss: 0.0292
Epoch 4/200
1074/1074 [=====] - 1s 863us/step - loss: 0.0290 - val_
loss: 0.0290
Epoch 5/200
1074/1074 [=====] - 1s 803us/step - loss: 0.0289 - val_
loss: 0.0291
Epoch 6/200
1074/1074 [=====] - 1s 822us/step - loss: 0.0289 - val_
loss: 0.0290
Epoch 7/200
1074/1074 [=====] - 1s 817us/step - loss: 0.0288 - val_
loss: 0.0289
Epoch 8/200
1074/1074 [=====] - 1s 842us/step - loss: 0.0288 - val_
loss: 0.0290
Epoch 9/200
1074/1074 [=====] - 1s 804us/step - loss: 0.0288 - val_
loss: 0.0289
Epoch 10/200
1074/1074 [=====] - 1s 814us/step - loss: 0.0288 - val_
loss: 0.0289
Epoch 11/200
1074/1074 [=====] - 1s 822us/step - loss: 0.0288 - val_
loss: 0.0288
Epoch 12/200
1074/1074 [=====] - 1s 847us/step - loss: 0.0287 - val_
loss: 0.0289
Epoch 13/200
1074/1074 [=====] - 1s 840us/step - loss: 0.0287 - val_
loss: 0.0289
Epoch 14/200
1074/1074 [=====] - 1s 812us/step - loss: 0.0287 - val_
loss: 0.0288
Epoch 15/200
1074/1074 [=====] - 1s 798us/step - loss: 0.0287 - val_
loss: 0.0287
Epoch 16/200
1074/1074 [=====] - 1s 826us/step - loss: 0.0287 - val_
loss: 0.0287
Epoch 17/200
1074/1074 [=====] - 1s 808us/step - loss: 0.0286 - val_
loss: 0.0287
Epoch 18/200
1074/1074 [=====] - 1s 811us/step - loss: 0.0286 - val_
loss: 0.0287
Epoch 19/200
1074/1074 [=====] - 1s 807us/step - loss: 0.0286 - val_
loss: 0.0287
Epoch 20/200
1074/1074 [=====] - 1s 847us/step - loss: 0.0286 - val_
loss: 0.0288
Epoch 21/200
1074/1074 [=====] - 1s 810us/step - loss: 0.0286 - val_
loss: 0.0287
Epoch 22/200
1074/1074 [=====] - 1s 818us/step - loss: 0.0286 - val_
loss: 0.0287
Epoch 23/200
1074/1074 [=====] - 1s 770us/step - loss: 0.0286 - val_
loss: 0.0287
Epoch 24/200
1074/1074 [=====] - 1s 805us/step - loss: 0.0285 - val_
loss: 0.0288
Epoch 25/200
```

```
Epoch 177/200
1074/1074 [=====] - 1s 814us/step - loss: 0.0273 - val_
loss: 0.0274
Epoch 178/200
1074/1074 [=====] - 1s 803us/step - loss: 0.0273 - val_
loss: 0.0276
Epoch 179/200
1074/1074 [=====] - 1s 812us/step - loss: 0.0273 - val_
loss: 0.0274
Epoch 180/200
1074/1074 [=====] - 1s 799us/step - loss: 0.0272 - val_
loss: 0.0273
Epoch 181/200
1074/1074 [=====] - 1s 839us/step - loss: 0.0272 - val_
loss: 0.0273
Epoch 182/200
1074/1074 [=====] - 1s 844us/step - loss: 0.0273 - val_
loss: 0.0273
Epoch 183/200
1074/1074 [=====] - 1s 804us/step - loss: 0.0272 - val_
loss: 0.0273
Epoch 184/200
1074/1074 [=====] - 1s 801us/step - loss: 0.0272 - val_
loss: 0.0273
Epoch 185/200
1074/1074 [=====] - 1s 814us/step - loss: 0.0272 - val_
loss: 0.0273
Epoch 186/200
1074/1074 [=====] - 1s 807us/step - loss: 0.0272 - val_
loss: 0.0273
Epoch 187/200
1074/1074 [=====] - 1s 811us/step - loss: 0.0272 - val_
loss: 0.0273
Epoch 188/200
1074/1074 [=====] - 1s 850us/step - loss: 0.0272 - val_
loss: 0.0273
Epoch 189/200
1074/1074 [=====] - 1s 836us/step - loss: 0.0272 - val_
loss: 0.0273
Epoch 190/200
1074/1074 [=====] - 1s 817us/step - loss: 0.0272 - val_
loss: 0.0273
Epoch 191/200
1074/1074 [=====] - 1s 801us/step - loss: 0.0272 - val_
loss: 0.0272
Epoch 192/200
1074/1074 [=====] - 1s 818us/step - loss: 0.0272 - val_
loss: 0.0276
Epoch 193/200
1074/1074 [=====] - 1s 840us/step - loss: 0.0272 - val_
loss: 0.0272
Epoch 194/200
1074/1074 [=====] - 1s 814us/step - loss: 0.0272 - val_
loss: 0.0273
Epoch 195/200
1074/1074 [=====] - 1s 808us/step - loss: 0.0272 - val_
loss: 0.0272
Epoch 196/200
1074/1074 [=====] - 1s 812us/step - loss: 0.0272 - val_
loss: 0.0272
Epoch 197/200
1074/1074 [=====] - 1s 810us/step - loss: 0.0271 - val_
loss: 0.0272
Epoch 198/200
1074/1074 [=====] - 1s 832us/step - loss: 0.0271 - val_
```

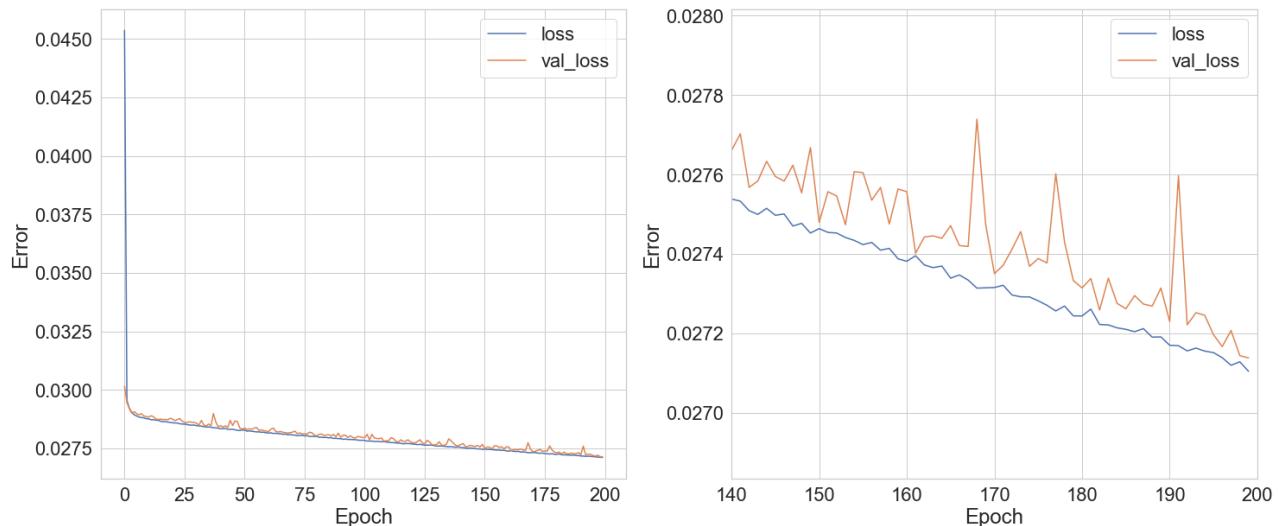
```
loss: 0.0272
Epoch 199/200
1074/1074 [=====] - 1s 807us/step - loss: 0.0271 - val_
loss: 0.0271
Epoch 200/200
1074/1074 [=====] - 1s 860us/step - loss: 0.0271 - val_
loss: 0.0271
```

## Plot loss vs. epoch

In [167]:

```
plot_loss(history_linear_12, 'loss_linear_12.pdf', 'Full linear model (predict 6
```

Full linear model (predict 6 forces from 12 input features)



In [40]:

```
# save model loss on test set for evaluation section below
test_results['linear_12'] = linear_model_12.evaluate(X_test_12, Y_test, verbose=
```

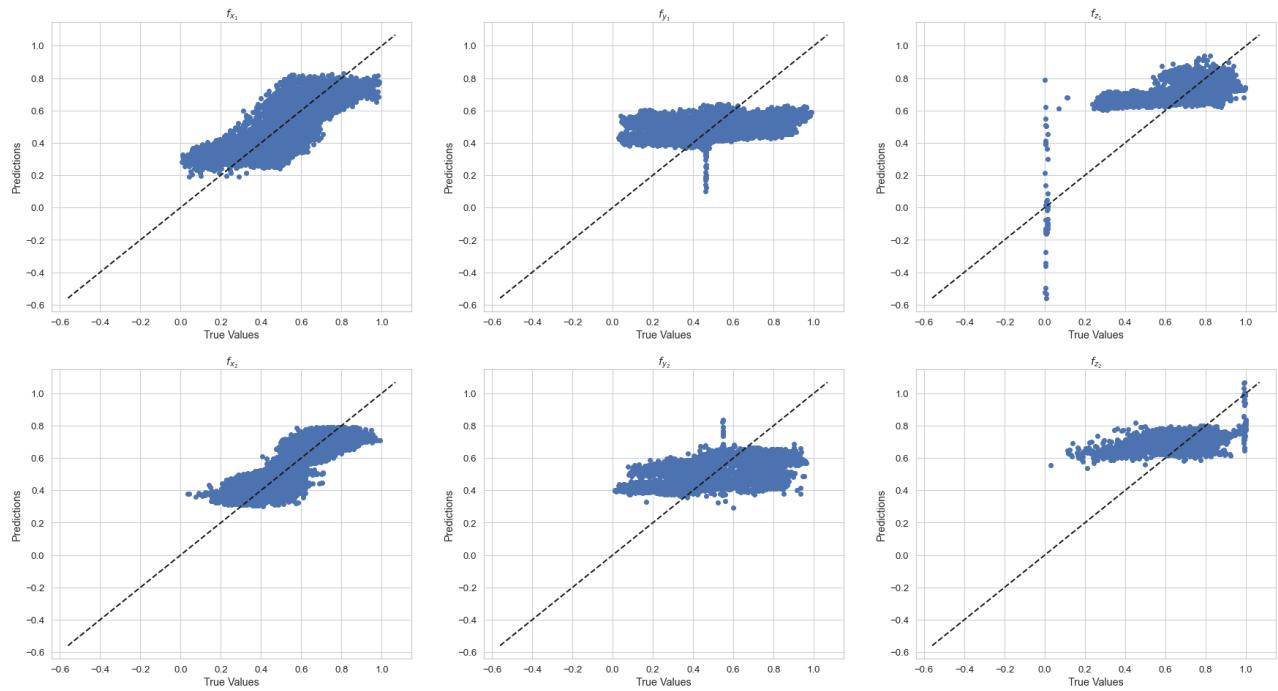
## Plot predictions vs. true values

In [41]:

```
def plot_pred_vs_true(pred, true, name, titles=None):
    fig = plt.figure(figsize=(30,16))
    for i in range(len(pred.T)):
        ax = fig.add_subplot(2,3,i+1)
        ax.scatter(true.T[i], pred.T[i])
        ax.plot([np.min(pred), np.max(pred)], [np.min(pred), np.max(pred)], lines
        ax.set_xlabel('True Values')
        ax.set_ylabel('Predictions')
        if len(titles)>=len(pred.T):
            ax.set_title(titles[i])
    plt.savefig(output_dir/'pred_vs_true_{}.pdf'.format(name))
```

In [42]:

```
Y_test_pred_linear_12 = linear_model_12.predict(X_test_12)
plot_pred_vs_true(Y_test_pred_linear_12, Y_test, 'pred_vs_true_linear_12', title
```



## 4.2. DNN regression

I experimented by varying the number of layers, dropout rate, learning rate, batch size, and adding batch normalization layers—the model below achieves pretty good performance.

In [43]:

```
# The Sequential model will do just fine here
# (functional API is more flexible though)

def setup_dnn_model(n_outputs):
    model = keras.Sequential([
        #layers.BatchNormalization(),
        layers.Dense(200, activation='relu'),
        layers.Dropout(0.05),
        layers.Dense(200, activation='relu'),
        layers.Dropout(0.05),
        layers.Dense(200, activation='relu'),
        layers.Dropout(0.05),
        layers.Dense(100, activation='relu'),
        layers.Dropout(0.05),
        layers.Dense(100, activation='relu'),
        layers.Dense(n_outputs)
    ])

    model.compile(loss='mean_squared_error',
                  optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3, decay=5e-
    return model
```

In [44]:

```
# dnn config
dnn_tag = "dnn_200x3_100x2_05dropout"
dnn_epochs = 1000
dnn_batch_size = 32
```

Predict all 6 forces from 12 input features

In [45]:

```
dnn_model_12 = setup_dnn_model(Y_train.shape[-1])
dnn_model_12_tag = "{}_12features".format(dnn_tag)
```

In [46]:

```
\%%time

early_stop = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=20)
save_every_epoch = tf.keras.callbacks.ModelCheckpoint(output_dir/'dnn_12_tmp.h5'

history_dnn_12 = dnn_model_12.fit(
    X_train_12, Y_train,
    validation_data=(X_val_12, Y_val),
    batch_size=dnn_batch_size,
    epochs=dnn_epochs,
    callbacks=[save_every_epoch]
    #callbacks=[early_stop, save_every_epoch]
    #verbose=0,
)
dnn_model_12.summary()
with open(output_dir/'history_dnn_12.pickle', 'wb') as f:
    pickle.dump(history_dnn_12.history, f)
```

```
Epoch 1/1000
1074/1074 [=====] - 3s 2ms/step - loss: 0.0176 - val_lo
ss: 0.0060
Epoch 2/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0062 - val_lo
ss: 0.0035
Epoch 3/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0043 - val_lo
ss: 0.0029
Epoch 4/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0037 - val_lo
ss: 0.0024
Epoch 5/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0032 - val_lo
ss: 0.0024
Epoch 6/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0031 - val_lo
ss: 0.0023
Epoch 7/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0029 - val_lo
ss: 0.0024
Epoch 8/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0029 - val_lo
ss: 0.0022
Epoch 9/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0027 - val_lo
ss: 0.0023
Epoch 10/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0027 - val_lo
ss: 0.0021
Epoch 11/1000
1074/1074 [=====] - 3s 2ms/step - loss: 0.0026 - val_lo
ss: 0.0023
Epoch 12/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0026 - val_lo
ss: 0.0019
Epoch 13/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0025 - val_lo
ss: 0.0019
Epoch 14/1000
```

```

1074/1074 [=====] - 2s 2ms/step - loss: 7.6267e-04 - va
l_loss: 9.5025e-04
Epoch 990/1000
1074/1074 [=====] - 2s 2ms/step - loss: 7.7188e-04 - va
l_loss: 9.3111e-04
Epoch 991/1000
1074/1074 [=====] - 2s 2ms/step - loss: 7.6617e-04 - va
l_loss: 9.4429e-04
Epoch 992/1000
1074/1074 [=====] - 2s 2ms/step - loss: 7.5996e-04 - va
l_loss: 8.3807e-04
Epoch 993/1000
1074/1074 [=====] - 2s 2ms/step - loss: 7.6634e-04 - va
l_loss: 9.0566e-04
Epoch 994/1000
1074/1074 [=====] - 2s 2ms/step - loss: 7.6265e-04 - va
l_loss: 9.0445e-04
Epoch 995/1000
1074/1074 [=====] - 2s 2ms/step - loss: 7.6349e-04 - va
l_loss: 9.5535e-04
Epoch 996/1000
1074/1074 [=====] - 2s 2ms/step - loss: 7.6452e-04 - va
l_loss: 8.6927e-04
Epoch 997/1000
1074/1074 [=====] - 2s 2ms/step - loss: 7.7106e-04 - va
l_loss: 8.5391e-04
Epoch 998/1000
1074/1074 [=====] - 2s 2ms/step - loss: 7.6998e-04 - va
l_loss: 9.2541e-04
Epoch 999/1000
1074/1074 [=====] - 2s 2ms/step - loss: 7.6056e-04 - va
l_loss: 9.1517e-04
Epoch 1000/1000
1074/1074 [=====] - 2s 2ms/step - loss: 7.5942e-04 - va
l_loss: 9.2650e-04
Model: "sequential"

```

Layer (type)	Output Shape	Param #
<hr/>		
dense (Dense)	(None, 200)	2600
dropout (Dropout)	(None, 200)	0
dense_1 (Dense)	(None, 200)	40200
dropout_1 (Dropout)	(None, 200)	0
dense_2 (Dense)	(None, 200)	40200
dropout_2 (Dropout)	(None, 200)	0
dense_3 (Dense)	(None, 100)	20100
dropout_3 (Dropout)	(None, 100)	0
dense_4 (Dense)	(None, 100)	10100
dense_5 (Dense)	(None, 6)	606
<hr/>		
Total params:	113,806	
Trainable params:	113,806	
Non-trainable params:	0	

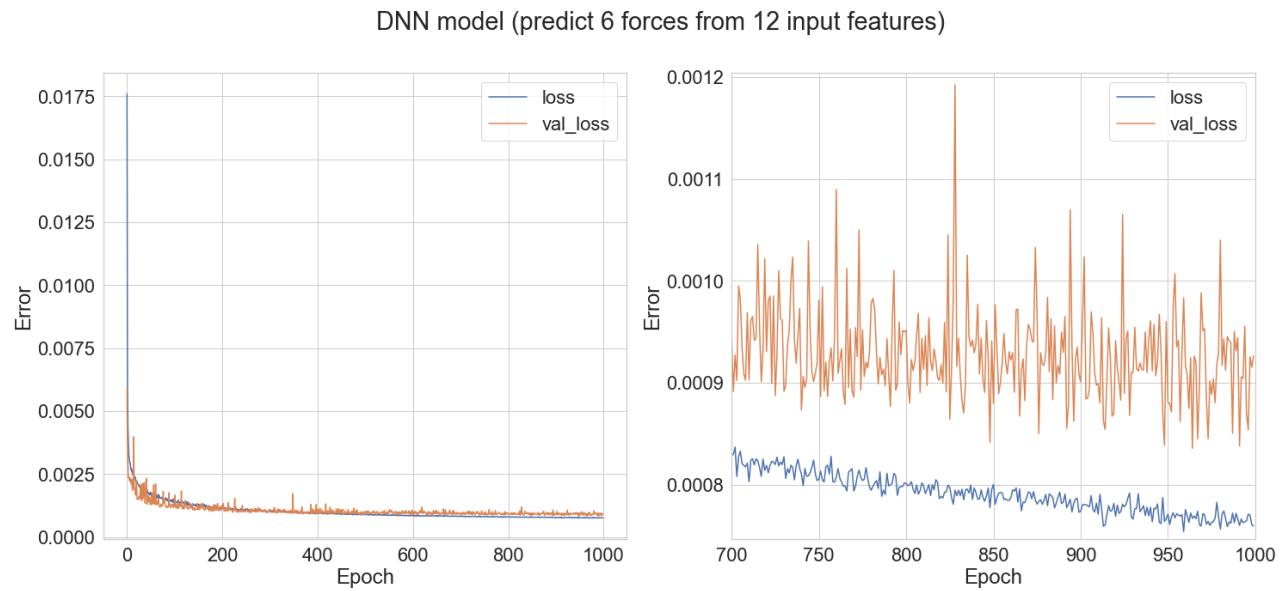
CPU times: user 1h 4min 14s, sys: 26min 52s, total: 1h 31min 7s  
Wall time: 34min 3s

## Save DNN model

```
In [47]: dnn_model_12.save(output_dir.format(dnn_model_12_tag, timestamp))
```

## Plot loss vs. epoch

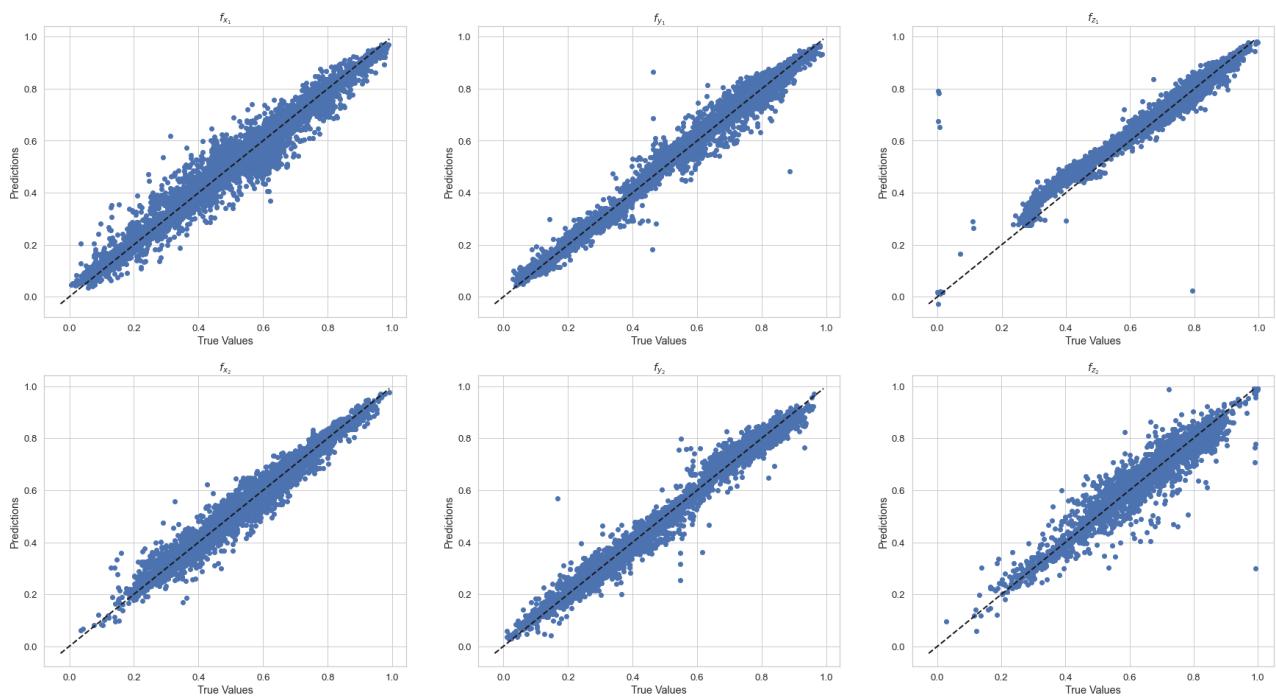
```
In [168...]: plot_loss(history_dnn_12, 'loss_{}.pdf'.format(dnn_model_12_tag), 'DNN model (pr
```



```
In [49]: # save model loss on test set for evaluation section below
test_results['dnn_12'] = dnn_model_12.evaluate(X_test_12, Y_test, verbose=0)
```

## Compare prediction vs. true values for the test set

```
In [50]: Y_test_pred_dnn_12 = dnn_model_12.predict(X_test_12)
plot_pred_vs_true(Y_test_pred_dnn_12, Y_test, 'pred_vs_true_{}'.format(dnn_model
```



## Predict all 6 forces from 36 input features (12+2x12)

Up to 2nd order derivatives for 12 input features.

```
In [51]: dnn_model_36 = setup_dnn_model(Y_train.shape[-1])
dnn_model_36_tag = "{}_36features".format(dnn_tag)
```

```
In [52]: %%time

early_stop = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=20)
save_every_epoch = tf.keras.callbacks.ModelCheckpoint(output_dir/'dnn_36_tmp.h5'

history_dnn_36 = dnn_model_36.fit(
    X_train_36, Y_train,
    validation_data=(X_val_36, Y_val),
    batch_size=dnn_batch_size,
    epochs=dnn_epochs,
    callbacks=[save_every_epoch]
    #callbacks=[early_stop, save_every_epoch]
    #verbose=0,
)
dnn_model_36.summary()
with open(output_dir/'history_dnn_36.pickle', 'wb') as f:
    pickle.dump(history_dnn_36.history, f)
```

```
Epoch 1/1000
1074/1074 [=====] - 3s 2ms/step - loss: 0.0194 - val_lo
ss: 0.0080
Epoch 2/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0074 - val_lo
ss: 0.0086
Epoch 3/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0052 - val_lo
ss: 0.0029
Epoch 4/1000
```

```
1074/1074 [=====] - 2s 2ms/step - loss: 3.4267e-04 - va  
l_loss: 3.4996e-04  
Epoch 980/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.3457e-04 - va  
l_loss: 3.6023e-04  
Epoch 981/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.4331e-04 - va  
l_loss: 3.3237e-04  
Epoch 982/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.3601e-04 - va  
l_loss: 3.5746e-04  
Epoch 983/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.3687e-04 - va  
l_loss: 3.2193e-04  
Epoch 984/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.3548e-04 - va  
l_loss: 3.4794e-04  
Epoch 985/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.3764e-04 - va  
l_loss: 3.5151e-04  
Epoch 986/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.3984e-04 - va  
l_loss: 3.2758e-04  
Epoch 987/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.3693e-04 - va  
l_loss: 3.2723e-04  
Epoch 988/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.4318e-04 - va  
l_loss: 3.5631e-04  
Epoch 989/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.3061e-04 - va  
l_loss: 3.0576e-04  
Epoch 990/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.3398e-04 - va  
l_loss: 3.3314e-04  
Epoch 991/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.3909e-04 - va  
l_loss: 3.3064e-04  
Epoch 992/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.3926e-04 - va  
l_loss: 3.4767e-04  
Epoch 993/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.4109e-04 - va  
l_loss: 3.5436e-04  
Epoch 994/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.5116e-04 - va  
l_loss: 3.8458e-04  
Epoch 995/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.2572e-04 - va  
l_loss: 3.2416e-04  
Epoch 996/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.2906e-04 - va  
l_loss: 3.3718e-04  
Epoch 997/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.3980e-04 - va  
l_loss: 3.3999e-04  
Epoch 998/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.3308e-04 - va  
l_loss: 3.2166e-04  
Epoch 999/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.3036e-04 - va  
l_loss: 3.4935e-04  
Epoch 1000/1000  
1074/1074 [=====] - 2s 2ms/step - loss: 3.4566e-04 - va  
l_loss: 3.8959e-04
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
<hr/>		
dense_6 (Dense)	(None, 200)	7400
dropout_4 (Dropout)	(None, 200)	0
dense_7 (Dense)	(None, 200)	40200
dropout_5 (Dropout)	(None, 200)	0
dense_8 (Dense)	(None, 200)	40200
dropout_6 (Dropout)	(None, 200)	0
dense_9 (Dense)	(None, 100)	20100
dropout_7 (Dropout)	(None, 100)	0
dense_10 (Dense)	(None, 100)	10100
dense_11 (Dense)	(None, 6)	606
<hr/>		
Total params:	118,606	
Trainable params:	118,606	
Non-trainable params:	0	

CPU times: user 1h 6min 19s, sys: 28min 27s, total: 1h 34min 46s  
Wall time: 35min 26s

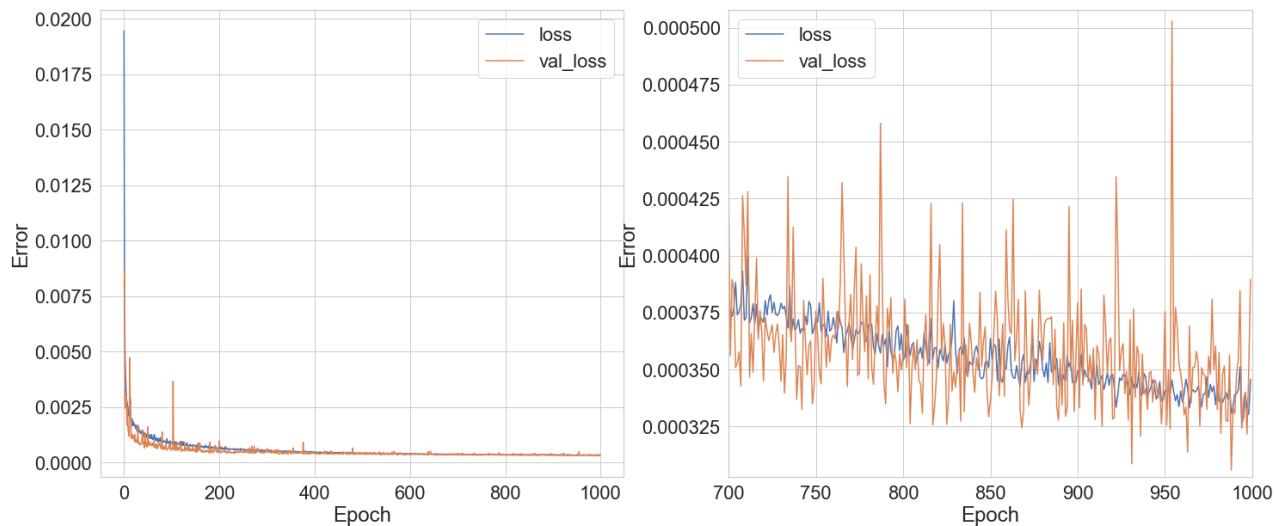
## Save DNN model

```
In [53]: dnn_model_36.save(output_dir.format(dnn_model_36_tag, timestamp))
```

## Plot loss vs. epoch

```
In [169]: plot_loss(history_dnn_36, 'loss_{0}.pdf'.format(dnn_model_36_tag), 'DNN model (pr
```

DNN model (predict 6 forces from 36 input features)



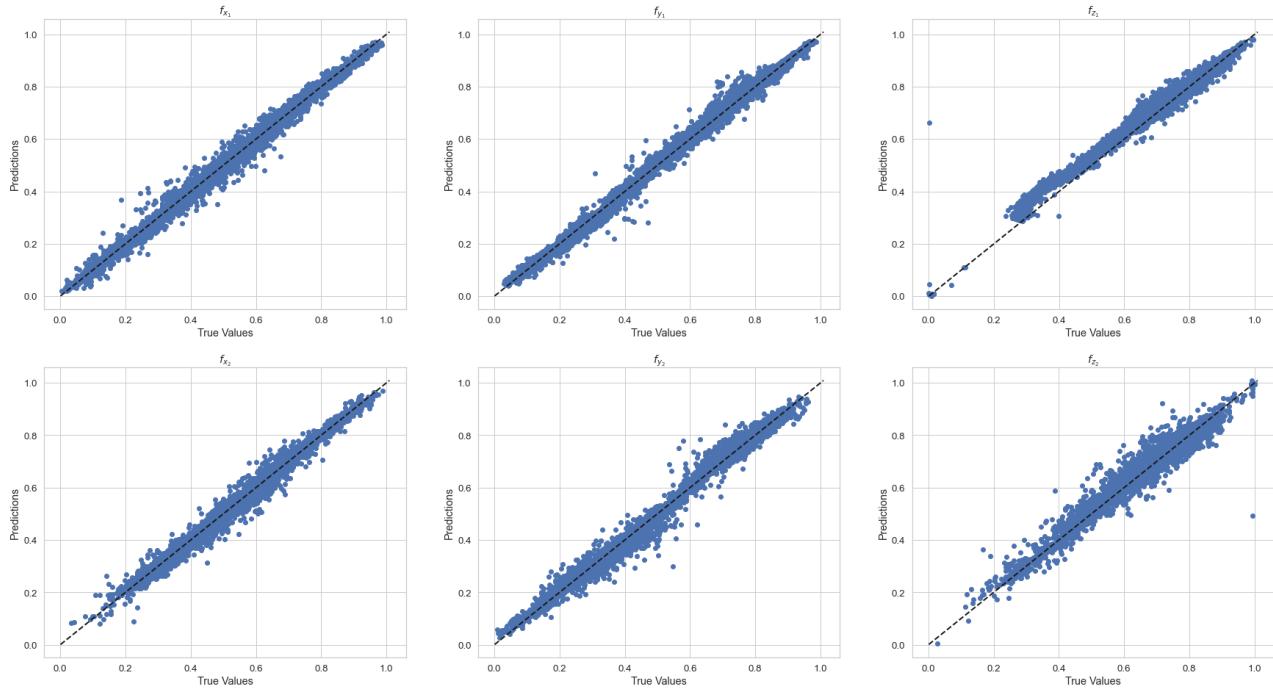
```
In [55]:
```

```
# save model loss on test set for evaluation section below
test_results['dnn_36'] = dnn_model_36.evaluate(X_test_36, Y_test, verbose=0)
```

## Compare prediction vs. true values for the test set

In [56]:

```
Y_test_pred_dnn_36 = dnn_model_36.predict(X_test_36)
plot_pred_vs_true(Y_test_pred_dnn_36, Y_test, 'pred_vs_true_{}'.format(dnn_model))
```



## Predict all 6 forces from 84 input features (12+6x12)

Up to 6th order derivatives for 12 input features.

In [57]:

```
dnn_model_84 = setup_dnn_model(Y_train.shape[-1])
dnn_model_84_tag = "{}_84features".format(dnn_tag)
```

In [58]:

```
%%time

early_stop = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=20)
save_every_epoch = tf.keras.callbacks.ModelCheckpoint(output_dir/'dnn_84_tmp.h5')

history_dnn_84 = dnn_model_84.fit(
    X_train, Y_train,
    validation_data=(X_val, Y_val),
    batch_size = dnn_batch_size,
    epochs=dnn_epochs,
    callbacks=[save_every_epoch]
    #callbacks=[early_stop, save_every_epoch]
    #verbose=0,
)
dnn_model_84.summary()
with open(output_dir/'history_dnn_84.pickle', 'wb') as f:
    pickle.dump(history_dnn_84.history, f)
```

```
Epoch 1/1000
1074/1074 [=====] - 3s 2ms/step - loss: 0.0219 - val_loss: 0.0087
Epoch 2/1000
1074/1074 [=====] - 3s 2ms/step - loss: 0.0099 - val_loss: 0.0075
Epoch 3/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0081 - val_loss: 0.0067
Epoch 4/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0066 - val_loss: 0.0052
Epoch 5/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0058 - val_loss: 0.0042
Epoch 6/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0051 - val_loss: 0.0040
Epoch 7/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0044 - val_loss: 0.0043
Epoch 8/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0045 - val_loss: 0.0025
Epoch 9/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0039 - val_loss: 0.0048
Epoch 10/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0034 - val_loss: 0.0026
Epoch 11/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0034 - val_loss: 0.0023
Epoch 12/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0033 - val_loss: 0.0018
Epoch 13/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0034 - val_loss: 0.0023
Epoch 14/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0032 - val_loss: 0.0019
Epoch 15/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0028 - val_loss: 0.0041
Epoch 16/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0027 - val_loss: 0.0018
Epoch 17/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0031 - val_loss: 0.0019
Epoch 18/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0024 - val_loss: 0.0018
Epoch 19/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0028 - val_loss: 0.0014
Epoch 20/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0027 - val_loss: 0.0031
Epoch 21/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0029 - val_loss: 0.0024
Epoch 22/1000
1074/1074 [=====] - 2s 2ms/step - loss: 0.0026 - val_loss:
```

```

l_loss: 4.7526e-04
Epoch 998/1000
1074/1074 [=====] - 2s 2ms/step - loss: 4.1401e-04 - va
l_loss: 4.6309e-04
Epoch 999/1000
1074/1074 [=====] - 2s 2ms/step - loss: 4.2859e-04 - va
l_loss: 4.5703e-04
Epoch 1000/1000
1074/1074 [=====] - 2s 2ms/step - loss: 4.3420e-04 - va
l_loss: 4.7985e-04
Model: "sequential_2"

```

Layer (type)	Output Shape	Param #
<hr/>		
dense_12 (Dense)	(None, 200)	17000
dropout_8 (Dropout)	(None, 200)	0
dense_13 (Dense)	(None, 200)	40200
dropout_9 (Dropout)	(None, 200)	0
dense_14 (Dense)	(None, 200)	40200
dropout_10 (Dropout)	(None, 200)	0
dense_15 (Dense)	(None, 100)	20100
dropout_11 (Dropout)	(None, 100)	0
dense_16 (Dense)	(None, 100)	10100
dense_17 (Dense)	(None, 6)	606
<hr/>		
Total params:	128,206	
Trainable params:	128,206	
Non-trainable params:	0	

CPU times: user 1h 6min 27s, sys: 27min 37s, total: 1h 34min 5s  
Wall time: 35min 49s

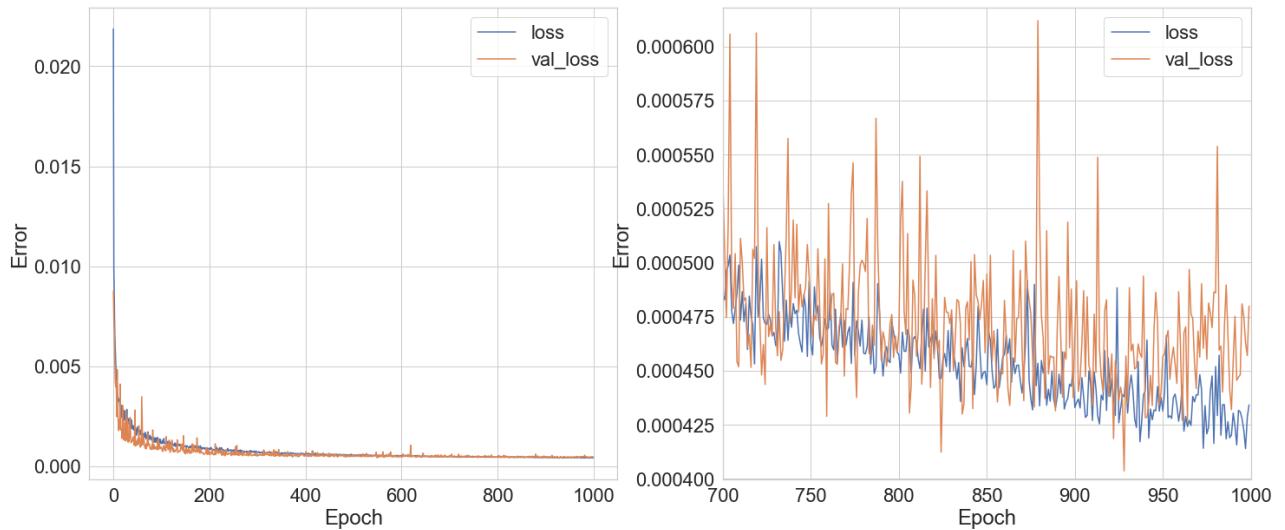
## Save DNN model

```
In [59]: dnn_model_84.save(output_dir.format("{}_{}").format(dnn_model_84_tag, timestamp))
```

## Plot loss vs. epoch

```
In [170...]: plot_loss(history_dnn_84, 'loss_{}.pdf'.format(dnn_model_84_tag), 'DNN model (pr
```

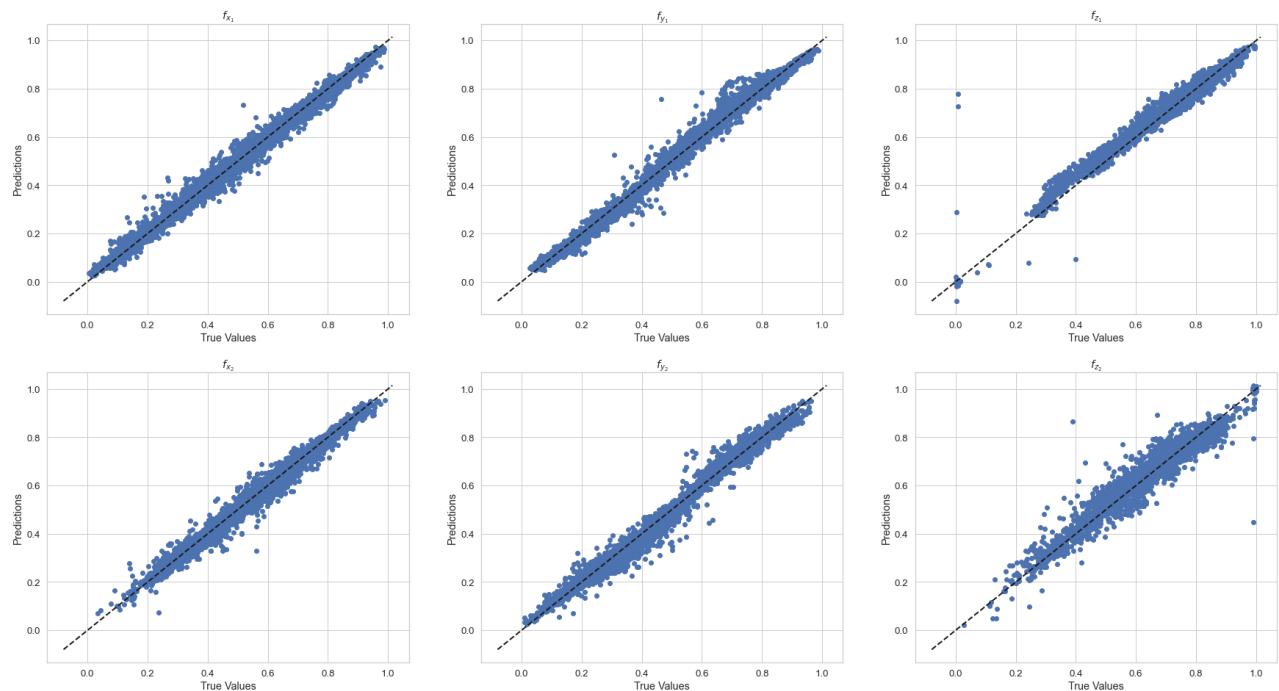
olsson\_solution  
DNN model (predict 6 forces from 84 input features)



```
In [61]: # save model loss on test set for evaluation section below
test_results['dnn_84'] = dnn_model_84.evaluate(X_test, Y_test, verbose=0)
```

### Compare prediction vs. true values for the test set

```
In [62]: Y_test_pred_dnn_84 = dnn_model_84.predict(X_test)
plot_pred_vs_true(Y_test_pred_dnn_84, Y_test, 'pred_vs_true_{}'.format(dnn_model
```



## 4.3 RNN regression

I experimented with LSTM, GRU, and SimpleRNN layers. Performance was relatively similar, but the LSTM seemed to do slightly better.

As is shown in section 5, an RNN did a significantly better job predicted unseen data than a DNN, especially for the dataset not seen during training.

```
In [63]: def setup_rnn_model(n_outputs):
    model = keras.Sequential([
        #layers.BatchNormalization(),
        layers.LSTM(100, activation='relu', return_sequences=True),
        #layers.SimpleRNN(100, activation='relu', return_sequences=True),
        #layers.GRU(100, activation='relu', return_sequences=True),
        layers.Dropout(0.05),
        layers.LSTM(100, activation='relu', return_sequences=False),
        #layers.SimpleRNN(100, activation='relu', return_sequences=False),
        #layers.GRU(100, activation='relu', return_sequences=False),
        layers.Dropout(0.05),
        layers.Dense(100, activation='relu'),
        layers.Dense(n_outputs)
    ])

    model.compile(loss='mean_squared_error',
                  optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3, decay=5e-5))
    return model
```

```
In [64]: # rnn config
rnn_tag = "rnn_lstm200x2_dense100x1_0p5drop"
rnn_epochs = 300
rnn_batch_size = 32
```

```
In [65]: rnn_model_12 = setup_rnn_model(Y_seq_train.shape[-1])
rnn_model_12_tag = "{}_12features".format(rnn_tag)
```

```
In [66]: rnn_model_36 = setup_rnn_model(Y_seq_train.shape[-1])
rnn_model_36_tag = "{}_36features".format(rnn_tag)
```

```
In [67]: rnn_model_84 = setup_rnn_model(Y_seq_train.shape[-1])
rnn_model_84_tag = "{}_84features".format(rnn_tag)
```

```
In [68]: %%time

early_stop = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=20)
save_every_epoch = tf.keras.callbacks.ModelCheckpoint(output_dir/'rnn_12_tmp.h5')

history_rnn_12 = rnn_model_12.fit(
    X_seq_train_12, Y_seq_train,
    validation_data=(X_seq_val_12, Y_seq_val),
    batch_size=rnn_batch_size,
    epochs=rnn_epochs,
    callbacks=[save_every_epoch]
    #callbacks=[early_stop, save_every_epoch]
    #verbose=0,
)
rnn_model_12.summary()
```

```
with open(output_dir/'history_rnn_12.pickle', 'wb') as f:  
    pickle.dump(history_rnn_12.history, f)
```

```
Epoch 1/300  
1074/1074 [=====] - 21s 18ms/step - loss: 0.0202 - val_loss: 0.0076  
Epoch 2/300  
1074/1074 [=====] - 19s 18ms/step - loss: 0.0069 - val_loss: 0.0050  
Epoch 3/300  
1074/1074 [=====] - 18s 16ms/step - loss: 0.0048 - val_loss: 0.0044  
Epoch 4/300  
1074/1074 [=====] - 18s 17ms/step - loss: 0.0039 - val_loss: 0.0035  
Epoch 5/300  
1074/1074 [=====] - 18s 17ms/step - loss: 0.0035 - val_loss: 0.0029  
Epoch 6/300  
1074/1074 [=====] - 18s 16ms/step - loss: 0.0031 - val_loss: 0.0027  
Epoch 7/300  
1074/1074 [=====] - 19s 18ms/step - loss: 0.0028 - val_loss: 0.0027  
Epoch 8/300  
1074/1074 [=====] - 19s 18ms/step - loss: 0.0027 - val_loss: 0.0025  
Epoch 9/300  
1074/1074 [=====] - 17s 16ms/step - loss: 0.0025 - val_loss: 0.0020  
Epoch 10/300  
1074/1074 [=====] - 17s 16ms/step - loss: 0.0023 - val_loss: 0.0023  
Epoch 11/300  
1074/1074 [=====] - 18s 17ms/step - loss: 0.0022 - val_loss: 0.0019  
Epoch 12/300  
1074/1074 [=====] - 18s 17ms/step - loss: 0.0021 - val_loss: 0.0016  
Epoch 13/300  
1074/1074 [=====] - 17s 16ms/step - loss: 0.0020 - val_loss: 0.0031  
Epoch 14/300  
1074/1074 [=====] - 19s 18ms/step - loss: 0.0020 - val_loss: 0.0017  
Epoch 15/300  
1074/1074 [=====] - 18s 17ms/step - loss: 0.0018 - val_loss: 0.0015  
Epoch 16/300  
1074/1074 [=====] - 18s 17ms/step - loss: 0.0018 - val_loss: 0.0021  
Epoch 17/300  
1074/1074 [=====] - 19s 17ms/step - loss: 0.0017 - val_loss: 0.0013  
Epoch 18/300  
1074/1074 [=====] - 20s 18ms/step - loss: 0.0016 - val_loss: 0.0014  
Epoch 19/300  
1074/1074 [=====] - 18s 17ms/step - loss: 0.0016 - val_loss: 0.0013  
Epoch 20/300  
1074/1074 [=====] - 21s 19ms/step - loss: 0.0015 - val_loss: 0.0012  
Epoch 21/300
```

```

olsson_solution

dropout_12 (Dropout)      (None, 20, 100)      0
lstm_1 (LSTM)             (None, 100)          80400
dropout_13 (Dropout)      (None, 100)          0
dense_18 (Dense)          (None, 100)          10100
dense_19 (Dense)          (None, 6)            606
=====
Total params: 136,306
Trainable params: 136,306
Non-trainable params: 0

CPU times: user 4h 5min 55s, sys: 57min 20s, total: 5h 3min 16s
Wall time: 1h 30min 53s

```

In [69]:

```

%%time

early_stop = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=20)
save_every_epoch = tf.keras.callbacks.ModelCheckpoint(output_dir/'rnn_36_tmp.h5'

history_rnn_36 = rnn_model_36.fit(
    X_seq_train_36, Y_seq_train,
    validation_data=(X_seq_val_36, Y_seq_val),
    batch_size=rnn_batch_size,
    epochs=rnn_epochs,
    callbacks=[save_every_epoch]
    #callbacks=[early_stop, save_every_epoch]
    #verbose=0,
)
rnn_model_36.summary()
with open(output_dir/'history_rnn_36.pickle', 'wb') as f:
    pickle.dump(history_rnn_36.history, f)

```

```

Epoch 1/300
1074/1074 [=====] - 20s 17ms/step - loss: 0.0171 - val_
loss: 0.0094
Epoch 2/300
1074/1074 [=====] - 18s 16ms/step - loss: 0.0068 - val_
loss: 0.0046
Epoch 3/300
1074/1074 [=====] - 17s 16ms/step - loss: 0.0044 - val_
loss: 0.0038
Epoch 4/300
1074/1074 [=====] - 18s 16ms/step - loss: 0.0033 - val_
loss: 0.0033
Epoch 5/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0028 - val_
loss: 0.0030
Epoch 6/300
1074/1074 [=====] - 18s 16ms/step - loss: 0.0026 - val_
loss: 0.0026
Epoch 7/300
1074/1074 [=====] - 18s 16ms/step - loss: 0.0024 - val_
loss: 0.0027
Epoch 8/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0022 - val_
loss: 0.0047
Epoch 9/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0021 - val_
loss: 0.0016

```

```

val_loss: 2.9235e-04
Epoch 292/300
1074/1074 [=====] - 18s 17ms/step - loss: 2.6069e-04 -
val_loss: 3.0082e-04
Epoch 293/300
1074/1074 [=====] - 18s 17ms/step - loss: 2.6264e-04 -
val_loss: 3.3003e-04
Epoch 294/300
1074/1074 [=====] - 18s 17ms/step - loss: 2.6386e-04 -
val_loss: 3.1733e-04
Epoch 295/300
1074/1074 [=====] - 18s 16ms/step - loss: 2.5697e-04 -
val_loss: 3.3751e-04
Epoch 296/300
1074/1074 [=====] - 17s 16ms/step - loss: 2.7905e-04 -
val_loss: 2.8409e-04
Epoch 297/300
1074/1074 [=====] - 19s 17ms/step - loss: 2.5475e-04 -
val_loss: 2.9780e-04
Epoch 298/300
1074/1074 [=====] - 18s 17ms/step - loss: 2.6043e-04 -
val_loss: 3.1433e-04
Epoch 299/300
1074/1074 [=====] - 17s 16ms/step - loss: 2.6849e-04 -
val_loss: 3.2881e-04
Epoch 300/300
1074/1074 [=====] - 18s 17ms/step - loss: 2.5034e-04 -
val_loss: 2.7966e-04
Model: "sequential_4"

```

Layer (type)	Output Shape	Param #
<hr/>		
lstm_2 (LSTM)	(None, 20, 100)	54800
dropout_14 (Dropout)	(None, 20, 100)	0
lstm_3 (LSTM)	(None, 100)	80400
dropout_15 (Dropout)	(None, 100)	0
dense_20 (Dense)	(None, 100)	10100
dense_21 (Dense)	(None, 6)	606
<hr/>		
Total params: 145,906		
Trainable params: 145,906		
Non-trainable params: 0		

---

```
CPU times: user 4h 7min 58s, sys: 56min 32s, total: 5h 4min 31s
Wall time: 1h 29min 17s
```

In [70]:

```
%time

early_stop = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=20)
save_every_epoch = tf.keras.callbacks.ModelCheckpoint(output_dir/'rnn_84_tmp.h5'

history_rnn_84 = rnn_model_84.fit(
    X_seq_train, Y_seq_train,
    validation_data=(X_seq_val, Y_seq_val),
    batch_size=rnn_batch_size,
    epochs=rnn_epochs,
    callbacks=[save_every_epoch]
```

```
#callbacks=[early_stop, save_every_epoch]
#verbose=0,
)
rnn_model_84.summary()
with open(output_dir/'history_rnn_84.pickle', 'wb') as f:
    pickle.dump(history_rnn_84.history, f)
```

```
Epoch 1/300
1074/1074 [=====] - 20s 17ms/step - loss: 0.0185 - val_
loss: 0.0110
Epoch 2/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0083 - val_
loss: 0.0058
Epoch 3/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0054 - val_
loss: 0.0082
Epoch 4/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0047 - val_
loss: 0.0033
Epoch 5/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0038 - val_
loss: 0.0030
Epoch 6/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0034 - val_
loss: 0.0042
Epoch 7/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0031 - val_
loss: 0.0085
Epoch 8/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0027 - val_
loss: 0.0021
Epoch 9/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0025 - val_
loss: 0.0035
Epoch 10/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0024 - val_
loss: 0.0039
Epoch 11/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0024 - val_
loss: 0.0025
Epoch 12/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0020 - val_
loss: 0.0016
Epoch 13/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0019 - val_
loss: 0.0019
Epoch 14/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0020 - val_
loss: 0.0016
Epoch 15/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0020 - val_
loss: 0.0014
Epoch 16/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0020 - val_
loss: 0.0017
Epoch 17/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0018 - val_
loss: 0.0020
Epoch 18/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0017 - val_
loss: 0.0015
Epoch 19/300
1074/1074 [=====] - 18s 17ms/step - loss: 0.0016 - val_
loss: 0.0013
```

Layer (type)	Output Shape	Param #
lstm_4 (LSTM)	(None, 20, 100)	74000
dropout_16 (Dropout)	(None, 20, 100)	0
lstm_5 (LSTM)	(None, 100)	80400
dropout_17 (Dropout)	(None, 100)	0
dense_22 (Dense)	(None, 100)	10100
dense_23 (Dense)	(None, 6)	606

---

Total params: 165,106  
Trainable params: 165,106  
Non-trainable params: 0

---

CPU times: user 4h 22min 8s, sys: 1h 27s, total: 5h 22min 35s  
Wall time: 1h 32min 51s

## Save RNN model

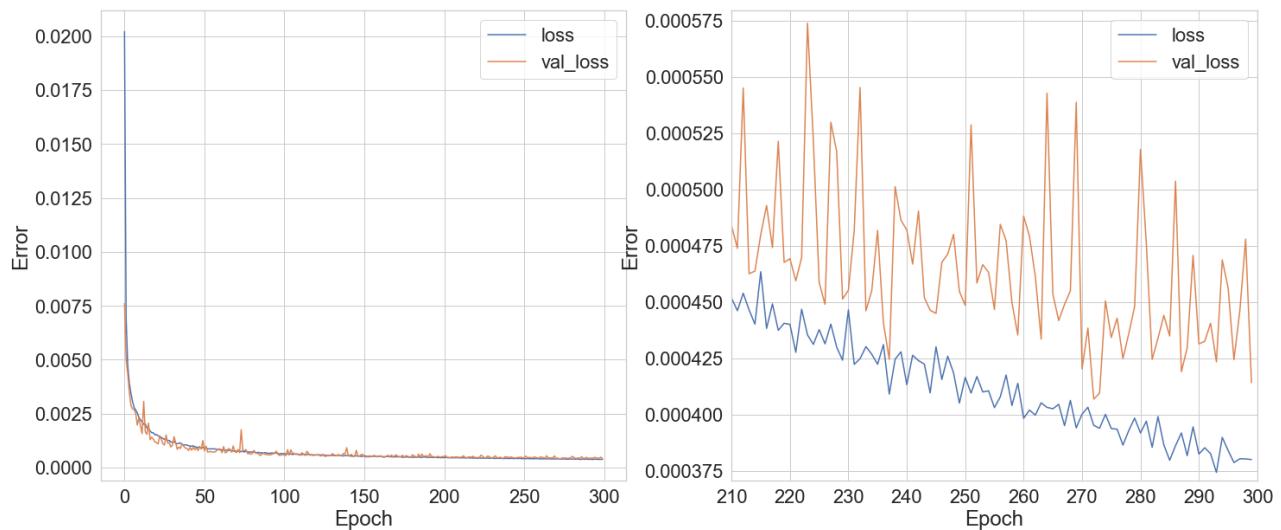
```
In [71]: rnn_model_12.save(output_dir.format(rnn_model_12_tag, timestamp), sav
In [72]: rnn_model_36.save(output_dir.format(rnn_model_36_tag, timestamp), sav
In [73]: rnn_model_84.save(output_dir.format(rnn_model_84_tag, timestamp), sav
```

## Plot loss vs. epoch

```
In [171... plot_loss(history_rnn_12, 'error_vs_epoch_{}.pdf'.format(rnn_model_12_tag), 'DNN

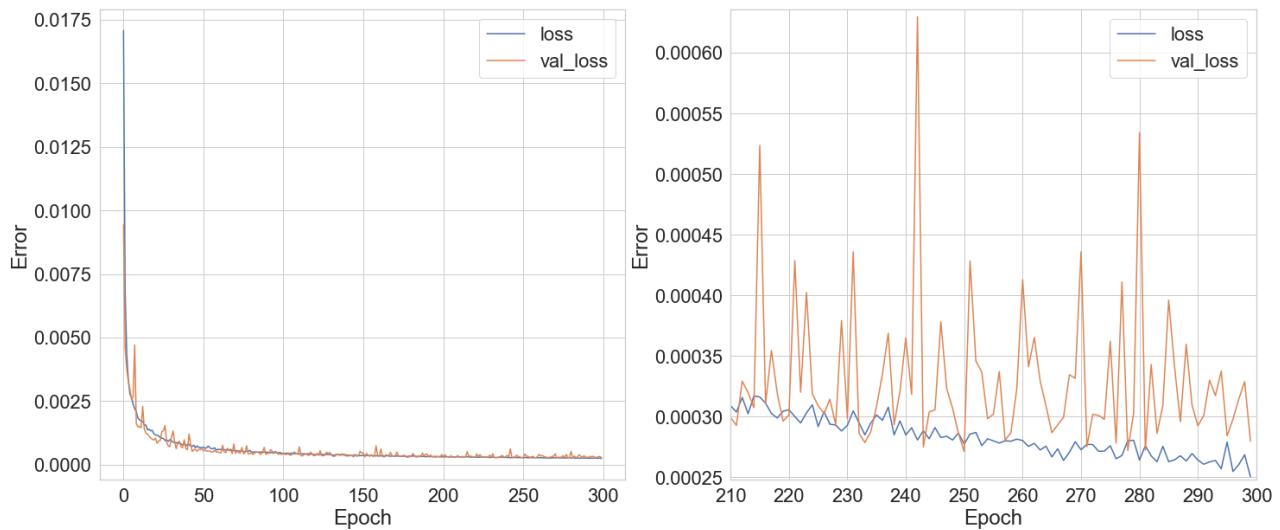
```

DNN model (predict 6 forces from 12 input features)



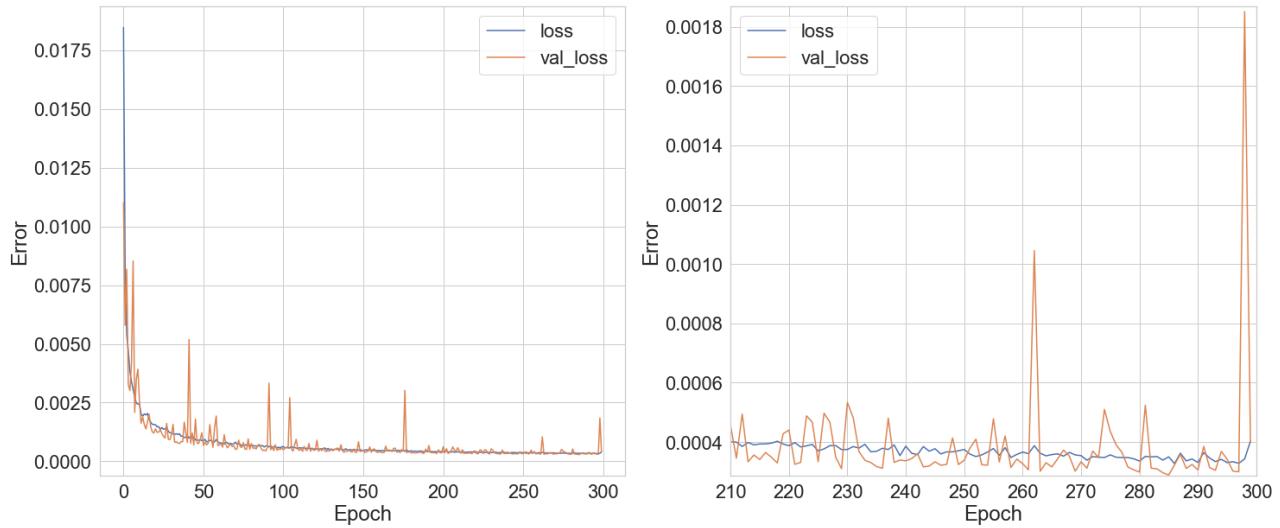
```
In [172... plot_loss(history_rnn_36, 'error_vs_epoch_{}.pdf'.format(rnn_model_36_tag), 'DNN
```

olsson\_solution  
DNN model (predict 6 forces from 36 input features)



```
In [173]: plot_loss(history_rnn_84, 'error_vs_epoch_{}.pdf'.format(rnn_model_84_tag), 'DNN')
```

DNN model (predict 6 forces from 84 input features)



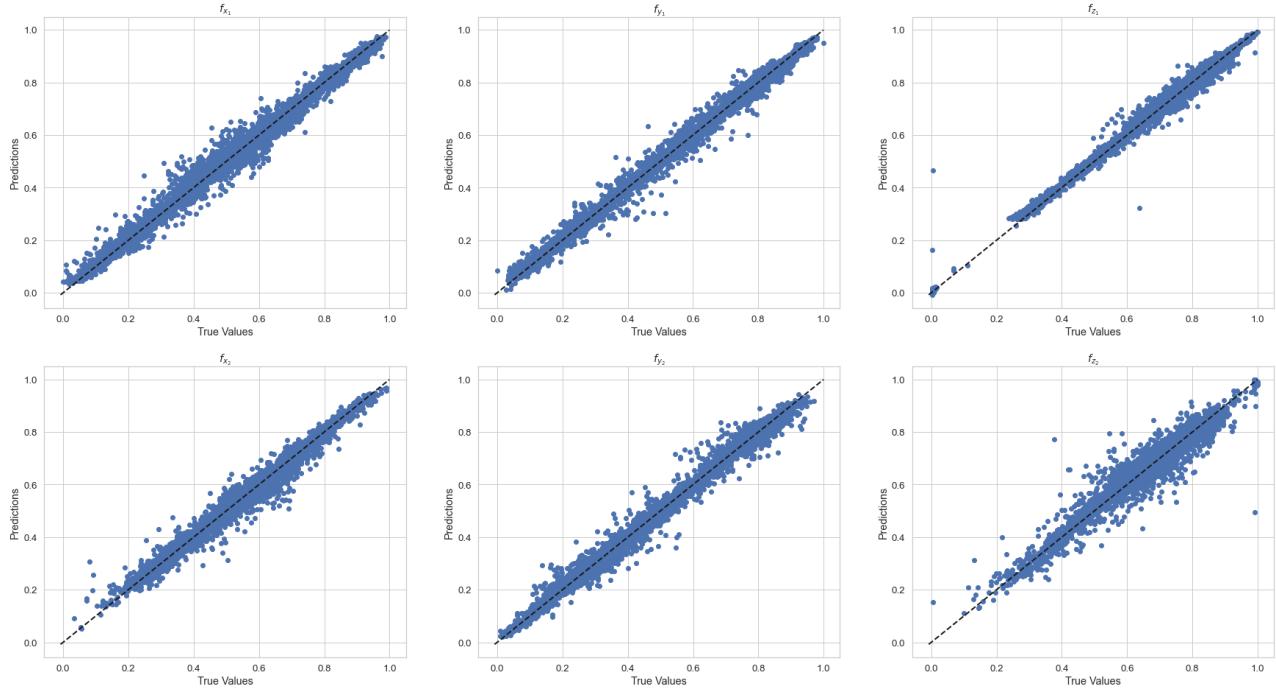
```
In [174]: # save model loss on test set for evaluation section below
test_results['rnn_12'] = rnn_model_12.evaluate(X_seq_test_12, Y_seq_test, verbose=0)
```

```
In [175]: # save model loss on test set for evaluation section below
test_results['rnn_36'] = rnn_model_36.evaluate(X_seq_test_36, Y_seq_test, verbose=0)
```

```
In [176]: # save model loss on test set for evaluation section below
test_results['rnn_84'] = rnn_model_84.evaluate(X_seq_test, Y_seq_test, verbose=0)
```

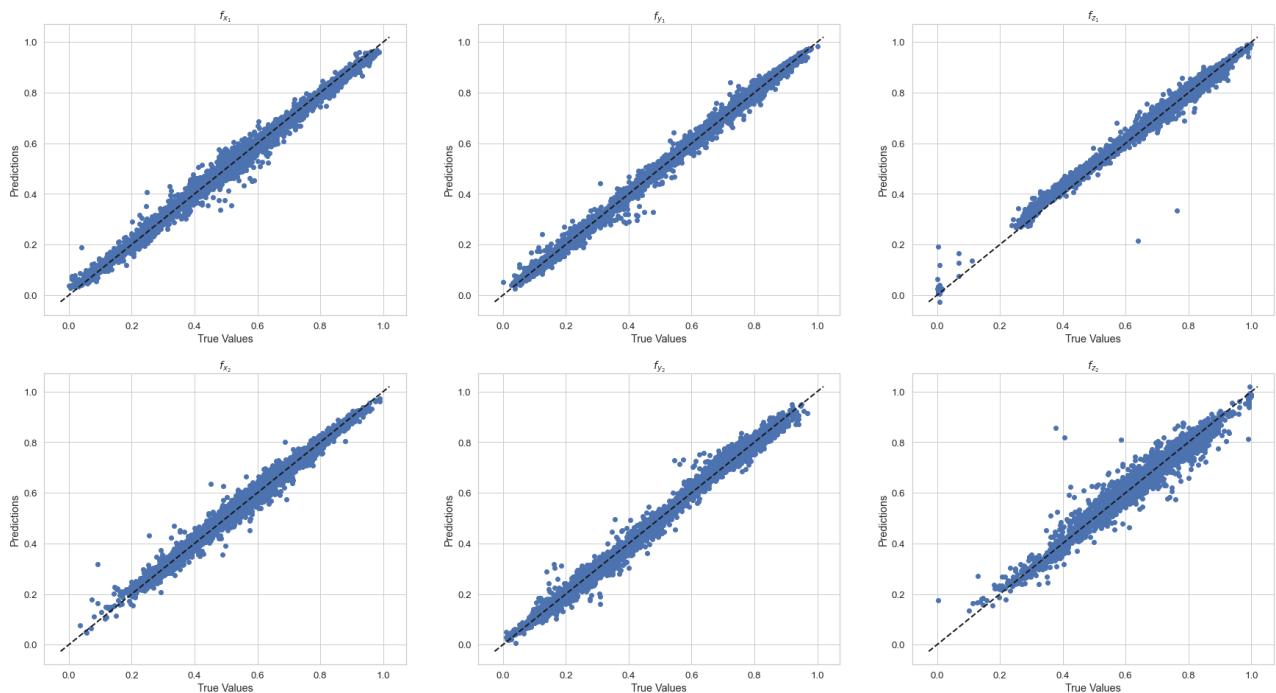
### Compare prediction vs. true values for the test set

```
In [80]: Y_seq_test_pred_12 = rnn_model_12.predict(X_seq_test_12)
plot_pred_vs_true(Y_seq_test_pred_12, Y_seq_test, 'pred_vs_true_{}'.format(rnn_m
```



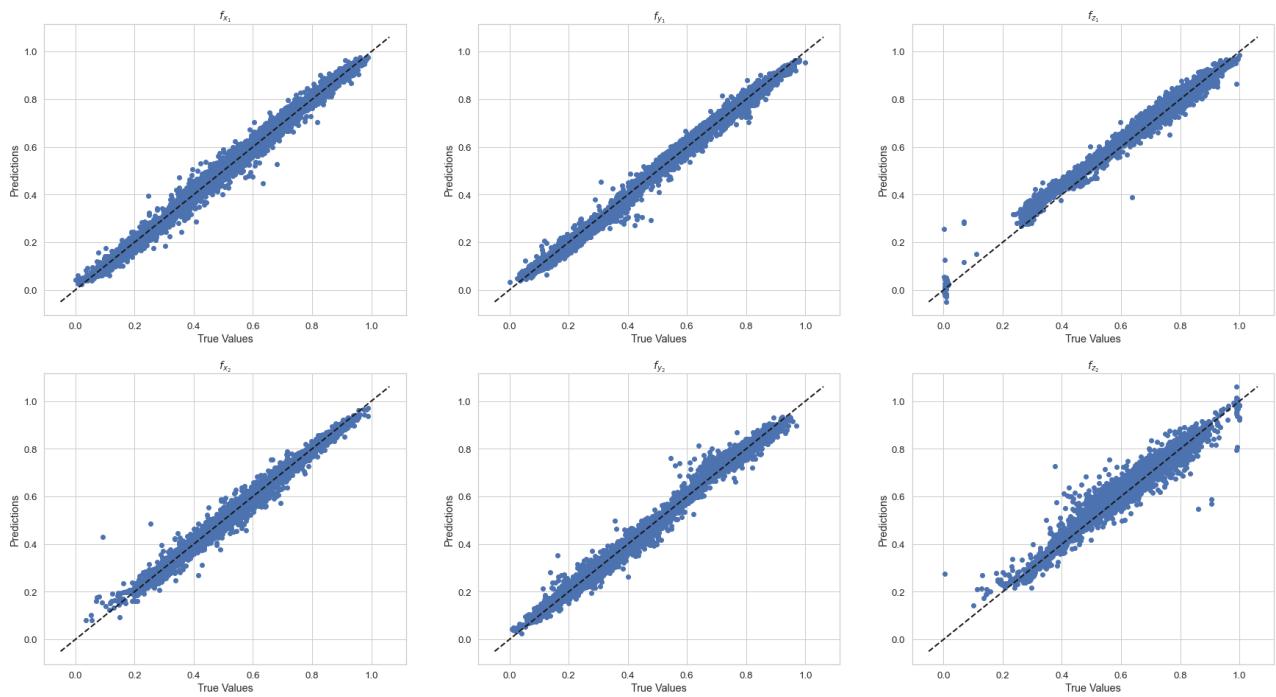
In [82]:

```
Y_seq_test_pred_36 = rnn_model_36.predict(X_seq_test_36)
plot_pred_vs_true(Y_seq_test_pred_36, Y_seq_test, 'pred_vs_true_{}'.format(rnn_m
```



In [83]:

```
Y_seq_test_pred = rnn_model_84.predict(X_seq_test)
plot_pred_vs_true(Y_seq_test_pred, Y_seq_test, 'pred_vs_true_{}'.format(rnn_mode
```



## 5. Evaluation

### 5.1 Loss on test sets

In [84]:

```
print("loss on test sets:")
for key, val in test_results.items():
    print("- {}: {:.2e}".format(key, val))
```

```
loss on test sets:
- linear_x1: 1.47e-02
- linear_12: 2.69e-02
- dnn_12: 8.80e-04
- dnn_36: 3.78e-04
- dnn_84: 4.47e-04
- rnn_12: 3.96e-04
- rnn_36: 2.87e-04
- rnn_84: 3.88e-04
```

Loss after 1000 and 300 epochs for DNN and RNN respectively:

- linear\_x1: 1.47e-02
- linear\_12: 2.69e-02
- dnn\_12: 8.80e-04
- dnn\_36: 3.78e-04
- dnn\_84: 4.47e-04
- rnn\_12: 3.96e-04
- rnn\_36: 2.87e-04
- rnn\_84: 3.88e-04

**Loss on Test1, Test2, Test4 (no data from Test2 was included in the training)**

```
In [85]: def create_separate_test_sets(df, features, outputs, n_steps=20, feature_idx=None):

    # select relevant features and outputs
    X = df[features].to_numpy()
    Y = df[outputs].to_numpy()

    # apply scaling
    X_normed = scaler_x.transform(X)
    Y_normed = scaler_y.transform(Y)

    # handle sequences
    X_seq, Y_seq = split_sequences(X_normed, Y_normed, n_steps)

    # select indices corresponding to desired feature
    if feature_idx:
        X = X[:,feature_idx]
        X_normed = X_normed[:,feature_idx]
        X_seq = X_seq[:, :, feature_idx]

    outputs = {
        'X': X, 'Y': Y,
        'X_normed': X_normed, 'Y_normed': Y_normed,
        'X_seq_normed': X_seq, 'Y_seq_normed': Y_seq,
        'Y_seq': scaler_y.inverse_transform(Y_seq)
    }

    # calculate model predictions
    if dnn_model:
        Y_pred_normed = dnn_model.predict(X_normed)
        Y_pred = scaler_y.inverse_transform(Y_pred_normed)
        outputs['Y_pred_normed'] = Y_pred_normed
        outputs['Y_pred'] = Y_pred
    if rnn_model:
        Y_seq_pred_normed = rnn_model.predict(X_seq)
        outputs['Y_seq_pred_normed'] = Y_seq_pred_normed
        Y_seq_pred = scaler_y.inverse_transform(Y_seq_pred_normed)
        outputs['Y_seq_pred'] = Y_seq_pred

    return outputs
```

```
In [86]: tests_12 = dict()
for i,df in enumerate(datasets):
    tests_12[dataset_filenames[i]] = create_separate_test_sets(
        df, features_nth, outputs, n_steps, feature_idx=feature_idx,
        dnn_model=dnn_model_12, rnn_model=rnn_model_12)
```

```
In [87]: tests_36 = dict()
for i,df in enumerate(datasets):
    tests_36[dataset_filenames[i]] = create_separate_test_sets(
        df, features_nth, outputs, n_steps, feature_idx=feature_idx_2nd,
        dnn_model=dnn_model_36, rnn_model=rnn_model_36)
```

```
In [88]: tests_84 = dict()
for i,df in enumerate(datasets):
    tests_84[dataset_filenames[i]] = create_separate_test_sets(
```

```
df, features_nth, outputs, n_steps, feature_idx=None,
dnn_model=dnn_model_84, rnn_model=rnn_model_84)
```

In [89]:

```
print("Loss on full Test1, Test2, Test4 datasets:")
for filename in dataset_filenames:
    loss_dnn12 = dnn_model_12.evaluate(
        tests_12[filename]['X_normed'], tests_12[filename]['Y_normed'], verbose=
    loss_dnn36 = dnn_model_36.evaluate(
        tests_36[filename]['X_normed'], tests_36[filename]['Y_normed'], verbose=
    loss_dnn84 = dnn_model_84.evaluate(
        tests_84[filename]['X_normed'], tests_84[filename]['Y_normed'], verbose=
    loss_rnn12 = rnn_model_12.evaluate(
        tests_12[filename]['X_seq_normed'], tests_12[filename]['Y_seq_normed'],
    loss_rnn36 = rnn_model_36.evaluate(
        tests_36[filename]['X_seq_normed'], tests_36[filename]['Y_seq_normed'],
    loss_rnn84 = rnn_model_84.evaluate(
        tests_84[filename]['X_seq_normed'], tests_84[filename]['Y_seq_normed'],
    print("- {}:\n DNN-12: {:.2e}\n DNN-36: {:.2e}\n DNN-84: {:.2e}\n RNN-12: {:.2e}\n RNN-36: {:.2e}\n RNN-84: {:.2e}\n".format(filename, loss_dnn12, loss_dnn36, loss_dnn84, loss_rnn12, loss_rnn36, loss_rnn84))
```

Loss on full Test1, Test2, Test4 datasets:

- Test1:
 

DNN-12:	5.62e-04
DNN-36:	2.34e-04
DNN-84:	3.00e-04
RNN-12:	1.76e-04
RNN-36:	1.53e-04
RNN-84:	2.54e-04
- Test2:
 

DNN-12:	7.01e-02
DNN-36:	3.22e-02
DNN-84:	3.06e-02
RNN-12:	3.60e-02
RNN-36:	4.03e-03
RNN-84:	4.52e-03
- Test4:
 

DNN-12:	8.17e-04
DNN-36:	3.62e-04
DNN-84:	4.27e-04
RNN-12:	3.73e-04
RNN-36:	2.46e-04
RNN-84:	3.68e-04

>> Note that no part of Test2 was included during training <<

### Observations:

- The RNN does significantly better than any of the DNNs for all datasets. It has about an order of magnitude lower loss on Test2.
- DNN/RNN-84 does a bit worse than DNN/RNN-36, indicating that adding higher-order derivatives beyond acceleration makes learning harder. It may also not be worth it adding the additional derivatives. Perhaps training for more epochs could help.
- Comparing the loss on Test1 and Test4 (70 percent of Test1+Test4 was seen during training) w.r.t. loss on Test2, it is clear that all models struggled to generalize well to Test2. **There is room for improvement here.**

## 5.2. Prediction error

### Model predictions on the different test sets

In [90]:

```
tmin = 0
tmax = -1
bins=50
linewidth=3

sns.set(font_scale = 2)
sns.color_palette()
sns.set_style("whitegrid")
for filename in dataset_filenames:

    Y_err_12 = tests_12[filename]['Y_normed'] - tests_12[filename]['Y_pred_norme']
    Y_err_36 = tests_36[filename]['Y_normed'] - tests_36[filename]['Y_pred_norme']
    Y_err_84 = tests_84[filename]['Y_normed'] - tests_84[filename]['Y_pred_norme']
    Y_seq_err_12 = tests_12[filename]['Y_seq_normed'] - tests_12[filename]['Y_se
    Y_seq_err_36 = tests_36[filename]['Y_seq_normed'] - tests_36[filename]['Y_se
    Y_seq_err_84 = tests_84[filename]['Y_seq_normed'] - tests_84[filename]['Y_se

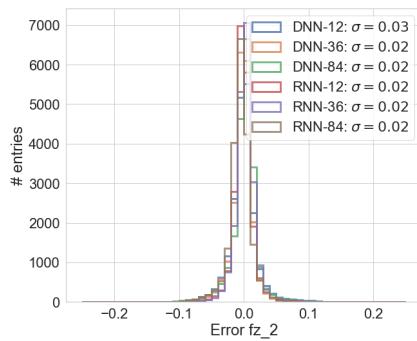
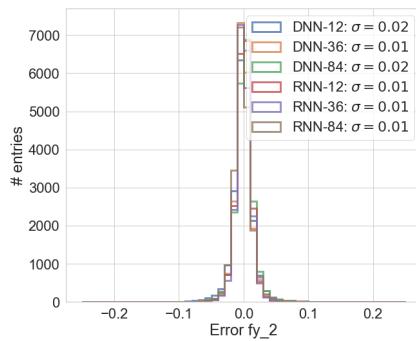
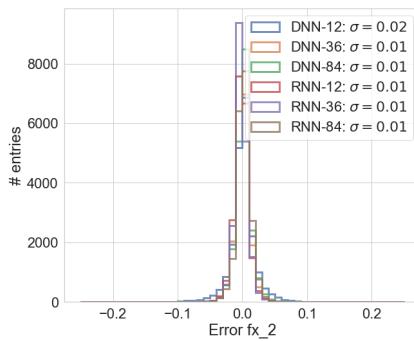
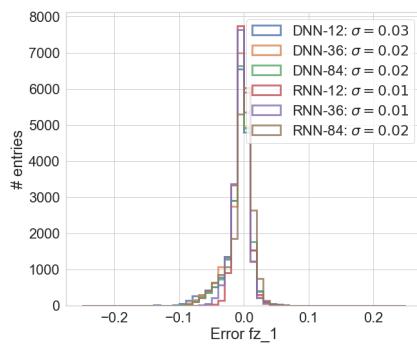
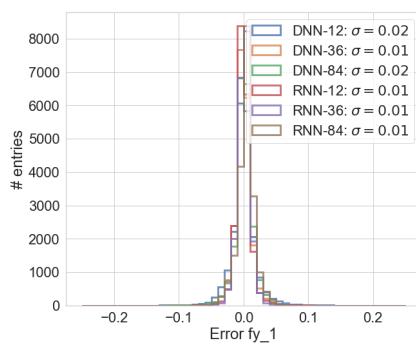
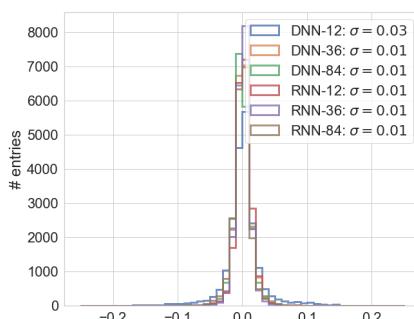
    fig = plt.figure(figsize=(28,16))
    fig.suptitle(filename, weight='bold').set_fontsize('24')
    for i in range(len(outputs)):

        label_dnn12 = "DNN-12: $\sigma={:.2f}$$".format(np.std(Y_err_12[:,i]))
        label_dnn36 = "DNN-36: $\sigma={:.2f}$$".format(np.std(Y_err_36[:,i]))
        label_dnn84 = "DNN-84: $\sigma={:.2f}$$".format(np.std(Y_err_84[:,i]))
        label_rnn12 = "RNN-12: $\sigma={:.2f}$$".format(np.std(Y_seq_err_12[:,i]))
        label_rnn36 = "RNN-36: $\sigma={:.2f}$$".format(np.std(Y_seq_err_36[:,i]))
        label_rnn84 = "RNN-84: $\sigma={:.2f}$$".format(np.std(Y_seq_err_84[:,i]))

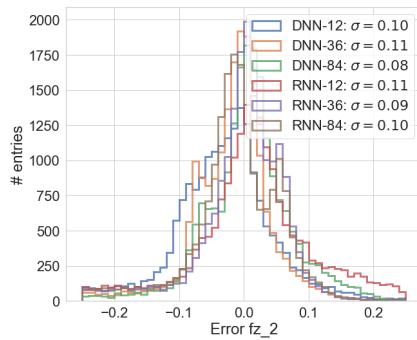
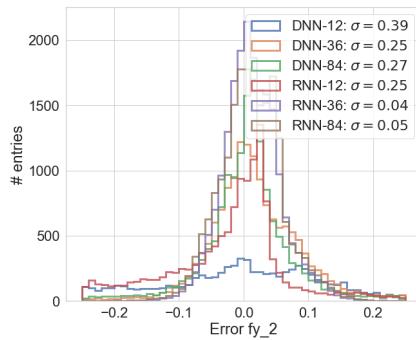
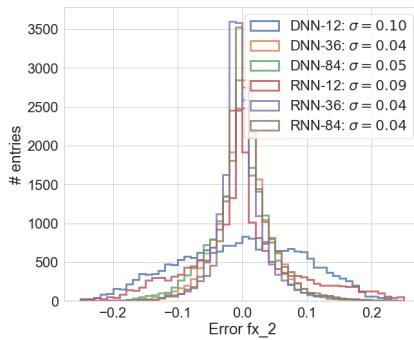
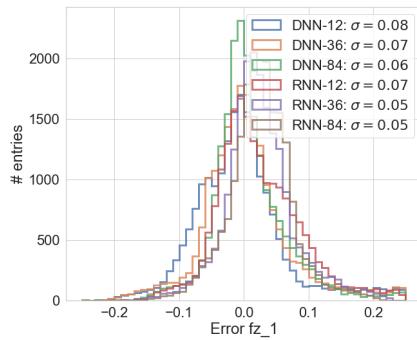
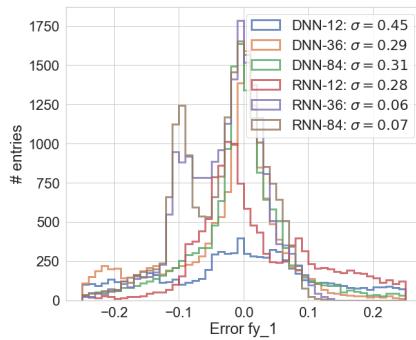
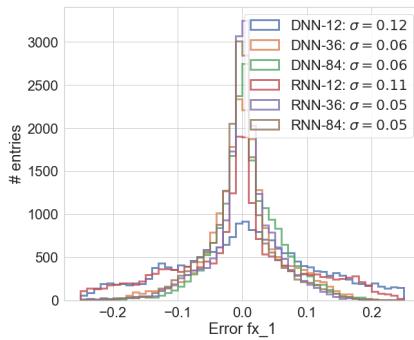
        ax = fig.add_subplot(2,3,i+1)
        ax.hist(Y_err_12[:,i], bins=bins, range=(-0.25, 0.25), alpha=0.8, histtype='stepfilled')
        ax.hist(Y_err_36[:,i], bins=bins, range=(-0.25, 0.25), alpha=0.8, histtype='stepfilled')
        ax.hist(Y_err_84[:,i], bins=bins, range=(-0.25, 0.25), alpha=0.8, histtype='stepfilled')
        ax.hist(Y_seq_err_12[:,i], bins=bins, range=(-0.25, 0.25), alpha=0.8, histtype='stepfilled')
        ax.hist(Y_seq_err_36[:,i], bins=bins, range=(-0.25, 0.25), alpha=0.8, histtype='stepfilled')
        ax.hist(Y_seq_err_84[:,i], bins=bins, range=(-0.25, 0.25), alpha=0.8, histtype='stepfilled')
        ax.set_xlabel('Error {}'.format(outputs[i]))
        ax.set_ylabel('# entries')
        ax.legend()
    plt.tight_layout()
    plt.savefig(output_dir/'{}_{}_error.pdf'.format(filename))
```

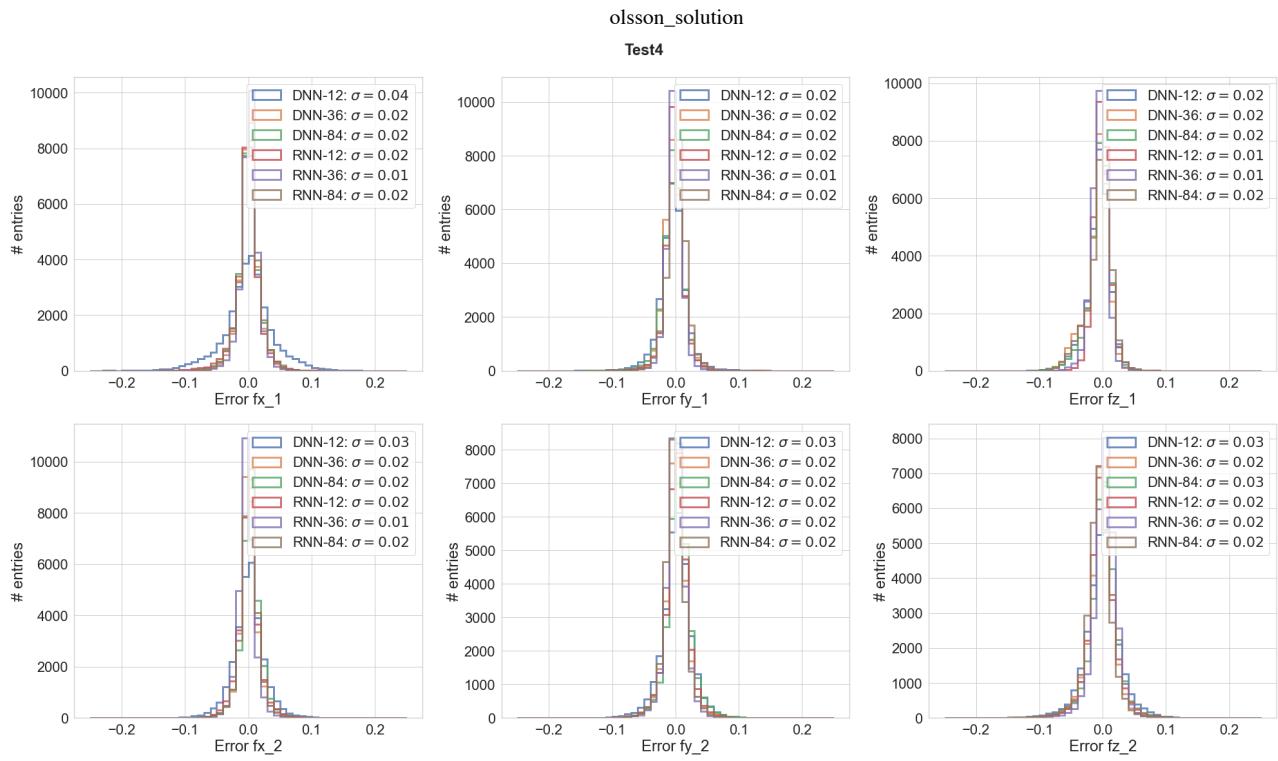
## olsson\_solution

## Test1



## Test2





### Observations:

- The RNN models outperforms the DNN models.
- The error is low (1-2% for RNN-36/84) on Test1 and Test4 (70% of Test1+Test4 was seen during training).
- The error is significantly higher (4-10%) for RNN-36/84 on Test2 (not seen during training).

## R2 score

In [91]:

```
from sklearn.metrics import r2_score

print("R2 score on Test1, Test2, Test4 datasets:")
print("(" + ".join(outputs) + ")")
for filename in dataset_filenames:

    r2_dnn12 = [ 'DNN-12:' ]
    r2_dnn36 = [ 'DNN-36:' ]
    r2_dnn84 = [ 'DNN-84:' ]
    r2_rnn12 = [ 'RNN-12:' ]
    r2_rnn36 = [ 'RNN-36:' ]
    r2_rnn84 = [ 'RNN-84:' ]
    for i in range(len(outputs)):
        r2_dnn12.append("{:.2f}".format(r2_score(tests_12[filename]['Y_normed'][i], tests_12[filename]['Y_seq_normed'][i])))
        r2_dnn36.append("{:.2f}".format(r2_score(tests_36[filename]['Y_normed'][i], tests_36[filename]['Y_seq_normed'][i])))
        r2_dnn84.append("{:.2f}".format(r2_score(tests_84[filename]['Y_normed'][i], tests_84[filename]['Y_seq_normed'][i])))
        r2_rnn12.append("{:.2f}".format(r2_score(tests_12[filename]['Y_normed'][i], tests_12[filename]['Y_seq_normed'][i])))
        r2_rnn36.append("{:.2f}".format(r2_score(tests_36[filename]['Y_normed'][i], tests_36[filename]['Y_seq_normed'][i])))
        r2_rnn84.append("{:.2f}".format(r2_score(tests_84[filename]['Y_normed'][i], tests_84[filename]['Y_seq_normed'][i])))

    print("- {}:".format(filename))
    print("\t".join(r2_dnn12))
    print("\t".join(r2_dnn36))
    print("\t".join(r2_dnn84))
    print("\t".join(r2_rnn12))
```

```
print("\t".join(r2_rnn36))
print("\t".join(r2_rnn84))
```

R2 score on Test1, Test2, Test4 datasets:  
(fx\_1 fy\_1 fz\_1 fx\_2 fy\_2 fz\_2)

- Test1:

	DNN-12:	0.97	0.99	0.97	0.98	0.99	0.96
DNN-36:	1.00	1.00	0.98	0.99	1.00	0.98	
DNN-84:	0.99	1.00	0.98	0.99	1.00	0.97	
RNN-12:	0.99	1.00	0.99	0.99	1.00	0.98	
RNN-36:	1.00	1.00	0.99	1.00	1.00	0.98	
RNN-84:	0.99	1.00	0.99	0.99	1.00	0.98	
- Test2:

	DNN-12:	0.50	-1.78	0.68	0.54	-1.85	0.07
DNN-36:	0.88	-0.26	0.76	0.91	-0.26	0.04	
DNN-84:	0.89	-0.22	0.82	0.88	-0.30	0.57	
RNN-12:	0.62	-0.27	0.76	0.61	-0.42	0.09	
RNN-36:	0.90	0.94	0.83	0.92	0.96	0.34	
RNN-84:	0.90	0.93	0.85	0.91	0.96	0.18	
- Test4:

	DNN-12:	0.95	0.99	0.98	0.97	0.99	0.90
DNN-36:	0.99	1.00	0.98	0.99	0.99	0.94	
DNN-84:	0.99	0.99	0.98	0.98	0.99	0.92	
RNN-12:	0.99	1.00	0.99	0.98	0.99	0.93	
RNN-36:	0.99	1.00	0.99	0.99	0.99	0.95	
RNN-84:	0.99	1.00	0.98	0.99	0.99	0.93	

The coefficient of determination (R2) is the proportion of the variation in the dependent variable (e.g., predicted forces) that is predictable from the independent variables (e.g., measured forces). The range is from negative infinity to +1.

- An R2 score of +1 indicates that the predictions match the observations perfectly.
- An R2 score of 0 indicates that the predictions are as good as random guesses around the mean of the observations.
- Negative R2 indicates that the predictions are worse than random.

### Observations:

- All models do pretty well on Test1 and Test4, of which a significant fraction (70%) was seen during training.
- The RNN models does significantly better on Test2, which was not seen during training.
- Although it did better than the DNN models, even the RNN struggled to predict forces in the  $z$ -direction on the Test2 dataset (especially for R2).

## Pearson correlation coefficient

In [92]:

```
def pearson(x, y):
    corr = np.corrcoef(x, y)
    return corr[0,1]
```

In [93]:

```
print("Pearson correlations for Test1, Test2, Test4 datasets:")
print("(" + ".join(outputs) + ")")
for filename in dataset_filenames:
    pearson_dnn12 = ['DNN-12:']
```

```

pearson_dnn36 = [ 'DNN-36:' ]
pearson_dnn84 = [ 'DNN-84:' ]
pearson_rnn12 = [ 'RNN-12:' ]
pearson_rnn36 = [ 'RNN-36:' ]
pearson_rnn84 = [ 'RNN-84:' ]
for i in range(len(outputs)):
    pearson_dnn12.append("{:.2f}".format(pearson(tests_12[filename])['Y_norme']))
    pearson_dnn36.append("{:.2f}".format(pearson(tests_36[filename])['Y_norme']))
    pearson_dnn84.append("{:.2f}".format(pearson(tests_84[filename])['Y_norme']))
    pearson_rnn12.append("{:.2f}".format(pearson(tests_12[filename])['Y_seq_norme']))
    pearson_rnn36.append("{:.2f}".format(pearson(tests_36[filename])['Y_seq_norme']))
    pearson_rnn84.append("{:.2f}".format(pearson(tests_84[filename])['Y_seq_norme']))

print("- {}".format(filename))
print("\t+\t".join(pearson_dnn12))
print("\t+\t".join(pearson_dnn36))
print("\t+\t".join(pearson_dnn84))
print("\t+\t".join(pearson_rnn12))
print("\t+\t".join(pearson_rnn36))
print("\t+\t".join(pearson_rnn84))

```

Pearson correlations for Test1, Test2, Test4 datasets:

(fx\_1 fy\_1 fz\_1 fx\_2 fy\_2 fz\_2)

- Test1:

	DNN-12:	0.99	1.00	0.99	0.99	1.00	0.98
DNN-12:	1.00	1.00	1.00	1.00	1.00	0.99	
DNN-36:	1.00	1.00	0.99	1.00	1.00	0.99	
DNN-84:	1.00	1.00	0.99	1.00	1.00	0.99	
RNN-12:	1.00	1.00	1.00	1.00	1.00	0.99	
RNN-36:	1.00	1.00	1.00	1.00	1.00	0.99	
RNN-84:	1.00	1.00	0.99	1.00	1.00	0.99	

- Test2:

	DNN-12:	0.75	-0.43	0.84	0.75	-0.51	0.49
DNN-12:	0.94	0.40	0.88	0.95	0.38	0.44	
DNN-36:	0.95	0.31	0.91	0.95	0.29	0.76	
DNN-84:	0.79	0.36	0.89	0.78	0.31	0.42	
RNN-12:	0.95	0.97	0.93	0.96	0.98	0.60	
RNN-36:	0.95	0.97	0.94	0.96	0.98	0.52	

- Test4:

	DNN-12:	0.98	1.00	0.99	0.98	0.99	0.95
DNN-12:	1.00	1.00	0.99	0.99	1.00	0.97	
DNN-36:	0.99	1.00	0.99	0.99	1.00	0.96	
DNN-84:	0.99	1.00	0.99	0.99	1.00	0.96	
RNN-12:	0.99	1.00	1.00	0.99	1.00	0.96	
RNN-36:	1.00	1.00	1.00	1.00	1.00	0.98	
RNN-84:	1.00	1.00	0.99	0.99	1.00	0.97	

The Pearson correlation coefficient ( $r$ ) can range from -1 to +1.

- An  $r$  of -1 indicates a perfect negative linear correlation.
- An  $r$  of 0 indicates no correlation.
- An  $r$  of +1 indicates a perfect positive linear correlation.

### Observations:

- Similar to the observations for the R2 score.

## 5.3. Time series plots of prediction vs. ground truth

In [114...]

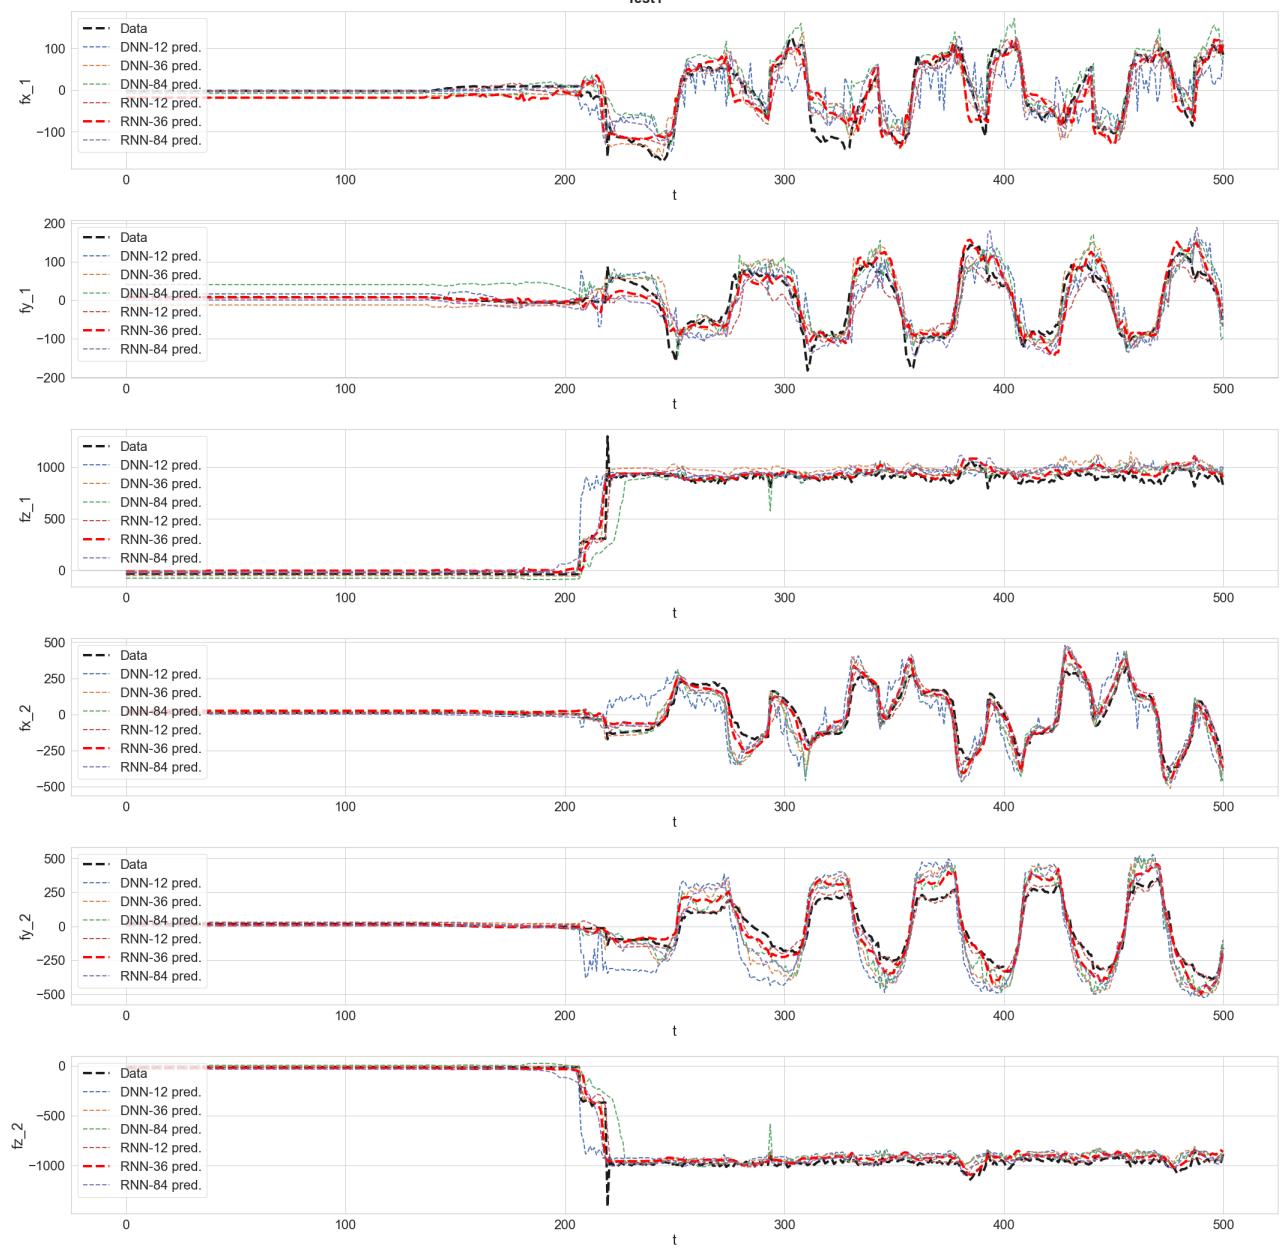
```
def plot_timeseries(tmin = 12000, tmax = 12500, linewidth=4):
```

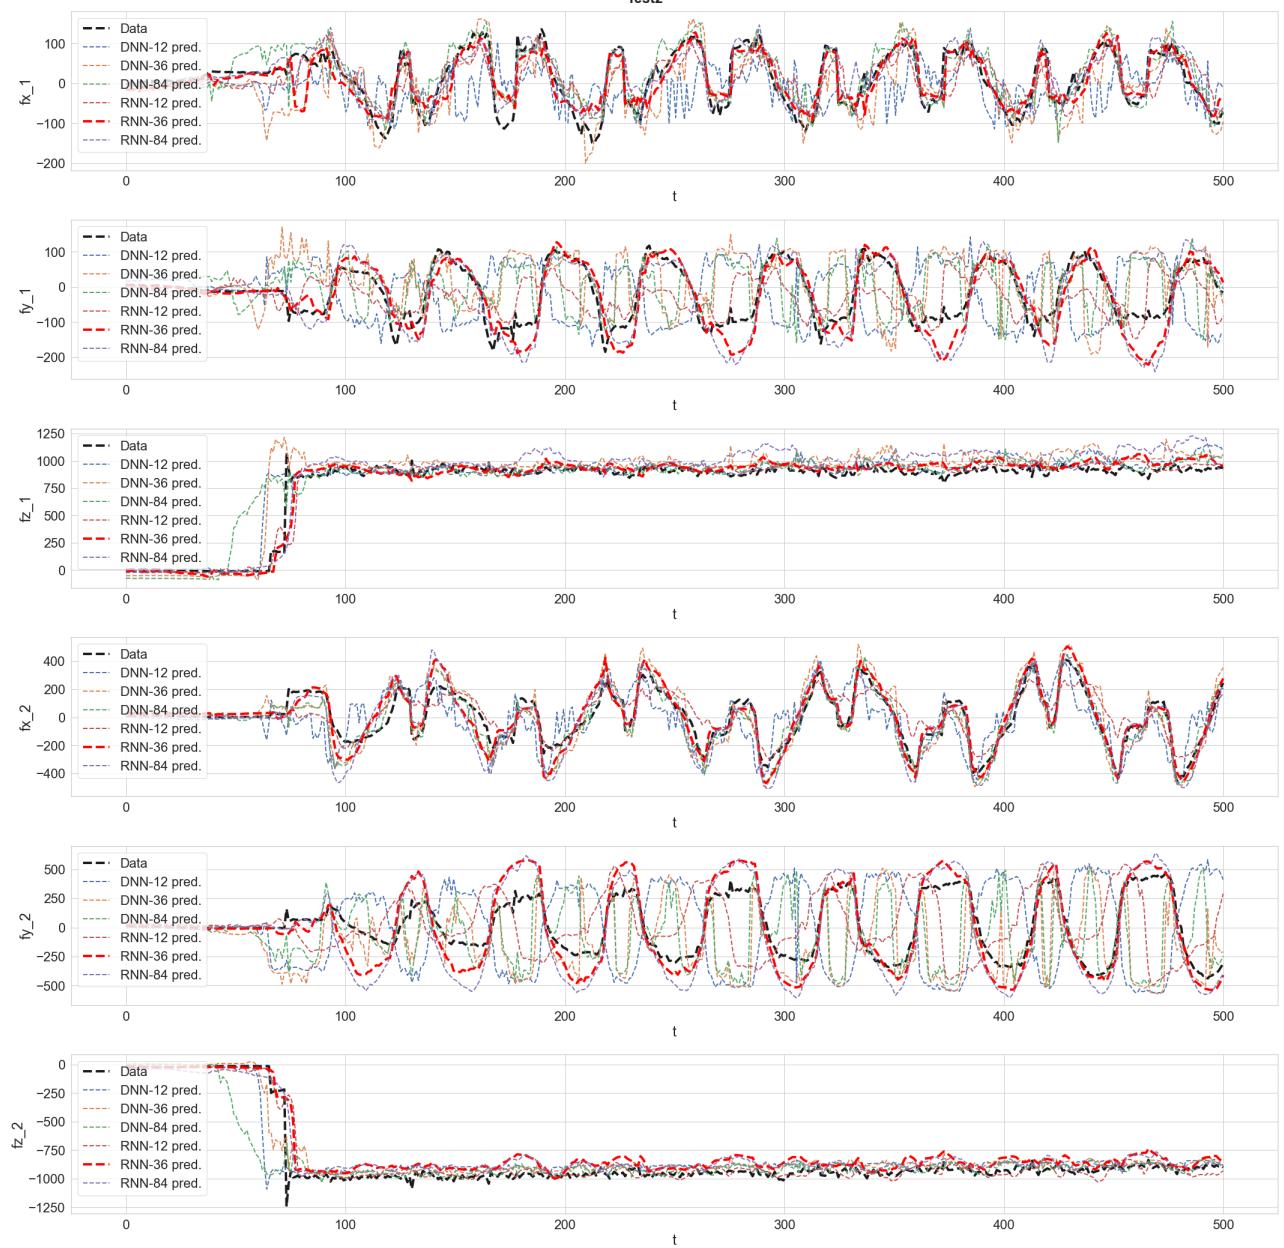
```
t = np.linspace(tmin, tmax, tmax-tmin)
sns.set(font_scale = 2)
sns.color_palette()
sns.set_style("whitegrid")
for filename in dataset_filenames:

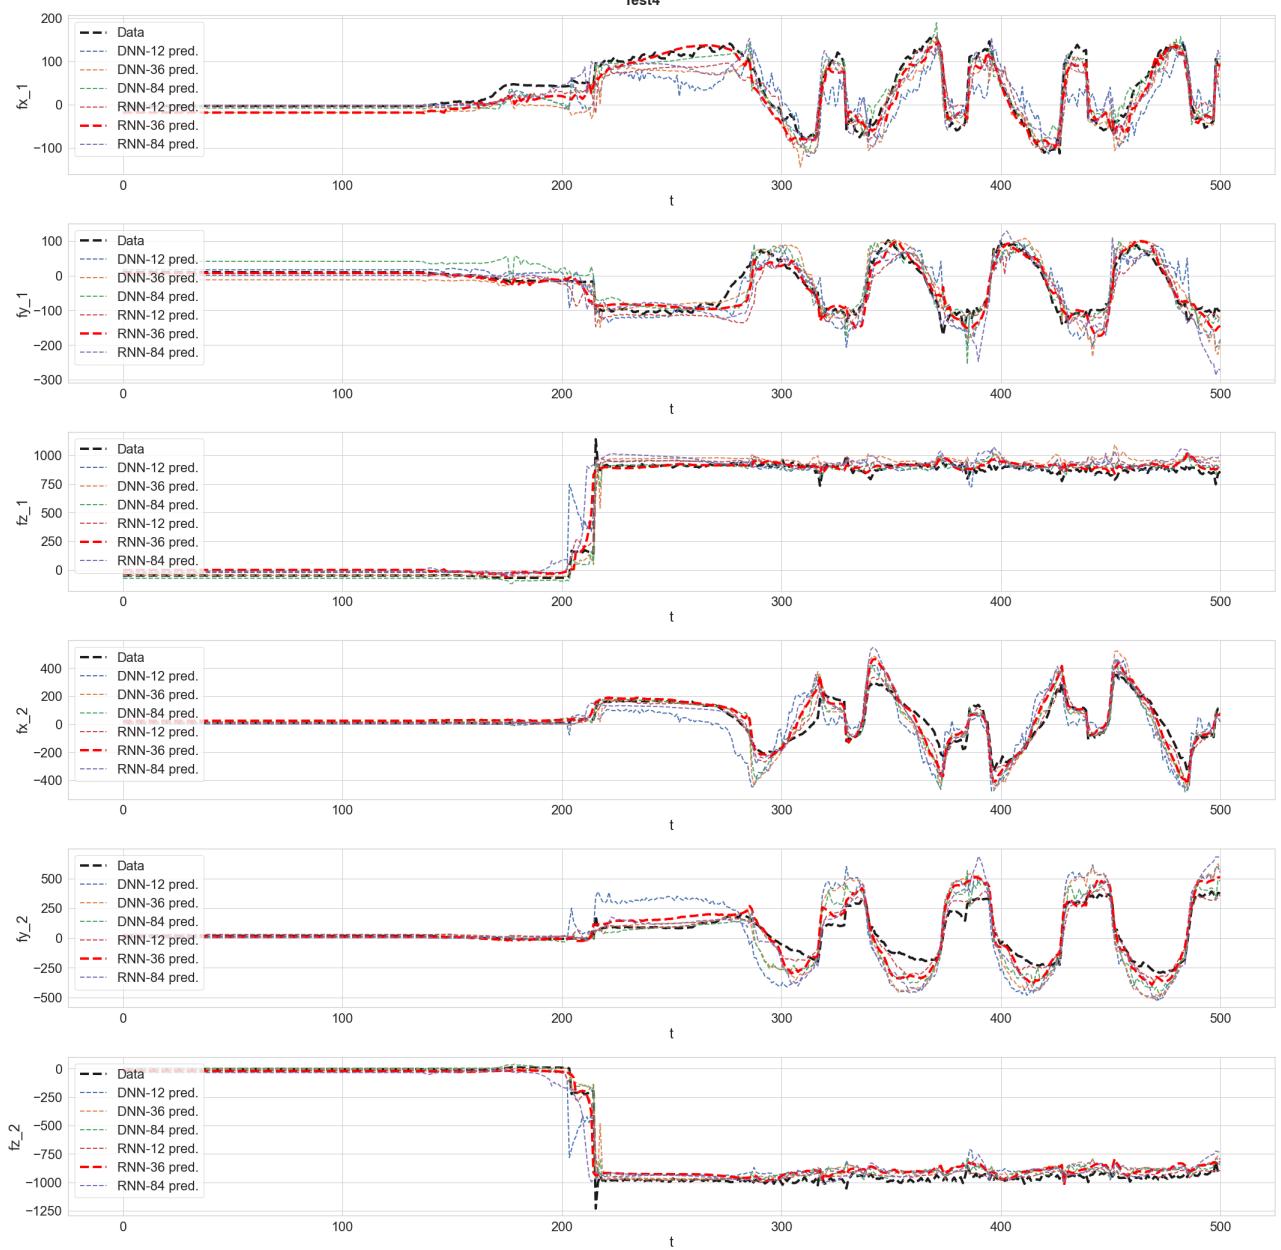
    fig = plt.figure(figsize=(30,30))
    fig.suptitle(filename, weight='bold').set_fontsize('24')
    for i in range(6):
        ax = fig.add_subplot(6, 1, i+1)
        ax.plot(t, tests_12[filename]['Y'].T[i][tmin+n_steps-1:tmax+n_steps-1])
        ax.plot(t, tests_12[filename]['Y_pred'].T[i][tmin+n_steps-1:tmax+n_steps-1])
        ax.plot(t, tests_36[filename]['Y_pred'].T[i][tmin+n_steps-1:tmax+n_steps-1])
        ax.plot(t, tests_84[filename]['Y_pred'].T[i][tmin+n_steps-1:tmax+n_steps-1])
        ax.plot(t, tests_12[filename]['Y_seq_pred'].T[i][tmin:tmax], label='Y')
        ax.plot(t, tests_36[filename]['Y_seq_pred'].T[i][tmin:tmax], label='Y')
        ax.plot(t, tests_84[filename]['Y_seq_pred'].T[i][tmin:tmax], label='Y')
        ax.set_xlabel('t')
        ax.set_ylabel(outputs[i])
        ax.legend(loc=2)
    plt.tight_layout()
    plt.savefig(output_dir/'{}_{:02d}_timeseries_t{}to{}.pdf'.format(filename, tmin,
```

In [115...]

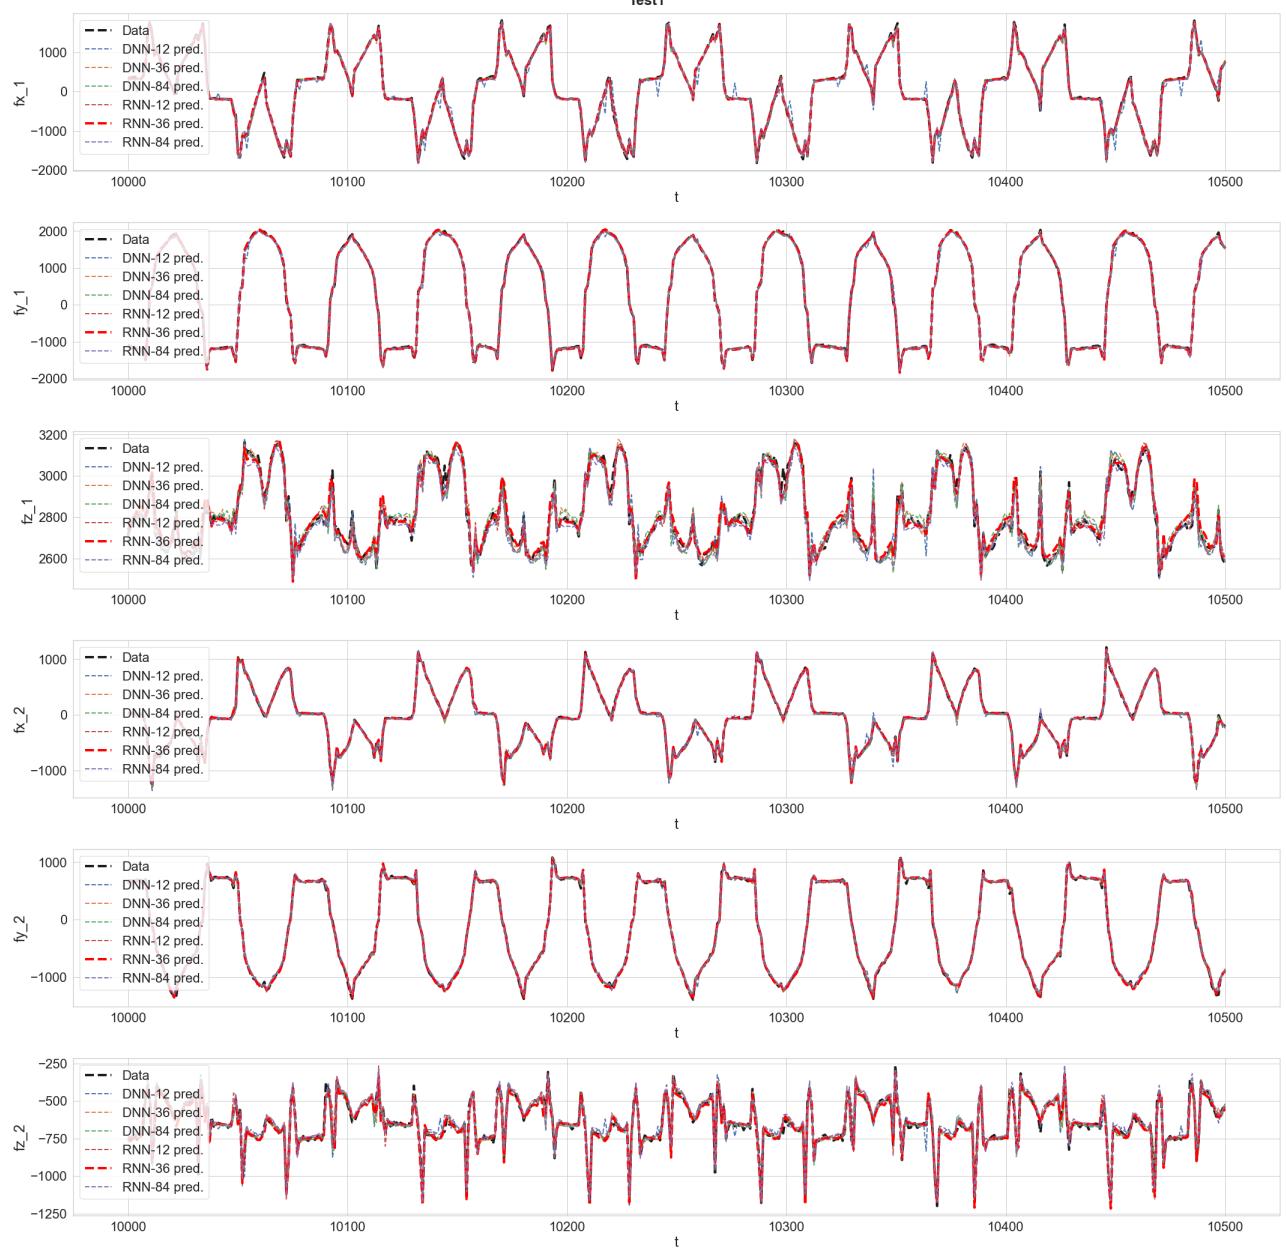
```
plot_timeseries(tmin=0, tmax=500)
```





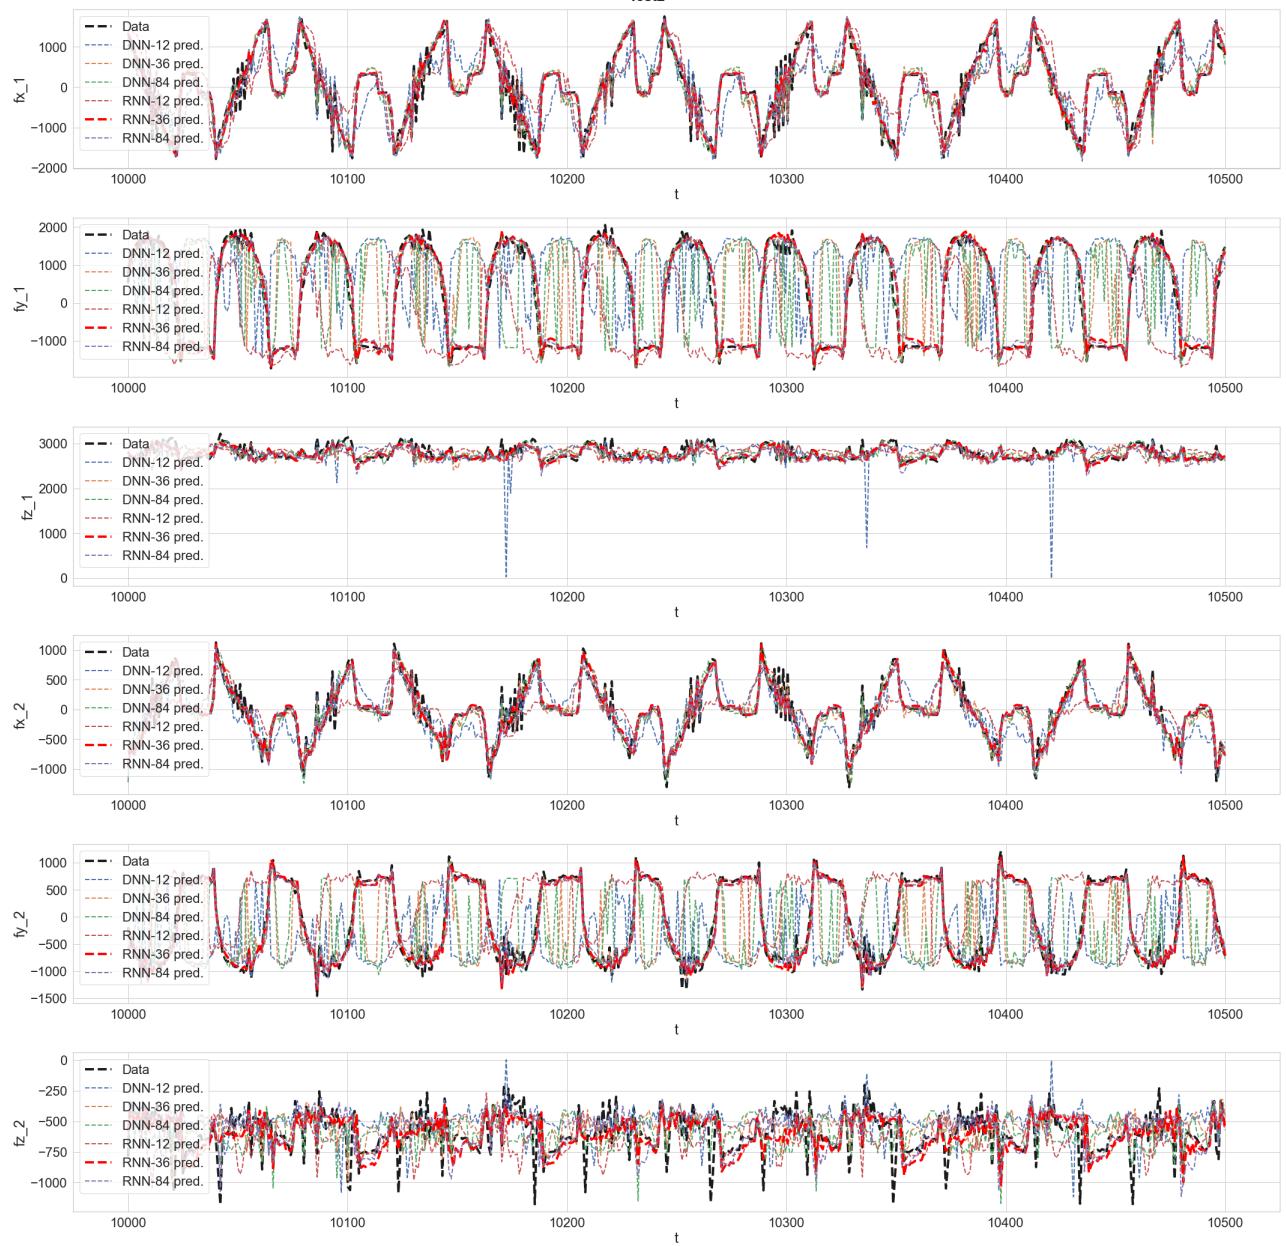


```
In [116]: plot_timeseries(tmin=10000, tmax=10500)
```



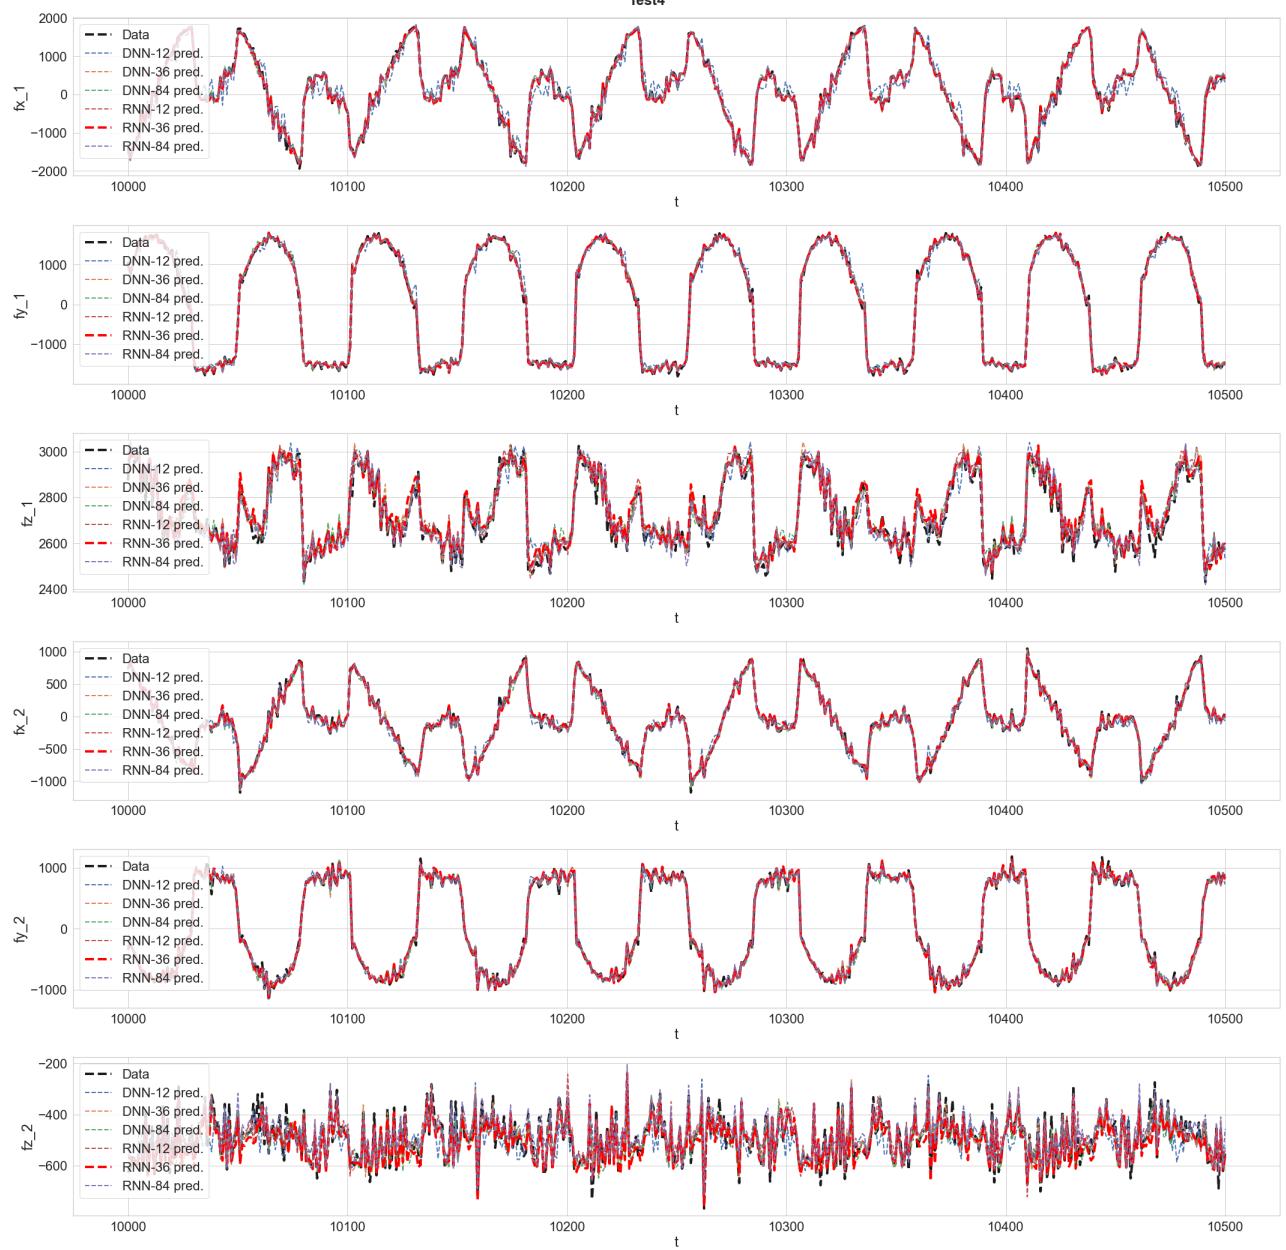
## olsson\_solution

Test2



## olsson\_solution

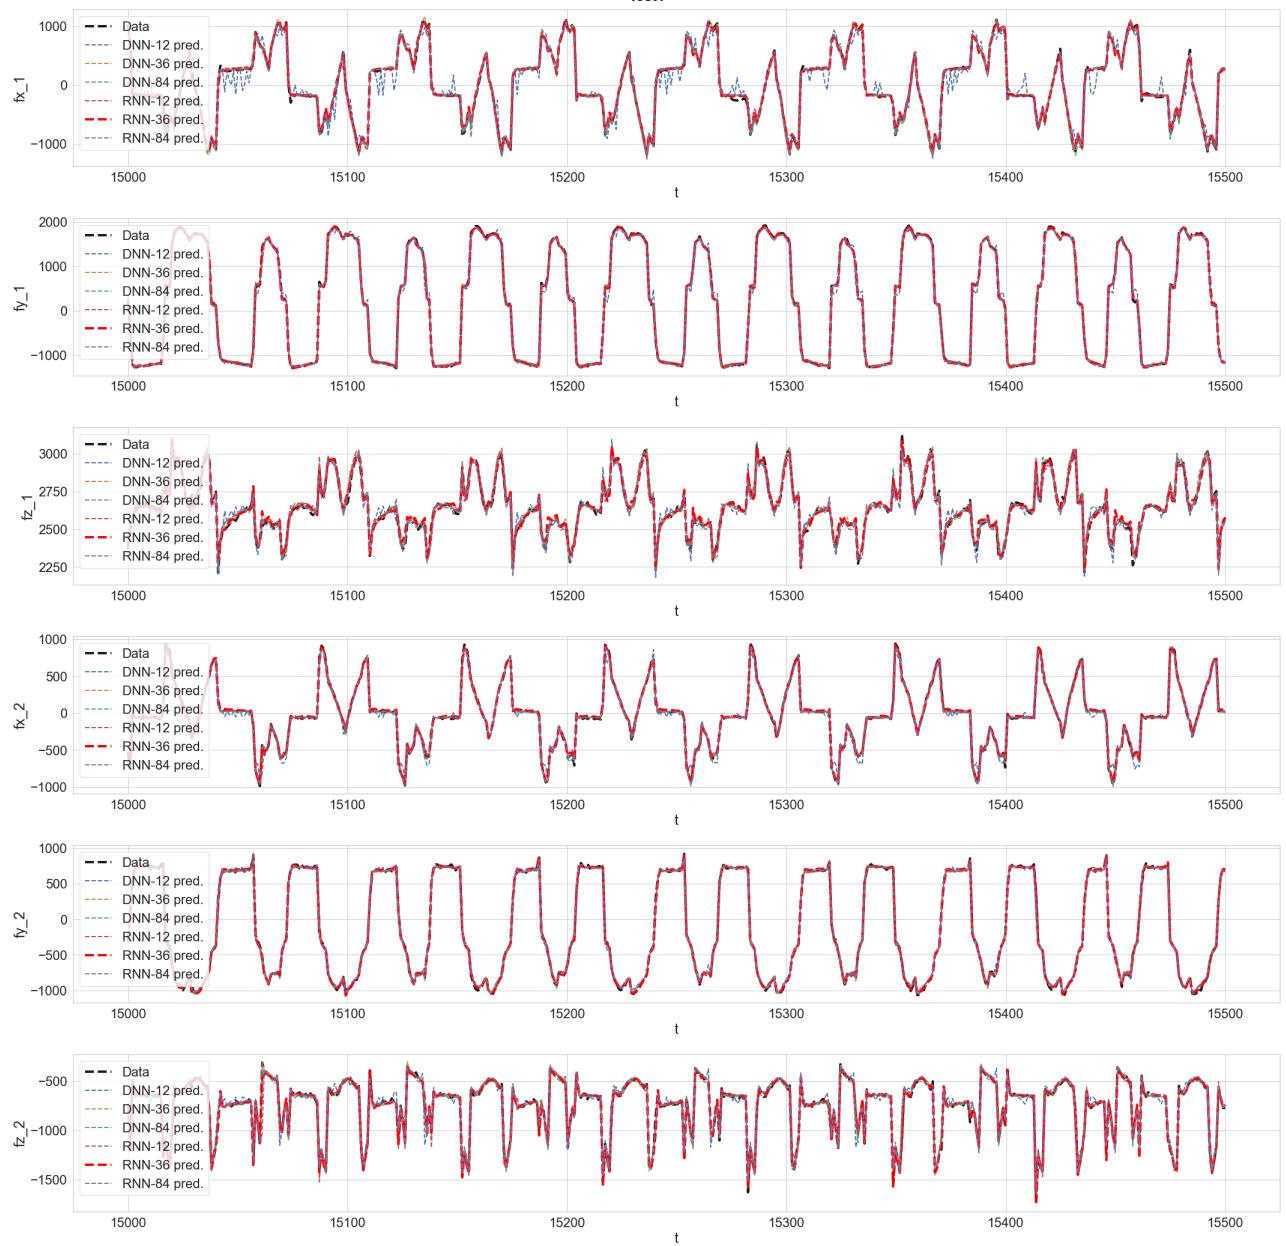
Test4

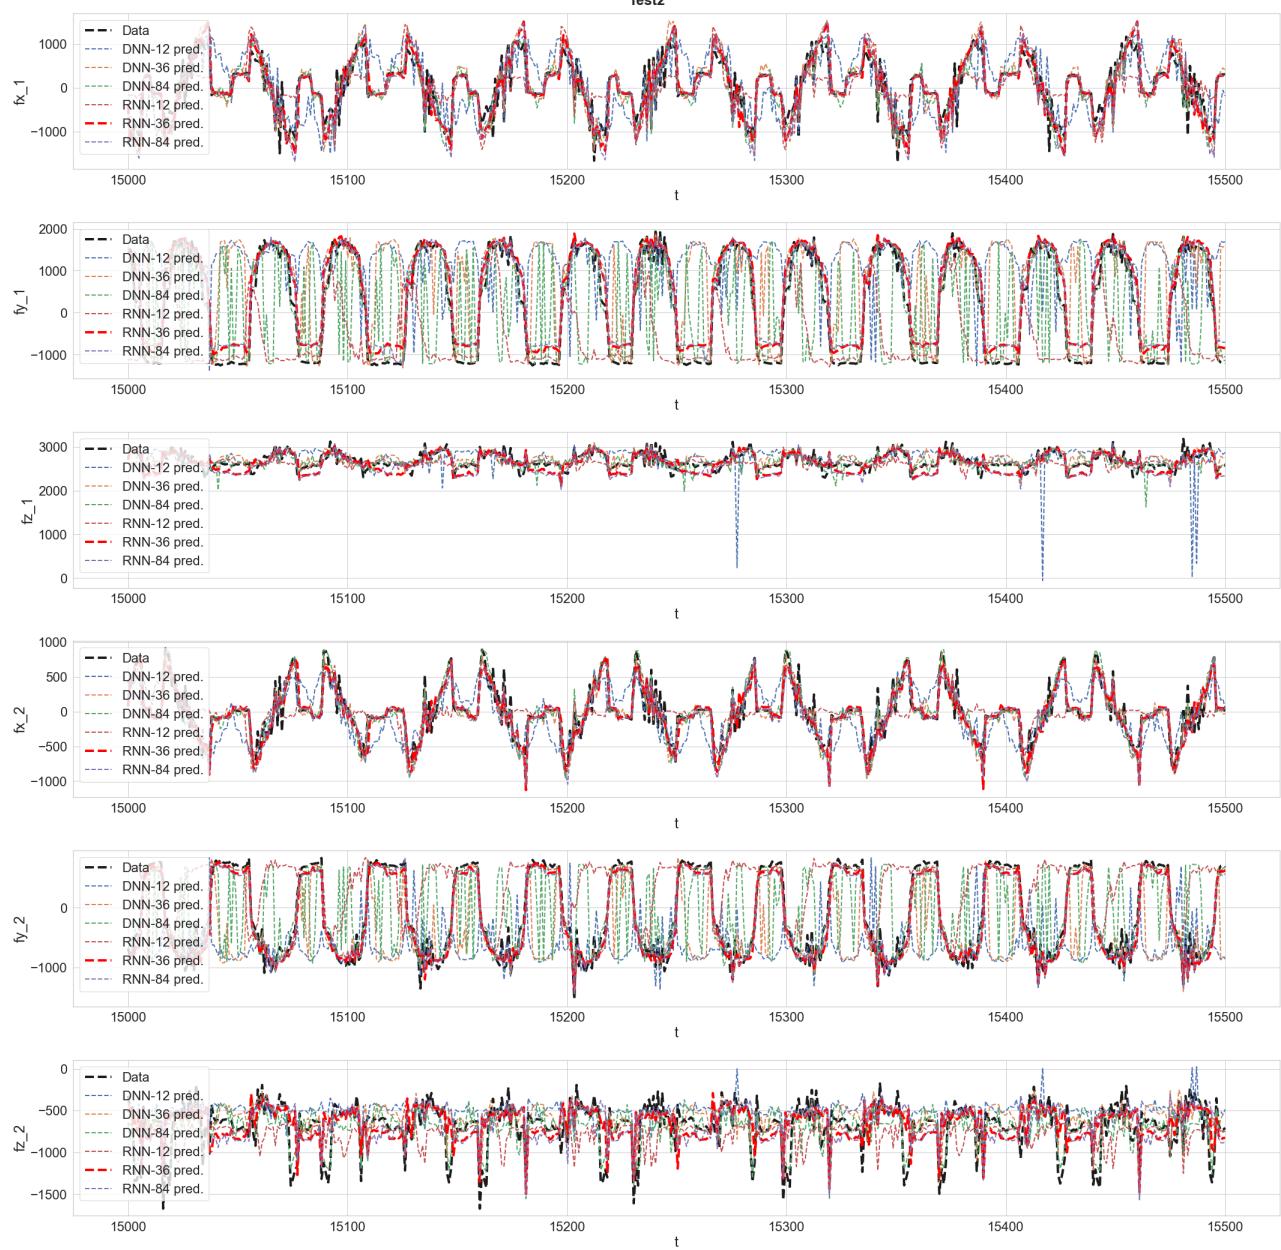


```
In [117]: plot_timeseries(tmin=15000, tmax=15500)
```

## olsson\_solution

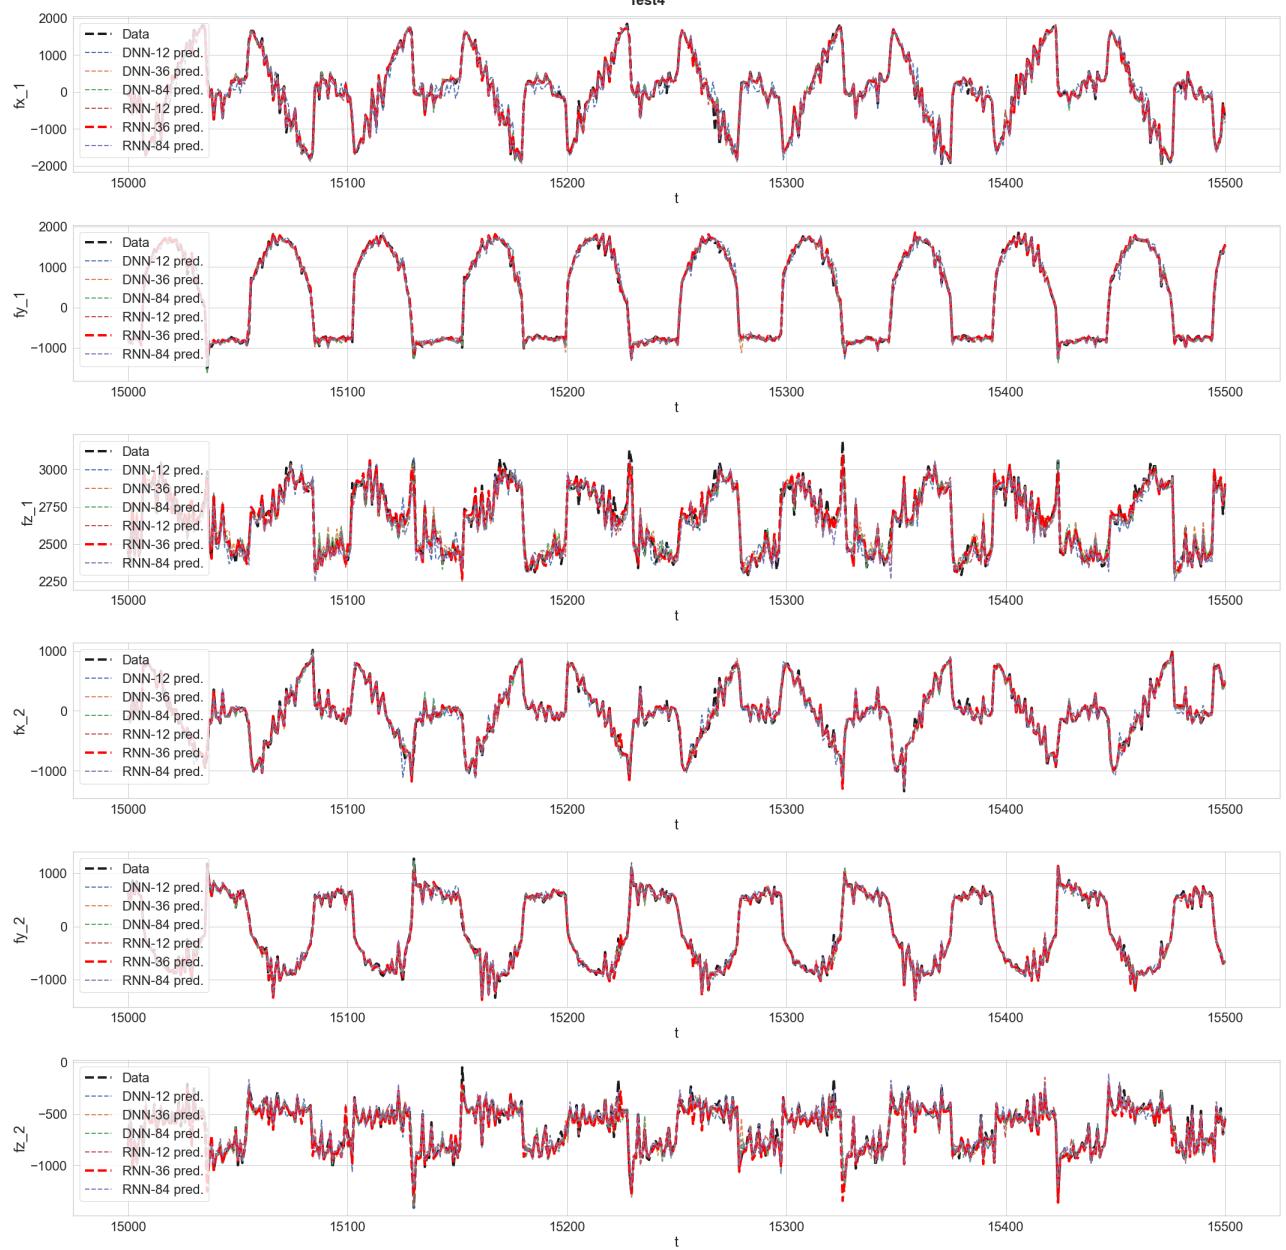
Test1





## olsson\_solution

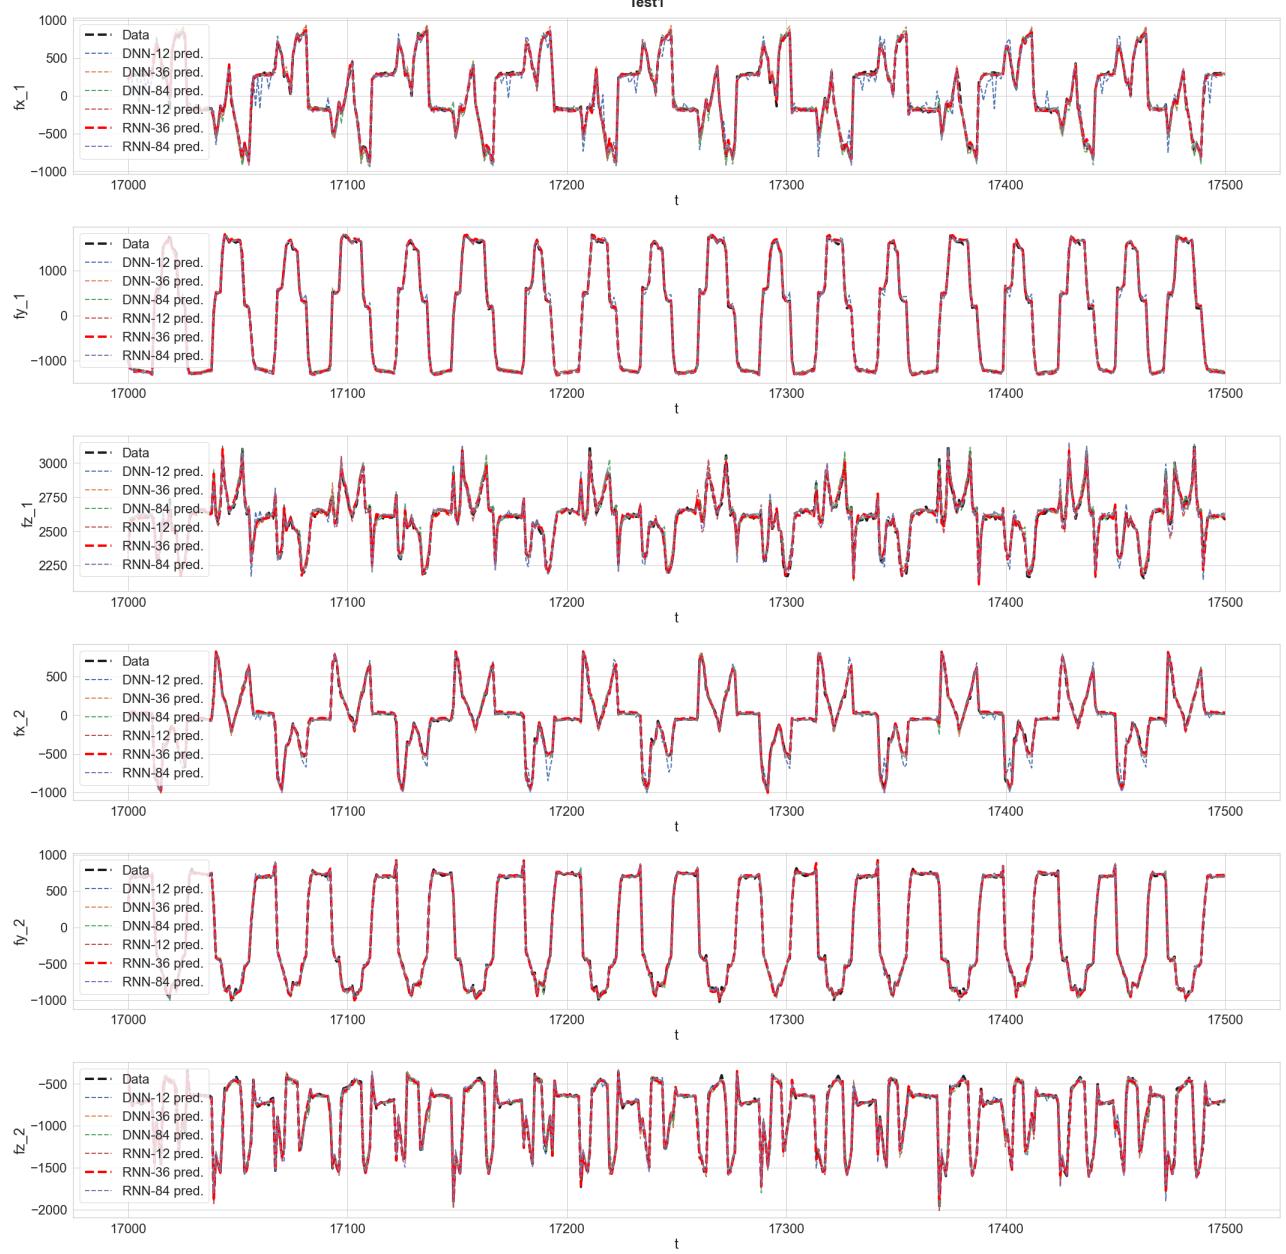
Test4

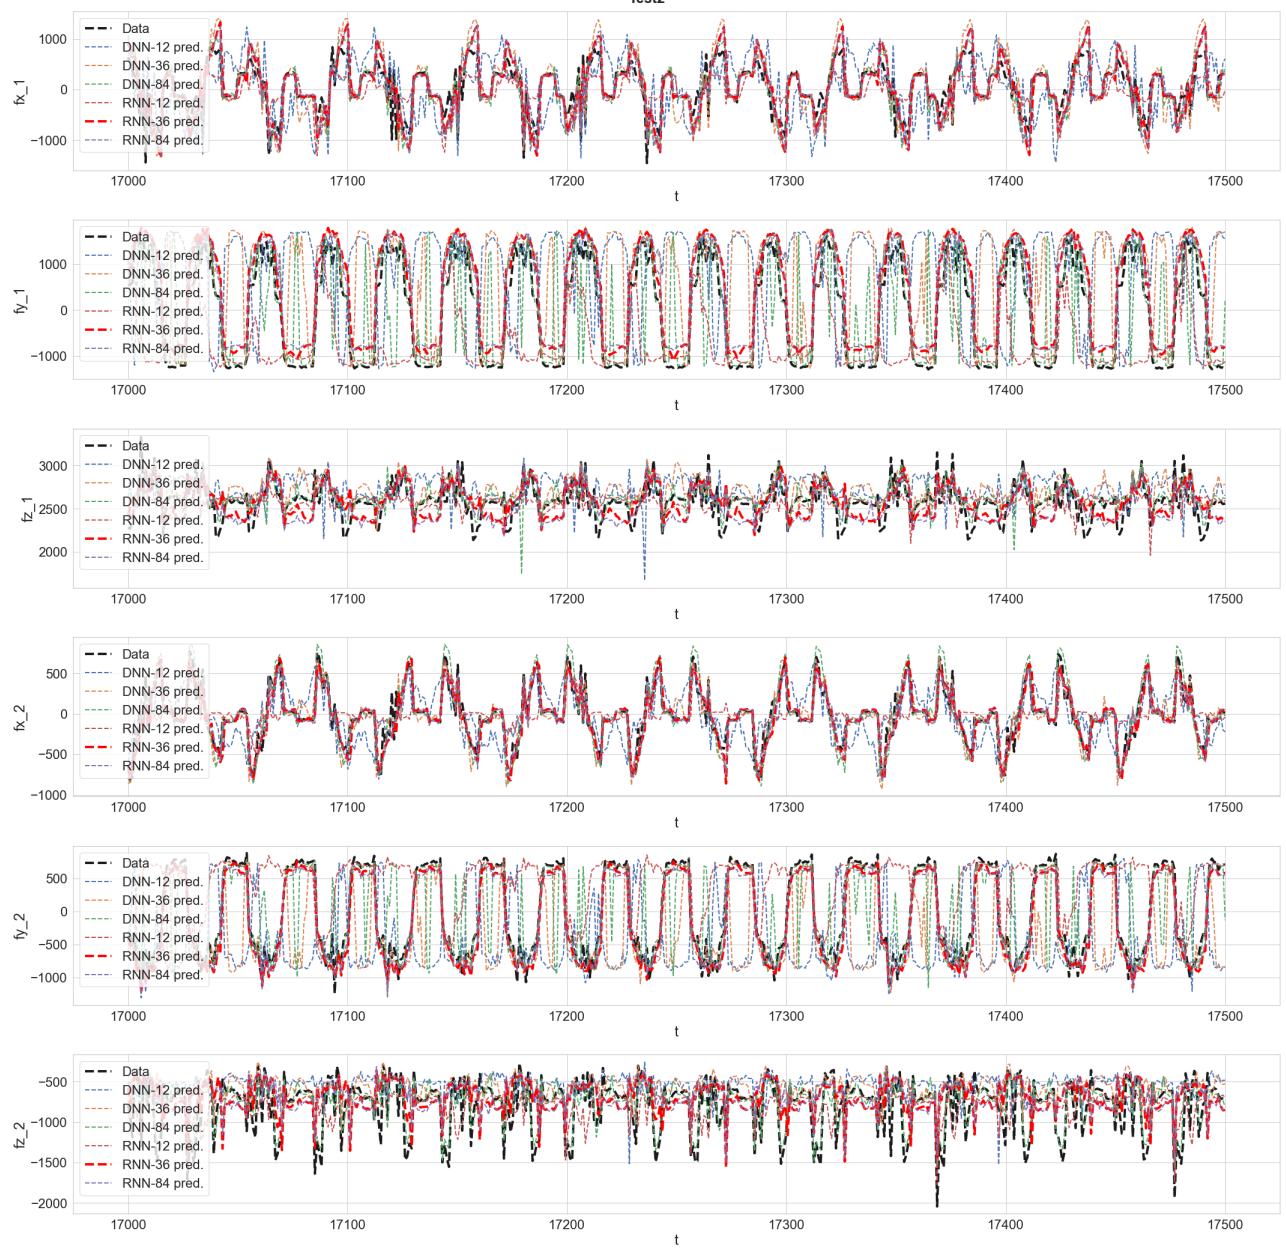


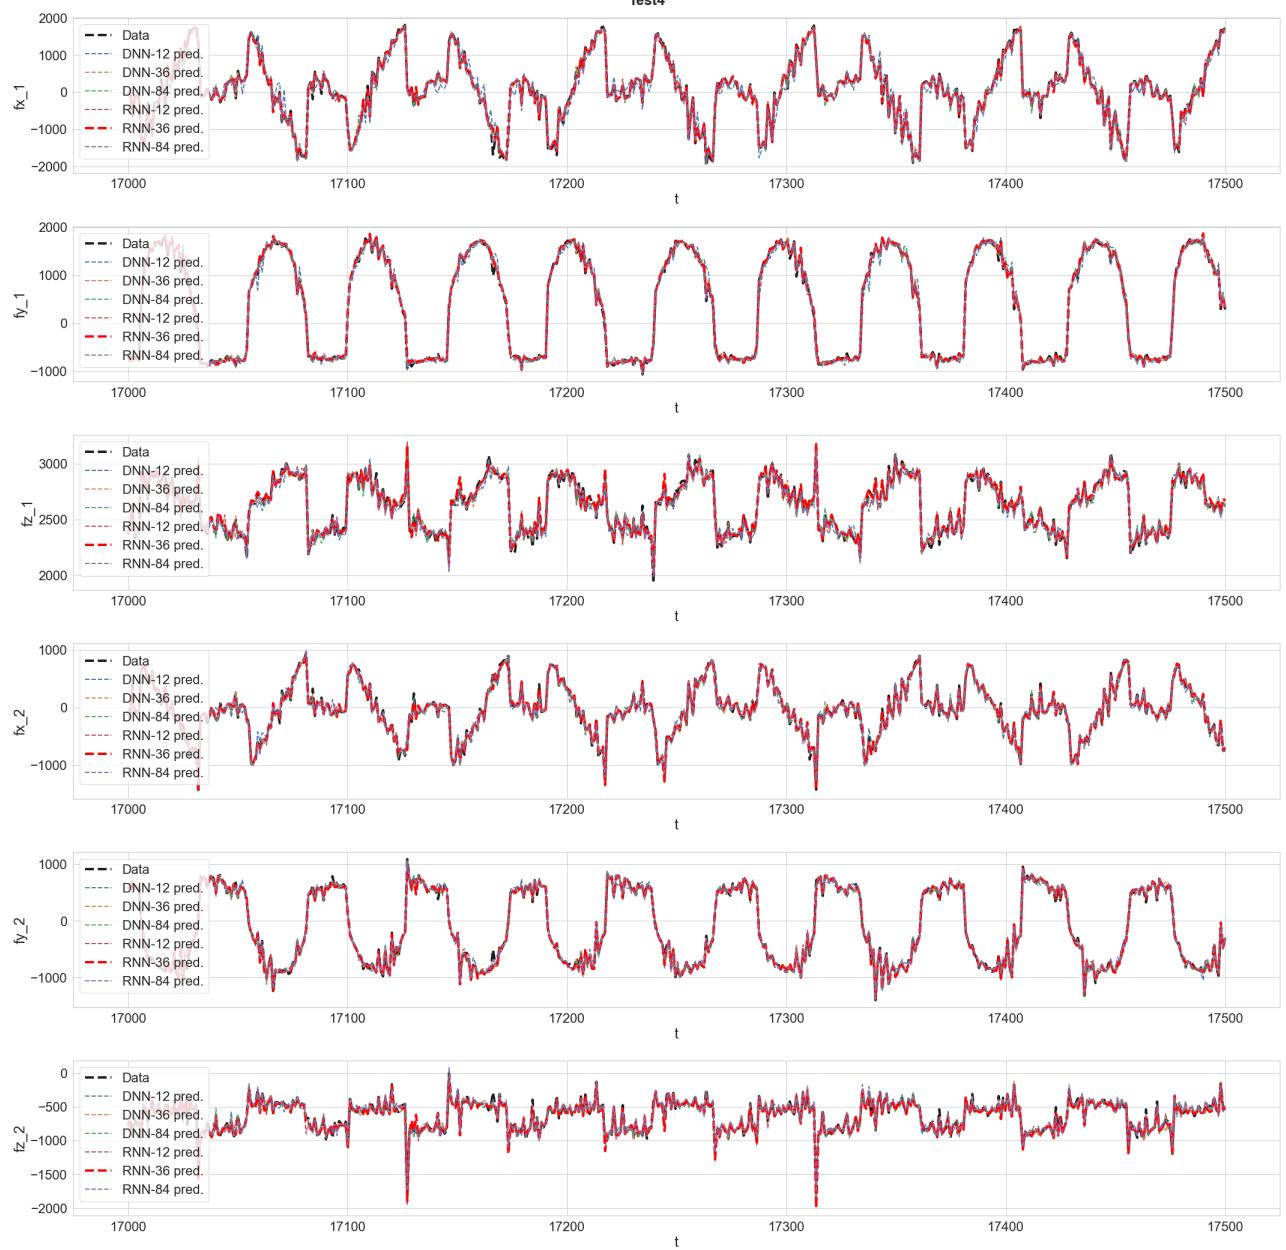
```
In [118]: plot_timeseries(tmin=17000, tmax=17500)
```

## olsson\_solution

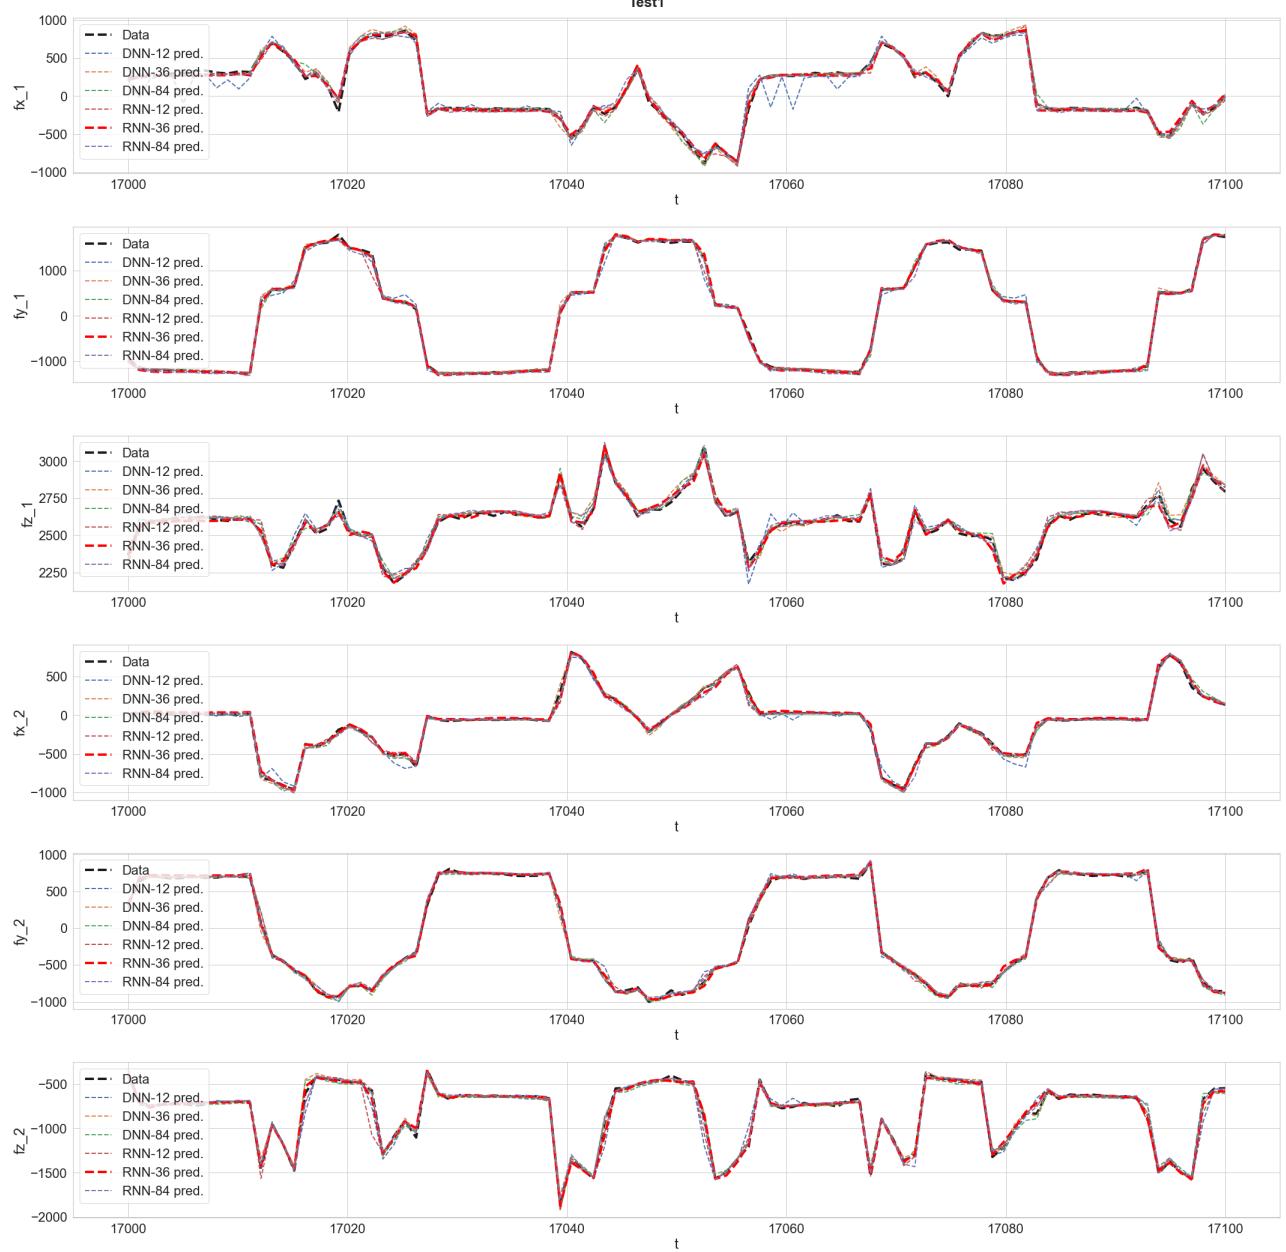
Test1

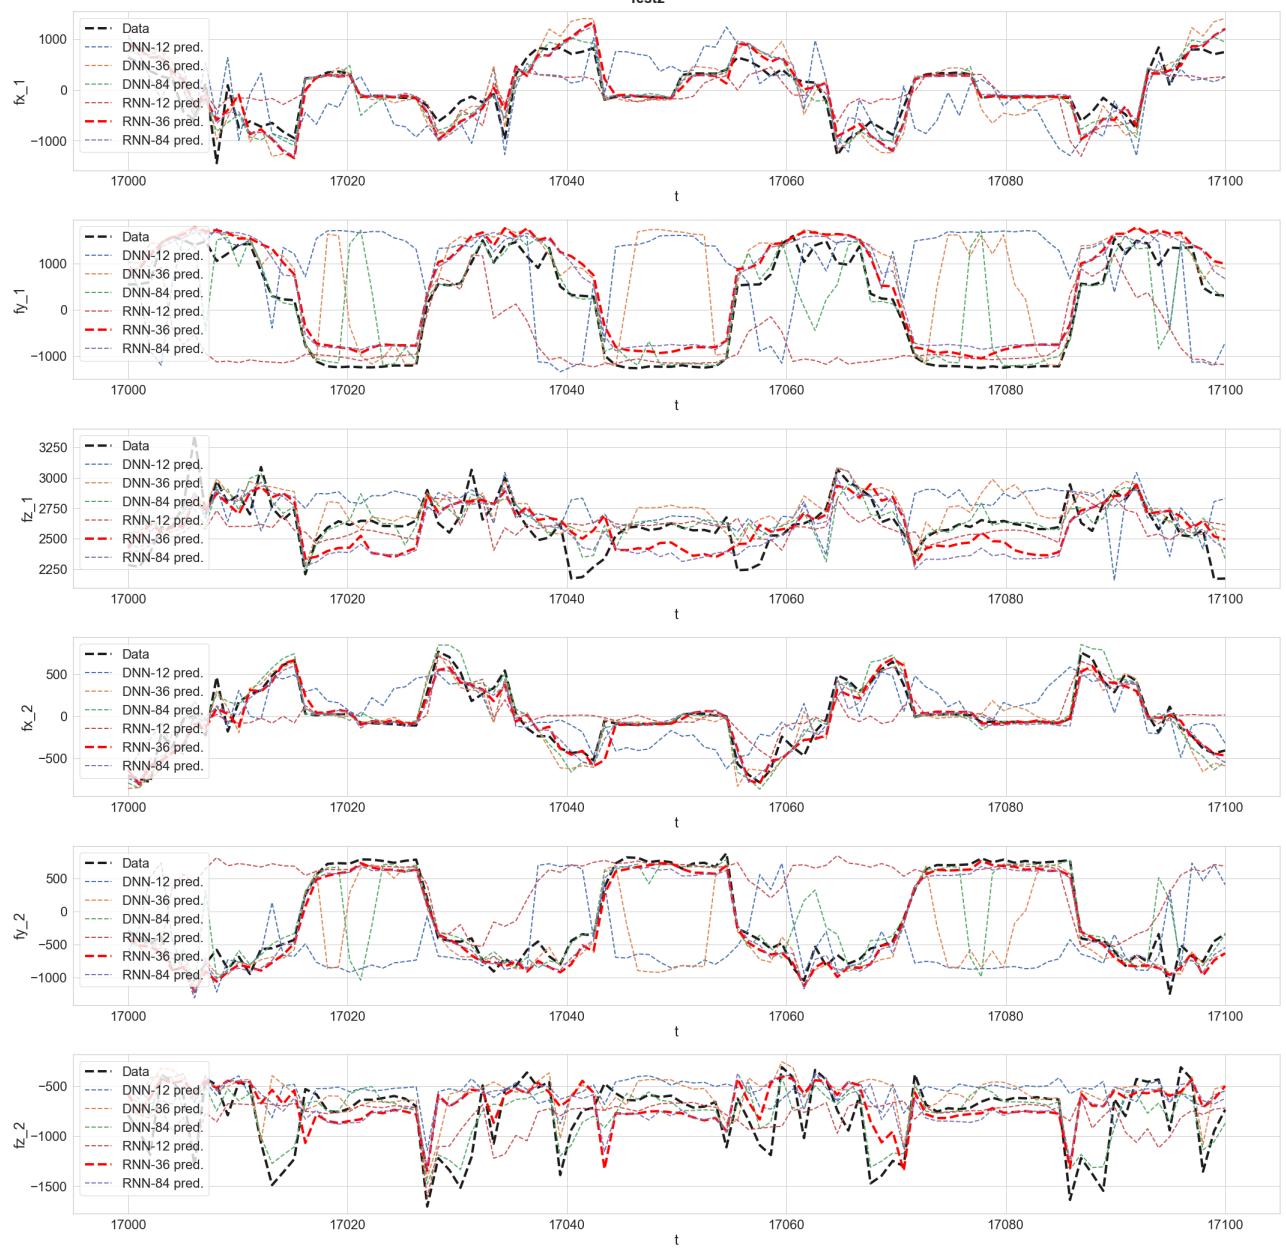


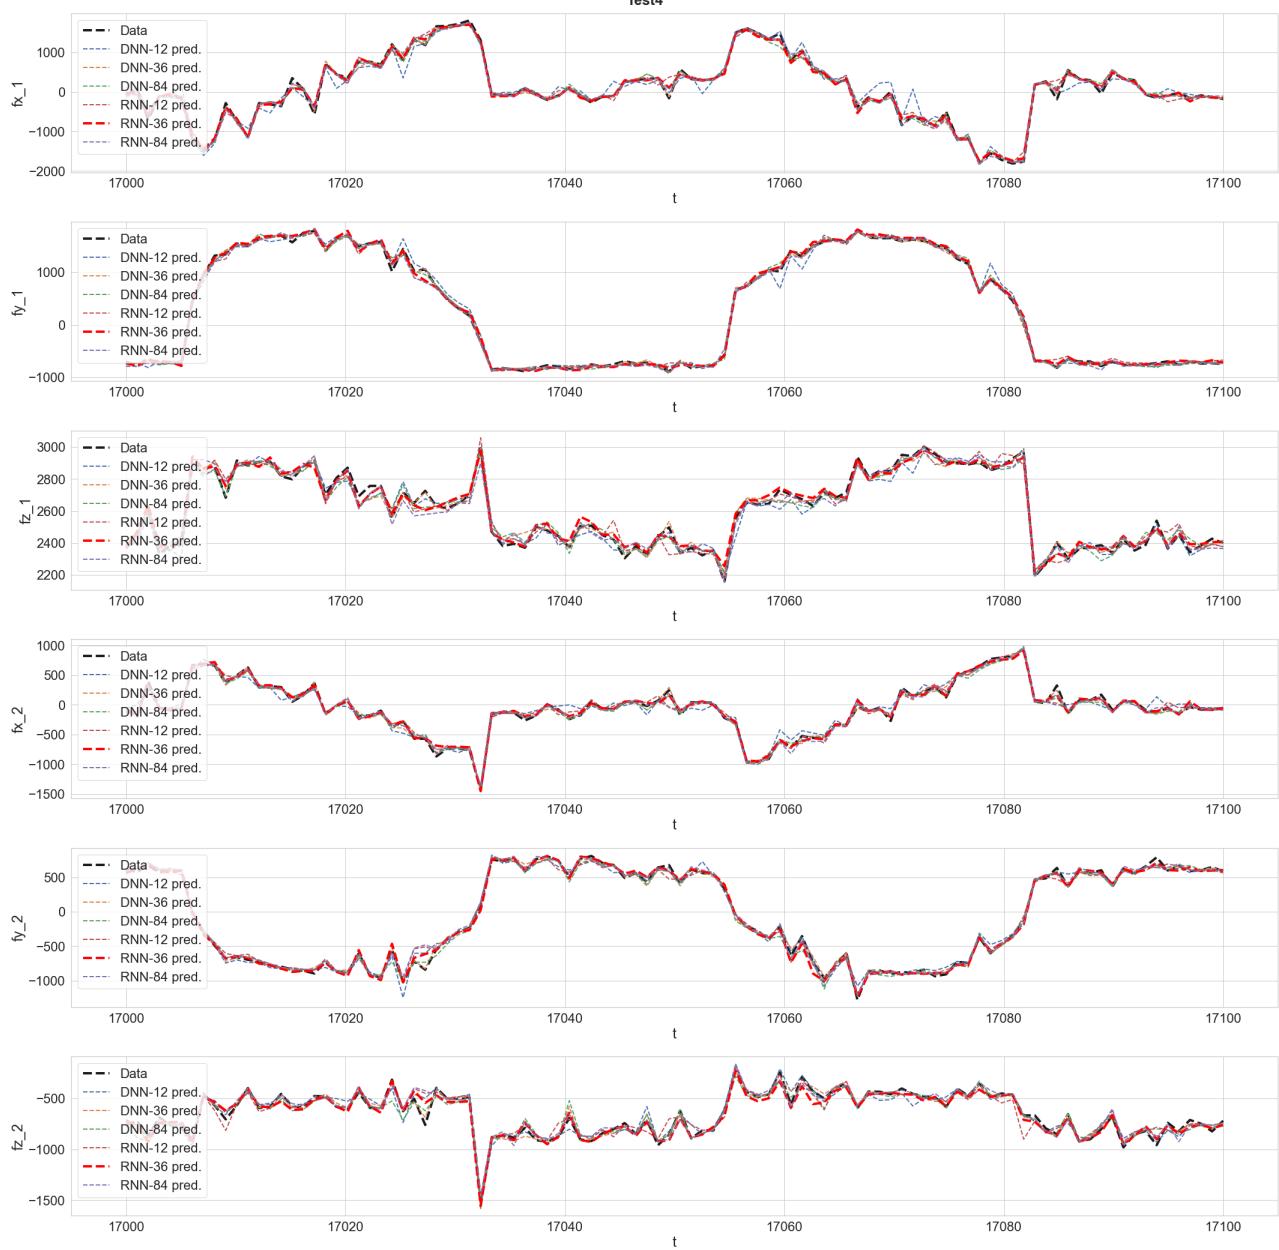




```
In [123]: plot_timeseries(tmin=17000, tmax=17100)
```







In [177]:

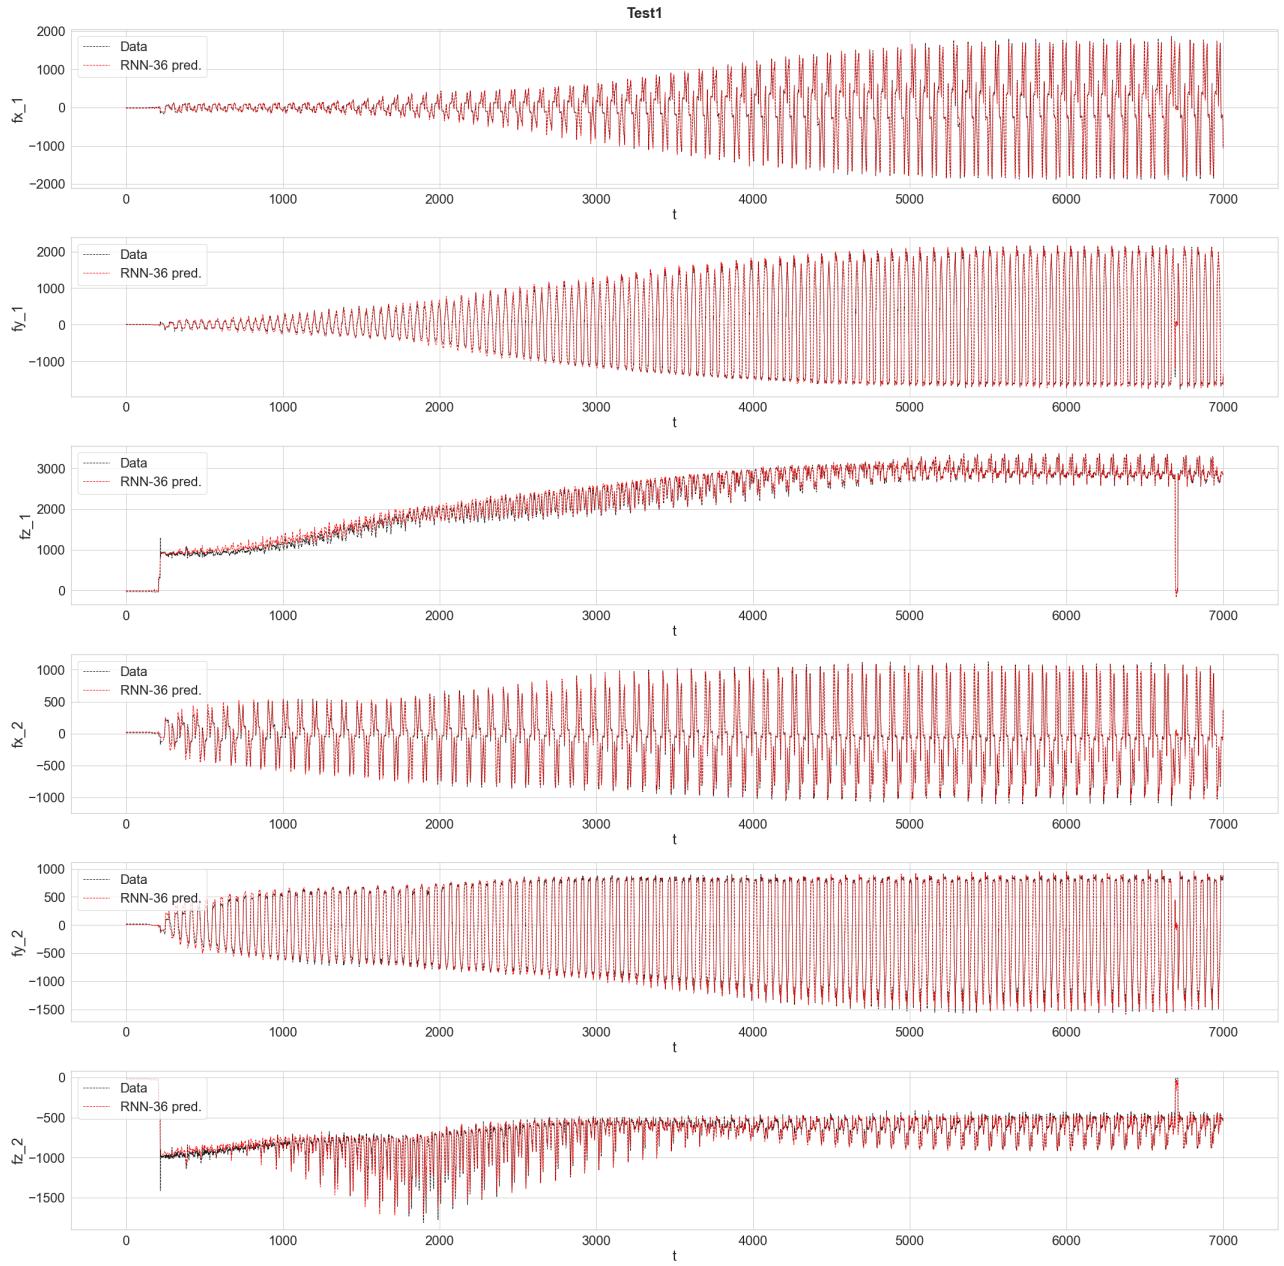
```
# plot large range for best model
tmin = 0
tmax = 7000
t = np.linspace(tmin, tmax, tmax-tmin)
sns.set(font_scale = 2)
sns.color_palette()
sns.set_style("whitegrid")
linewidth=1
for filename in dataset_filenames:

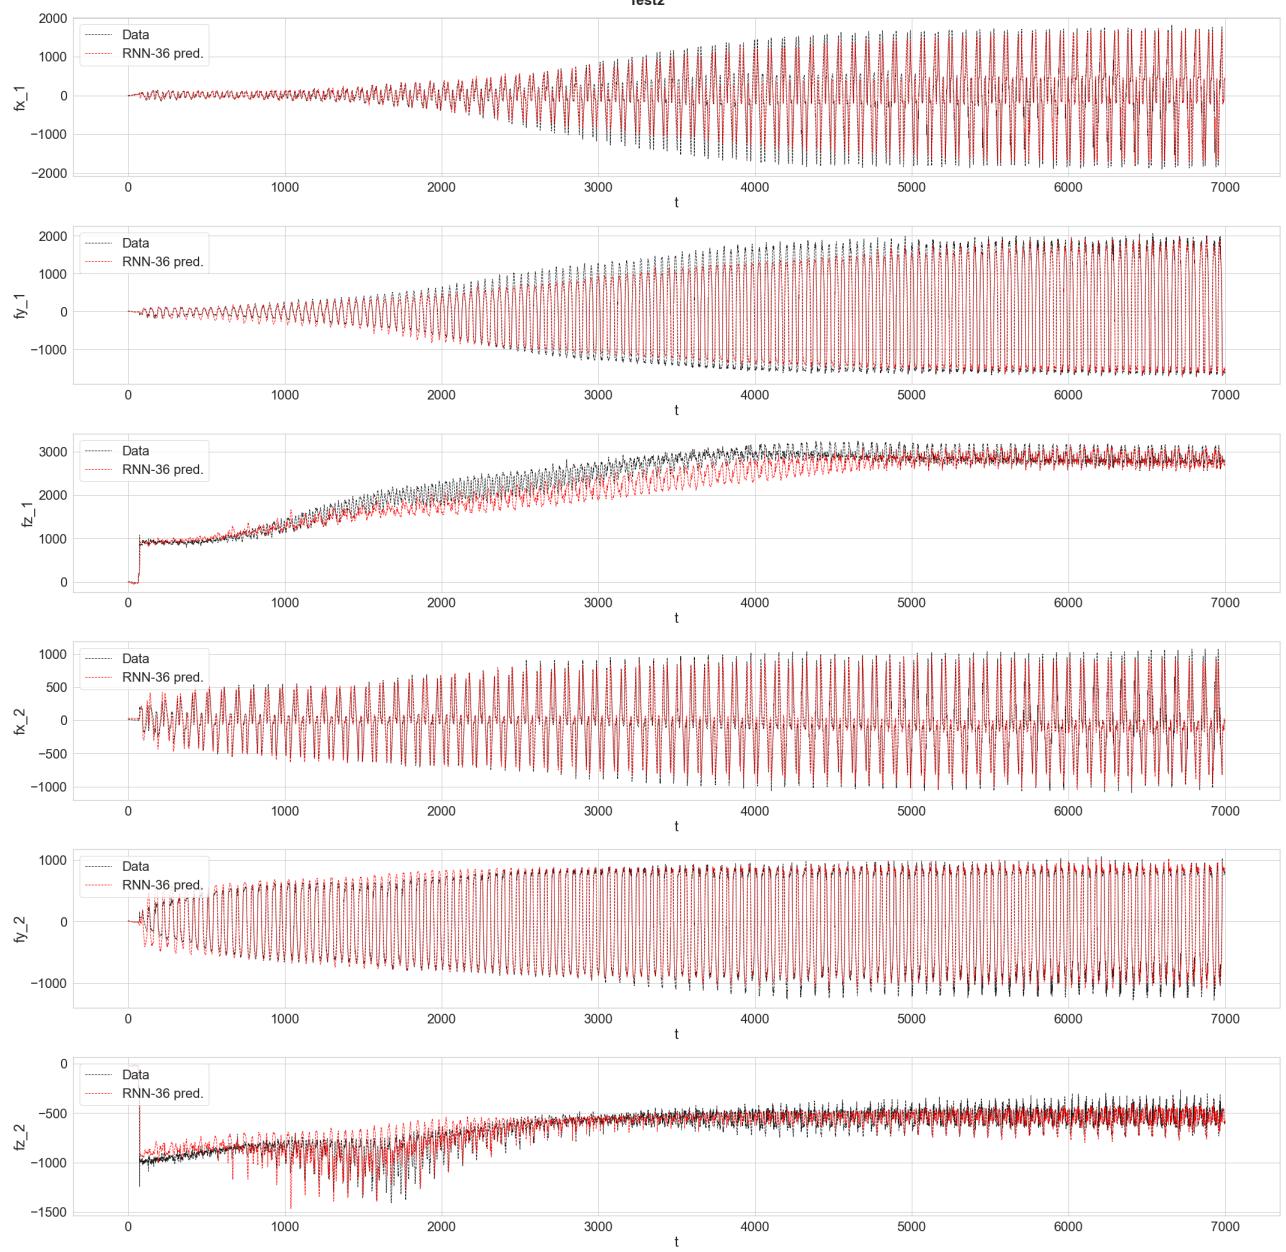
    fig = plt.figure(figsize=(30,30))
    fig.suptitle(filename, weight='bold').set_fontsize('24')
    for i in range(6):
        ax = fig.add_subplot(6, 1, i+1)
        ax.plot(t, tests_12[filename]['Y'].T[i][tmin+n_steps-1:tmax+n_steps-1],
                #ax.plot(t, tests_12[filename]['Y_pred'].T[i][tmin+n_steps-1:tmax+n_step
                #ax.plot(t, tests_36[filename]['Y_pred'].T[i][tmin+n_steps-1:tmax+n_step
                #ax.plot(t, tests_84[filename]['Y_pred'].T[i][tmin+n_steps-1:tmax+n_step
                #ax.plot(t, tests_12[filename]['Y_seq_pred'].T[i][tmin:tmax], label='RNN'
                #ax.plot(t, tests_12[filename]['Y'].T[i][tmin:tmax], label='Data'
```

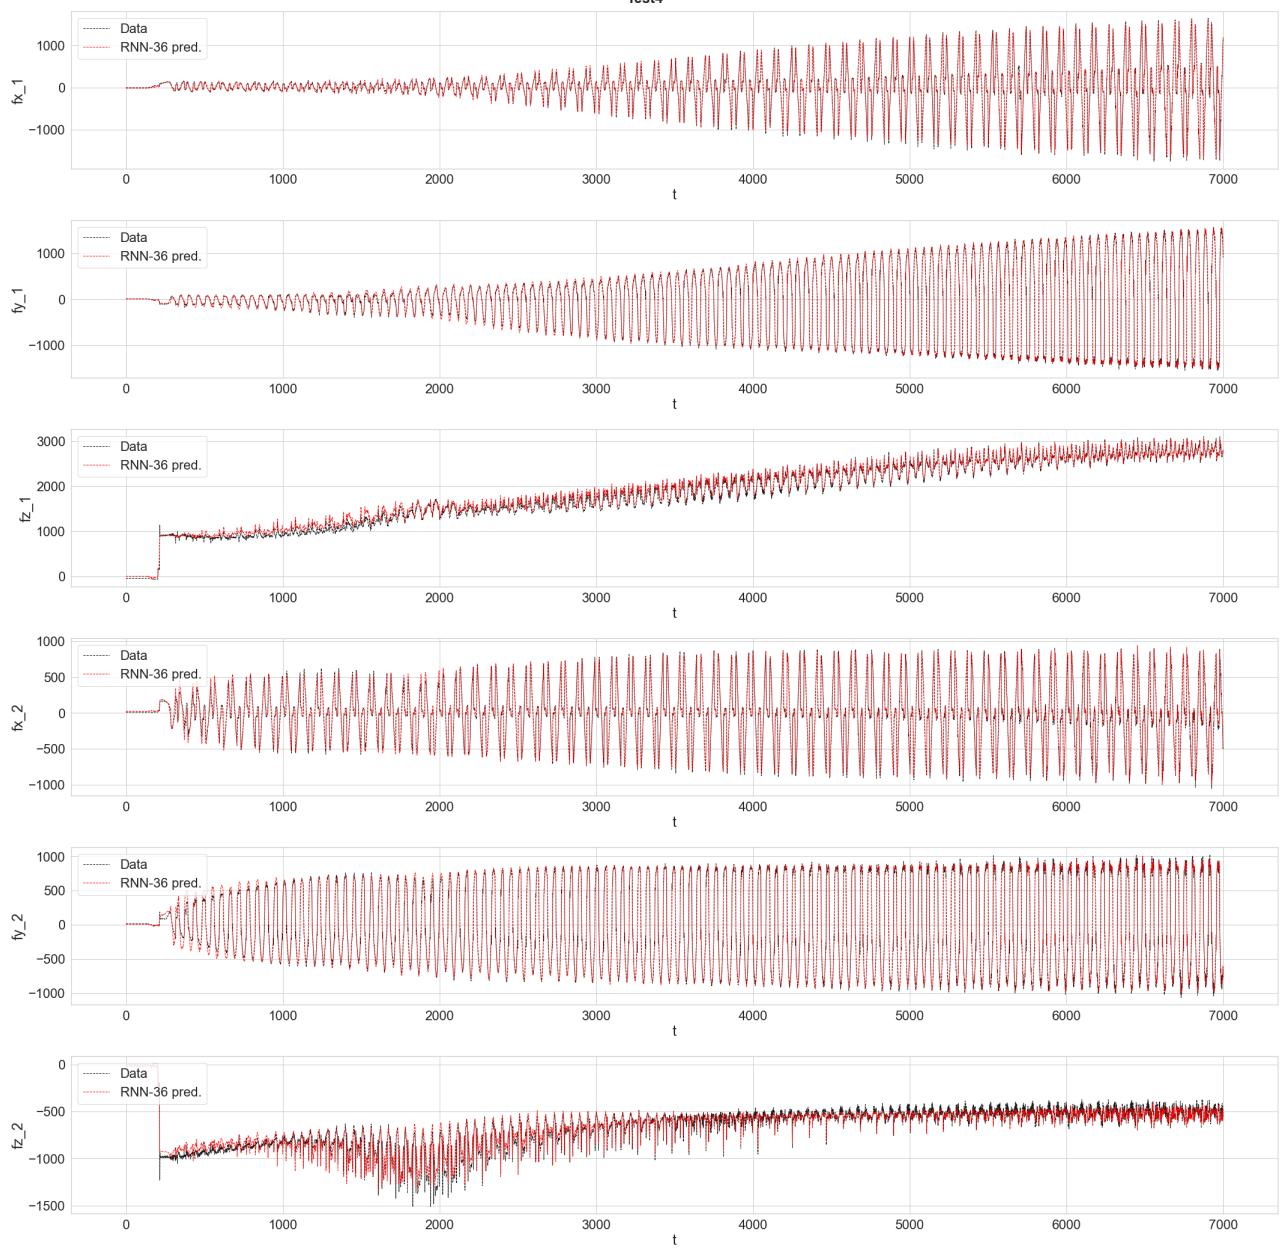
```

ax.plot(t, tests_36[filename]['Y_seq_pred'].T[i][tmin:tmax], label='RNN-36 pred')
#ax.plot(t, tests_84[filename]['Y_seq_pred'].T[i][tmin:tmax], label='RNN-84 pred')
ax.set_xlabel('t')
ax.set_ylabel(outputs[i])
ax.legend(loc=2)
plt.tight_layout()
plt.savefig(output_dir/'{}_timeseries_t{}to{}.pdf'.format(filename, tmin, tmax))

```







## 6. Summary and prospects

Six neural networks (three DNNs and three RNNs) were trained to predict tool-tip forces from input positions and angles (and their higher-order derivatives). The models were trained on 70% of the combined Test1 and Test4 datasets. Test2 was left out of the training entirely to evaluate the performance of the models on unseen runs of the robots.

The RNN models clearly outperformed the DNN models. This is especially clear from looking at the predicted tool-tip forces as a function of time for the Test2 dataset (unseen during training) above.

Adding the first-order (velocity) and second-order (acceleration) derivatives of the positions and angles as additional input features reduced the loss and significantly improved performance. Adding up to 6th order derivatives seemed helpful in some earlier DNN models (with fewer layers) that I trained (not shown in this notebook). However, the performance

comparisons between DNN-84 and DNN-36 in this notebook indicate that the impact is marginal.

Although the RNN did a decent job, all models struggled in generalizing to Test2. A more comprehensive hyperparameter optimization could be done to improve the performance. Adding training data from more runs with robots would probably help a lot.

You would probably want to further optimize and tailor the model to the application. A deeper (and recurrent) model will likely perform better if the goal is to achieve optimal accuracy. However, if the goal is to run fast inference on resource-constrained hardware, you'd want to optimize a smaller model that is *good enough* for the job.